

Please upload your solutions to the designated Moodle assignment on or before the due date as a single zip file using your group id as the file name. Include some brief instructions on how to run your solution to each problem in a file called Problem_X.txt. All solutions may be submitted in groups of up to 3 students.

INVERTED INDEXING & QUERY PROCESSING IN MAPREDUCE

Problem 1.

12 POINTS

- (a) Implement a basic form of an *Inverted Index* (based on the pseudo code provided on Slide #64 of Chapter 3 from the lecture slides) in the Java APIs of Apache Hadoop. Consider the following points:

- Use the basic WordCount.java example as a template for your inverted-index implementation.
- Turn all input lines into lower cases and split the lines using


```
value.toString().split("[^\\w']+");
```

 in order to obtain reasonable word tokens in your map method.
- Enable the job.setCombinerClass(...) option in the driver function (i.e., the main method in the Java class) to enable a local aggregation.
- Consider the Wikipedia-En-41784-Articles.tar.gz archive from Moodle as input to your inverted index. You may assume that each line in the wiki_*.1lineperdoc files corresponds to exactly one Wikipedia article, where the <doc id="X" ...> opening tag denotes an article with id X and the substring between the <doc ...> opening and </doc> closing tags contains the entire text content of that article.
- Extract all article ids X and split the text content of each Wikipedia article into reasonable word tokens as described above.
- The output of your inverted index (at the reduce method) should be emitted by the reducers into one or more part-r-XXXXX files in your Hadoop output directory. The format of these files should be as follows:

```
token<tab>postings
```

where each token is a word token extracted by your tokenizer, and postings is a sorted list of article-id, score pairs of Wikipedia articles in which the token occurs, while score is the frequency of the token in that article. Note that all postings under a same token should be sorted in increasing order of article-id.

5 POINTS

- (b) Implement a simple form of a search engine in the Java APIs of Apache Hadoop by implementing a *Reduce-Side Join* over the part-r-XXXXX files you implemented in (a). Also here, consider the following points:

- This time, use the WordCount2.java example to see how command-line parameters can be passed to the Mapper and Reducer classes, respectively.
- Specifically, pass a user-provided list of keywords (i.e., the “search strings”) from the command-line to the Mapper class.
- Implement a map method which reads one line from each part-r-XXXXX file at a time and extracts the corresponding token, postings fields from each such line.
- For each line that contains a token which occurs in the provided list of keywords (stored in the Mapper class), extract the article-id, score pairs of the corresponding postings list and emit the article-id as the key and the score as the value of the map method.
- In the reduce method, make sure that you then sum up all the score values received under a same article-id as key, and finally emit this article-id together with its summed up score values as one result of your join operation. The format of the final part-r-XXXXX files should be as follows:

```
article-id<tab>sum_of_scores
```

5 POINTS

- (c) Repeat steps (a) and (b) of this exercise by using the *Secondary Sorting* optimization (based on the pseudo code provided on Slide #65 of Chapter 3 from the lectures slides) in the Java APIs of Apache Hadoop. How does this optimization affect the usage of a Combiner and additional Partitioner class? What would be the best join technique for this MapReduce setup?

Also here, the format of the final part-r-XXXXX files should be as follows:

article-id<tab>sum_of_scores

2 POINTS

ANALYTICAL QUERIES IN MAPREDUCE

Problem 2. **12 POINTS**

Consider again the TSV files we downloaded from IMDB for the first exercise sheet. This time, implement the first three analytical queries from Problem 3 (a)–(c) of Exercise Sheet #1 in MapReduce:

- (a) Select the *most popular directors* based on how many movies they directed, stop after the top 25 directors with the most movies. **4 POINTS**
- (b) Select the *top 25 pairs of actors* that occur together in a same movie, ordered by the number of movies in which they co-occur. **4 POINTS**
- (c) Find *frequent 2-itemsets of actors or directors* who occurred together in at least 4 movies. **4 POINTS**

Note that you need to implement a suitable function to parse the various fields from each line of a TSV file into separate strings inside your `map` methods in Java. Splitting each line by `[\t]+` should suffice for our purpose. Also make sure to sort the output of your `reduce` methods in a proper way (e.g., by using an external sorting step as shown in the last exercise sheet). You do not need to consider the A-Priori optimization for solving (c).

ANALYTICAL QUERIES IN APACHE PIG

Problem 3. **12 POINTS**

Also here, consider the TSV files we downloaded from IMDB for the first exercise sheet. This time, implement the first three analytical queries from Problem 3 (a)–(c) of Exercise Sheet #1 in Apache Pig by using the Pig Latin operators presented in the lecture slides:

- (a) Select the *most popular directors* based on how many movies they directed, stop after the top 25 directors with the most movies. **4 POINTS**
- (b) Select the *top 25 pairs of actors* that occur together in a same movie, ordered by the number of movies in which they co-occur. **4 POINTS**
- (c) Find *frequent 2-itemsets of actors or directors* who occurred together in at least 4 movies. **4 POINTS**

Also here, you do not need to consider the A-Priori optimization for solving (c).