

Big Data Analytics

Chapter 4: Text Processing with Latent Semantic Analysis

Following: [3] "**Advanced Analytics with Spark**", Chapter 6

Prof. Dr. Martin Theobald
Faculty of Science, Technology & Communication
University of Luxembourg



How to Analyze Text Data?

Unlike the previous data sets, text data is often considered as "**unstructured data**" – although this is far from actually being true:

- ▶ In fact, natural language follows a very fine-grained structure, with a well-defined grammar and usually a clear intention of an author in describing a possibly complex matter. However, this structure still is quite difficult to process for a machine...

Thus, in the following we focus on a very **basic semantic analysis** of words occurring in documents based on **large-scale co-occurrence statistics**.

We will represent these statistics in the form of a **large** (and usually very **sparse**) **matrix**, where the rows denote words and the columns denote documents.

The cell entries of the matrix are usually based on **TF-IDF weights**:

- ▶ The **term frequency** (TF) of a word is the number of times the word (or its lemma) occurs within a document.
- ▶ The **document frequency** (DF) of a word is the number of documents that contain the word (or its lemma).
- ▶ This coincides with a so-called **bag-of-words representation** for documents in many Information Retrieval and Data Mining applications.

The Wikipedia Data Set

- ▶ We have extracted 500 MB of **English Wikipedia articles** from a recent snapshot and converted the resulting 41,784 articles into a plain-text (UTF-8) format.
- ▶ The dump is split into 1,060 individual files, each of about 500 KB size.
- ▶ Each such split is of the following format:

```
<doc id="38298" url="https://en.wikipedia.org/wiki?curid=38298"  
      title="George Boole">  
      George Boole was an English mathematician, educator, philosopher  
      and logician. ....  
<doc>
```

4.1 Latent Semantic Analysis

Latent Semantic Analysis vs. Latent Semantic Indexing

Latent Semantic Analysis (LSA) and -Indexing (LSI) are mostly synonymous.

- ▶ The general technique of applying a Singular Value Decomposition (SVD) to a **word-document matrix A** is called Latent Semantic Analysis (LSA).
- ▶ By applying SVD to A , similar words are mapped to their **implicit** (i.e., "latent") **meanings**, which are represented as a k -dimensional vector for each word.
- ▶ Conversely, documents are also represented by k -dimensional vectors, which each represent the strength of the **latent concepts** the documents contains.
- ▶ The usage of LSA for **indexing documents** in the context of **Information Retrieval** and/or **Data Mining** is usually called Latent Semantic Indexing.

Basic LSI Idea:

- ▶ When given the following documents:
 - $d_1 = "Jaguar animal cat"$
 - $d_2 = "Jaguar car brand"$
 - $d_3 = "Jaguar luxury vehicle"$
- ▶ The query $q = "car brand"$ should return d_2, d_3, d_1 (in that particular order).

Singular Value Decomposition

Generally, an **SVD of a word-document matrix A** is of the following form:

$$A_{m \times n} = U_{m \times m} S_{m \times n} V^T_{n \times n}$$

- ▶ U is an orthogonal matrix (i.e., $U U^T = I$) whose columns are orthonormal eigenvectors of $A A^T$.
- ▶ S is a diagonal matrix containing the square roots of the eigenvalues of U and V , respectively, *in descending order*.
- ▶ V^T is the transpose of an orthogonal matrix (i.e., $V V^T = I$) whose rows are orthonormal eigenvectors of $A^T A$.

Suppose that A is an $m \times n$ matrix, then a **full SVD decomposition** yields an $m \times m$ matrix U , an $m \times n$ matrix S and an $n \times n$ matrix V^T .

However, by **reducing S** to its k largest singular values, we may achieve the desired **dimensionality reduction** also for U and V^T :

$$A_{m \times n} \approx U_{m \times k} S_{k \times k} V^T_{k \times n}$$

The goal here usually is not to reconstruct A , but to find word-to-word or document-to-document similarities:

- ▶ **Word-to-word similarities**: compare rows $U_{m \times k} S_{k \times k}$
- ▶ **Document-to-document similarities**: compare rows $[S_{k \times k} V^T_{k \times n}]^T = V_{n \times k} S_{k \times k}$

SVD Example: Full Decomposition (I)

$$A = \begin{bmatrix} 2 & 0 & 8 & 6 & 0 \\ 1 & 6 & 0 & 1 & 7 \\ 5 & 0 & 7 & 4 & 0 \\ 7 & 0 & 8 & 5 & 0 \\ 0 & 10 & 0 & 0 & 7 \end{bmatrix} = \begin{bmatrix} -0.54 & 0.07 & 0.82 & -0.11 & 0.12 \\ -0.10 & -0.59 & -0.11 & -0.79 & -0.06 \\ -0.53 & 0.06 & -0.21 & 0.12 & -0.81 \\ -0.65 & 0.07 & -0.51 & 0.06 & 0.56 \\ -0.06 & -0.80 & 0.09 & 0.59 & 0.04 \end{bmatrix}$$
$$\times \begin{bmatrix} 17.92 & 0 & 0 & 0 & 0 \\ 0 & 15.17 & 0 & 0 & 0 \\ 0 & 0 & 3.56 & 0 & 0 \\ 0 & 0 & 0 & 1.98 & 0 \\ 0 & 0 & 0 & 0 & 0.35 \end{bmatrix}$$
$$\times \begin{bmatrix} -0.46 & 0.02 & -0.87 & 0.00 & 0.17 \\ -0.07 & -0.76 & 0.06 & 0.60 & 0.23 \\ -0.74 & 0.10 & 0.28 & 0.22 & -0.56 \\ -0.48 & 0.03 & 0.40 & -0.33 & 0.70 \\ -0.07 & -0.64 & -0.04 & -0.69 & -0.32 \end{bmatrix}$$

- ▶ This full matrix decomposition is indeed **lossless**, i.e., the product $U_{m \times m} S_{m \times n} V^T_{n \times n}$ returns again $A_{m \times n}$.

SVD Example: Full Decomposition (II)

- ▶ Computing **word-to-word similarities** based on

$$W_{m \times n} = U_{m \times m} S_{m \times n} = \begin{bmatrix} -9.72 & 0.99 & 2.93 & -0.21 & 0.04 \\ -1.82 & -9.00 & -0.40 & -1.56 & -0.02 \\ -9.41 & 0.90 & -0.76 & 0.23 & -0.28 \\ -11.56 & 1.07 & -1.81 & 0.12 & 0.20 \\ -1.16 & -12.09 & 0.32 & 1.18 & 0.02 \end{bmatrix}$$

- ▶ and using, e.g., **Cosine similarity**

$$\text{CosSim}(\mathbf{w}_i, \mathbf{w}_j) = \frac{\mathbf{w}_i \cdot \mathbf{w}_j}{\|\mathbf{w}_i\|_2 \|\mathbf{w}_j\|_2}$$

- ▶ we get:

$$\begin{aligned}\text{CosSim}(\mathbf{w}_1, \mathbf{w}_1) &= 1.00 \\ \text{CosSim}(\mathbf{w}_1, \mathbf{w}_2) &= 0.08 \\ \text{CosSim}(\mathbf{w}_1, \mathbf{w}_3) &= 0.93 \\ \text{CosSim}(\mathbf{w}_1, \mathbf{w}_4) &= 0.90 \\ \text{CosSim}(\mathbf{w}_1, \mathbf{w}_5) &= 0.00\end{aligned}$$

- ▶ These values coincide with the pairwise row distances in A .

SVD Example: Full Decomposition (III)

- ▶ Computing **document-to-document similarities** based on

$$\mathbf{D}_{n \times m} = [\mathbf{S}_{m \times n} \mathbf{V}^T_{n \times n}]^T = \begin{bmatrix} -8.33 & -0.33 & -3.10 & -0.00 & 0.06 \\ -1.26 & -11.53 & 0.22 & 1.19 & 0.08 \\ -13.17 & 1.50 & 1.01 & 0.44 & -0.20 \\ -8.68 & 0.39 & 1.42 & -0.66 & 0.25 \\ -1.16 & -9.73 & -0.16 & -1.37 & -0.11 \end{bmatrix}$$

- ▶ and again using **Cosine similarity**

$$\text{CosSim}(\mathbf{d}_i, \mathbf{d}_j) = \frac{\mathbf{d}_i \cdot \mathbf{d}_j}{\|\mathbf{d}_i\|_2 \|\mathbf{d}_j\|_2}$$

- ▶ we get:

$$\begin{aligned}\text{CosSim}(\mathbf{d}_1, \mathbf{d}_1) &= 1.00 \\ \text{CosSim}(\mathbf{d}_1, \mathbf{d}_2) &= 0.06 \\ \text{CosSim}(\mathbf{d}_1, \mathbf{d}_3) &= 0.90 \\ \text{CosSim}(\mathbf{d}_1, \mathbf{d}_4) &= 0.87 \\ \text{CosSim}(\mathbf{d}_1, \mathbf{d}_5) &= 0.08\end{aligned}$$

- ▶ These values in turn coincide with the pairwise column distances in \mathbf{A} .

SVD Example: Reduced Decomposition (I)

$$A = \begin{bmatrix} 2 & 0 & 8 & 6 & 0 \\ 1 & 6 & 0 & 1 & 7 \\ 5 & 0 & 7 & 4 & 0 \\ 7 & 0 & 8 & 5 & 0 \\ 0 & 10 & 0 & 0 & 7 \end{bmatrix} \approx \left[\begin{array}{ccccc|cc} -0.54 & 0.07 & 0.82 & -0.11 & 0.12 \\ -0.10 & -0.59 & -0.11 & -0.79 & -0.06 \\ -0.53 & 0.06 & -0.21 & 0.12 & -0.81 \\ -0.65 & 0.07 & -0.51 & 0.06 & 0.56 \\ -0.06 & -0.80 & 0.09 & 0.59 & 0.04 \end{array} \right] \quad k = 3$$
$$\times \begin{bmatrix} 17.92 & 0 & 0 & 0 & 0 \\ 0 & 15.17 & 0 & 0 & 0 \\ 0 & 0 & 3.56 & 0 & 0 \\ \hline 0 & 0 & 0 & 1.98 & 0 \\ 0 & 0 & 0 & 0 & 0.35 \end{bmatrix}$$
$$\times \begin{bmatrix} -0.46 & 0.02 & -0.87 & 0.00 & 0.17 \\ -0.07 & -0.76 & 0.06 & 0.60 & 0.23 \\ \hline -0.74 & 0.10 & 0.28 & 0.22 & -0.56 \\ -0.48 & 0.03 & 0.40 & -0.33 & 0.70 \\ -0.07 & -0.64 & -0.04 & -0.69 & -0.32 \end{bmatrix}$$

- This reduced matrix decomposition now becomes "**lossy**", i.e., the product $U_{m \times k} S_{k \times k} V^T_{k \times n}$ returns only an approximation of $A_{m \times n}$.

SVD Example: Reduced Decomposition (II)

- ▶ Computing **word-to-word similarities** based on

$k = 3$

$$\mathbf{W}_{m \times k} = \mathbf{U}_{m \times k} \mathbf{S}_{k \times k} = \begin{bmatrix} -9.72 & 0.99 & 2.93 \\ -1.82 & -9.00 & -0.40 \\ -9.41 & 0.90 & -0.76 \\ -11.56 & 1.07 & -1.81 \\ -1.16 & -12.09 & 0.32 \end{bmatrix}$$

- ▶ and using, e.g., **Cosine similarity**

$$\text{CosSim}(\mathbf{w}_i, \mathbf{w}_j) = \frac{\mathbf{w}_i \cdot \mathbf{w}_j}{\|\mathbf{w}_i\|_2 \|\mathbf{w}_j\|_2}$$

- ▶ we get:

$$\begin{aligned}\text{CosSim}(\mathbf{w}_1, \mathbf{w}_1) &= 1.00 \\ \text{CosSim}(\mathbf{w}_1, \mathbf{w}_2) &= 0.08 \\ \text{CosSim}(\mathbf{w}_1, \mathbf{w}_3) &= 0.93 \\ \text{CosSim}(\mathbf{w}_1, \mathbf{w}_4) &= 0.90 \\ \text{CosSim}(\mathbf{w}_1, \mathbf{w}_5) &= 0.00\end{aligned}$$

- ▶ These values now approximate the pairwise row distances in \mathbf{A} .

SVD Example: Reduced Decomposition (III)

- ▶ Computing **document-to-document similarities** based on

$k = 3$

$$\mathbf{D}_{n \times k} = [\mathbf{S}_{k \times k} \ \mathbf{V}^T_{k \times n}]^T = \begin{bmatrix} -8.33 & 1.06 & -2.62 \\ 0.39 & -11.53 & 0.35 \\ -15.56 & 0.96 & 1.01 \\ -0.01 & 9.12 & 0.80 \\ 3.07 & 3.46 & -2.01 \end{bmatrix}$$

- ▶ and again using **Cosine similarity**

$$\text{CosSim}(\mathbf{d}_i, \mathbf{d}_j) = \frac{\mathbf{d}_i \cdot \mathbf{d}_j}{\|\mathbf{d}_i\|_2 \|\mathbf{d}_j\|_2}$$

- ▶ we get:

$$\text{CosSim}(\mathbf{d}_1, \mathbf{d}_1) = 1.00$$

$$\text{CosSim}(\mathbf{d}_1, \mathbf{d}_2) = 0.08$$

$$\text{CosSim}(\mathbf{d}_1, \mathbf{d}_3) = 0.92$$

$$\text{CosSim}(\mathbf{d}_1, \mathbf{d}_4) = -0.14$$

$$\text{CosSim}(\mathbf{d}_1, \mathbf{d}_5) = -0.54$$

- ▶ These values now approximate the pairwise column distances in \mathbf{A} .

Processing Queries

- ▶ A **new query vector** $\mathbf{q}_{1 \times m}$ can be processed against the reduced decomposition by transforming it in a similar manner, namely by computing

$$\mathbf{q}'_{1 \times k} = \mathbf{q}_{1 \times m} \times \mathbf{U}_{m \times k} \times \mathbf{S}_{k \times k}^{-1}$$

- ▶ and then by computing the pairwise Cosine similarities of $\mathbf{q}'_{1 \times k}$ with the row vectors in $\mathbf{D}_{n \times k}$ which each represents a reduced document vector.

Example:

- ▶ For the query vector $\mathbf{q}_{1 \times m} = [1, 3, 0, 2, 0]$, we obtain the **transformed and reduced query vector** $\mathbf{q}'_{1 \times k} = [-0.12, -0.10, -0.15]$ with the following Cosine similarities to the row vectors in $\mathbf{D}_{n \times k} = [\mathbf{S}_{k \times k} \mathbf{V}_{k \times n}^T]^T$ (see previous slide):

$$\begin{aligned}\text{CosSim}(\mathbf{q}', \mathbf{d}_1) &= 0.78 \\ \text{CosSim}(\mathbf{q}', \mathbf{d}_2) &= 0.44 \\ \text{CosSim}(\mathbf{q}', \mathbf{d}_3) &= 0.47 \\ \text{CosSim}(\mathbf{q}', \mathbf{d}_4) &= -0.53 \\ \text{CosSim}(\mathbf{q}', \mathbf{d}_5) &= -0.39\end{aligned}$$

$k = 3$

Finding an SVD of a Word-Document Matrix (I)

That is, just like the matrix factorization shown in Chapter 3, we can compute the SVD of a matrix A by **solving a set of linear equations**:

1. Compute the k largest (non-zero) eigenvalues of either $A A^T$ or $A^T A$.
2. To find U :
 - Use these eigenvalues to compute k eigenvectors of $A A^T$.
 - Turn each eigenvector into a column vector of a new matrix U' .
 - Convert U' into the orthogonal matrix U (e.g., using the [Gram–Schmidt](#) procedure)
3. To find V^T :
 - Use these eigenvalues to compute k eigenvectors of $A^T A$.
 - Turn each eigenvector into a row vector of a new matrix $V^{T'}$.
 - Convert $V^{T'}$ into the orthogonal matrix V^T (e.g., using the [Gram–Schmidt](#) procedure)
4. To find S :
 - Populate a new $k \times k$ diagonal matrix with the square roots of the non-zero eigenvalues computed in step 1.

Finding an SVD of a Word-Document Matrix (II)

- ▶ Spark imports a number of low-level Fortran libraries for matrix operations ([ARPACK](#)) that provide different numerical optimizations also for **approximate eigenvalue decompositions**, including SVD.
- ▶ See also [here](#) for a **tutorial on LSA** via SVD and eigenvalue decomposition.
- ▶ See the previous chapter on possible ways to implement **distributed matrix multiplications** (Methods 1 & 2) in Spark.

4.2 Text Analysis & Singular-Value-Decomposition in Spark's MLlib

Loading the Data Set (I)

- ▶ First of all, we need to define a special loading function for our new Wikipedia file format.
- ▶ Recall that a Wikipedia article is delimited by a pair of `<doc ...>` and `</doc>` tags, such that a simple line-based input format would not guarantee that consecutive lines in the input documents are also stored as consecutive lines generated by the `sc.textFile(...)` RDD.
- ▶ On the other hand, we do not want to assume that all Wikipedia files fit into the main-memory of our driver node, hence we should keep working with partitioned RDDs from the beginning.
- ▶ Fortunately, there is an alternative function that loads whole text files into an RDD:

```
sc.wholeTextFiles("./pathToTextFiles")
```

Loading the Data Set (II)

- ▶ We can then transform the RDD with the input files (each containing many Wikipedia articles) into a new RDD containing **one entry per Wikipedia article** by a single `flatMap` transformation:

```
val textFiles = sc.wholeTextFiles("./path-to-wiki-articles")
val plainText = textFiles.flatMap{ case (uri, text) =>
    parse(text.split("\n")) } // contains one entry for each
                           Wikipedia article
```

- ▶ Note that the `parse` function (see Moodle) is executed on all partitions of input files **in parallel**.
- ▶ This way, we do **not need to assume that all articles fit into the main memory of our driver node** at any time. Every processing step is performed only via RDD transformations!

Create a Default NLP Pipeline

- ▶ Our basic **natural-language-processing** (NLP) **pipeline** consists of:
 - ▶ Sentence splitting
 - ▶ Tokenization
 - ▶ Part-Of-Speech (POS) tagging
 - ▶ Lemmatization
- ▶ The **Stanford CoreNLP tools** provide a very comprehensive API for these (and other) NLP tasks:

```
import edu.stanford.nlp.pipeline._  
import edu.stanford.nlp.ling.CoreAnnotations._  
  
def createNLPPipeline(): StanfordCoreNLP = {  
    val props = new Properties()  
    props.put("annotators", "tokenize, ssplit, pos, lemma")  
    new StanfordCoreNLP(props)  
}
```

Filtering Out Non-Letter Tokens & Stopwords

- ▶ For a more compact dictionary of words, we may still want to filter out tokens that are **not proper sequences of letters** and very frequent words, the so-called **stopwords**.

```
def isOnlyLetters(str: String): Boolean = {  
    str.forall(c => Character.isLetter(c)) }
```

- ▶ Stopwords are extremely frequent words of a language, such as **prepositions**, **articles**, **auxiliary verbs**, etc., that do not add much information to our semantic indexing approach.

```
import scala.io.Source._  
val bStopWords = sc.broadcast(  
    fromFile("stopwords.txt").getLines().toSet)
```

- ▶ A respective file with English stopwords is available on Moodle.

From Text to Word Lemmas

- ▶ An input text of type `String` should be tokenized and the words should be reduced to their lemmatized form:

```
def plainTextToLemmas(text: String,  
                      pipeline: StanfordCoreNLP): Seq[String] = {  
    val doc = new Annotation(text)  
    pipeline.annotate(doc)  
    val lemmas = new ArrayBuffer[String]()  
    val sentences = doc.get(classOf[SentencesAnnotation])  
    for (sentence <- sentences;  
         token <- sentence.get(classOf[TokensAnnotation])) {  
        val lemma = token.get(classOf[LemmaAnnotation])  
        if (lemma.length > 2 && !bStopWords.value.contains(lemma)  
            && isOnlyLetters(lemma)) {  
            lemmas += lemma.toLowerCase } }  
    lemmas }
```

Initialize the NLP Pipelines

- ▶ Use a `mapPartitions` transformation on the `plainText` RDD so that we only **initialize** the NLP pipeline object **once per partition** instead of once per document:

```
val lemmatized: RDD[Seq[String]] =  
    plainText.mapPartitions(it => {  
        val pipeline = createNLPPipeline()  
        it.map { case(title, contents) =>  
            plainTextToLemmas(contents, pipeline)  
        }  
    })
```

Computing TF Weights

- ▶ The term frequency of a term in a document is a local property of that document. So we can compute it via a **local aggregation** over the documents.

```
import scala.collection.mutable.HashMap

val docTermFreqs = lemmatized.map(terms => {
    val termFreqs = terms.foldLeft(new HashMap[String, Int]()) {
        (map, term) => {
            map += term -> (map.getOrElse(term, 0) + 1)
            map
        }
    }
    termFreqs
})
```

- ▶ For frequent access, the **term-frequency map** should be **cached**.

```
docTermFreqs.cache()
```

Remove Infrequent Terms

- ▶ To further reduce the term space, we may run a **standard word count** to filter out infrequent terms (in this case with a document frequency of less than 12):

```
val docFreqs = docTermFreqs.flatMap(_.keySet).map((_, 1)).  
    reduceByKey(_ + _, 12)
```

- ▶ Keep only the **top 50,000 terms** from within all documents:

```
val ordering = Ordering.by[(String, Int), Int](_.._2)  
val topDocFreqs = docFreqs.top(50000)(ordering)
```

Computing DF Weights (I)

The document frequency of a term across all documents is a global property of the collection. Since we are operating over the partitioned data structure `docTermFreqs`, we need to aggregate it in two steps:

1. We count the **document frequencies locally for each partition**.

```
val zero = new HashMap[String, Int]()

def merge(dfs: HashMap[String, Int],
          tfs: (String, HashMap[String, Int])
        : HashMap[String, Int] = {
  tfs._2.keySet.foreach { term =>
    dfs += term -> (dfs.getOrDefault(term, 0) + 1)
  }
  dfs
}
```

Computing DF Weights (II)

The document frequency of a term across all documents is a global property of the collection. Since we are operating over the partitioned data structure `docTermFreqs`, we need to aggregate it in two steps:

2. We combine the **document frequencies globally across all partitions**.

```
def comb(dfs1: HashMap[String, Int], dfs2: HashMap[String, Int])
  : HashMap[String, Int] = {
  for ((term, count) <- dfs2) {
    dfs1 += term -> (dfs1.getOrDefault(term, 0) + count)
  }
  dfs1
}
```

Computing DF Weights (III)

The document frequency of a term across all documents is a global property of the collection. Since we are operating over the partitioned data structure `docTermFreqs`, we need to aggregate it in two steps:

```
docTermFreqs.aggregate(zero)(merge, comb)
```

And finally invert the DF values into their **IDF counterparts**:

```
val idfs = topDocFreqs.map {  
    case (term, count) =>  
        (term, math.log(numDocs.toDouble / count))  
}.toMap
```

And **broadcast** this map:

```
bIdfs = sc.broadcast(idfs)
```

Broadcast the Term Dictionary

- ▶ Do not forget to **broadcast the static term-id dictionary**:

```
val termIds = idfs.keys.zipWithIndex.toMap  
val bTermIds = sc.broadcast(termIds)
```

Finally Merge the TF and IDF Weights into Sparse Vectors

- ▶ The last transformation **combines the TF and IDF weights** for all terms in a document into a **sparse vector** representation:

```
import scala.collection.JavaConversions._  
import org.apache.spark.mllib.linalg.Vectors  
  
val vecs = docTermFreqs.map(termFreqs => {  
    val docTotalTerms = termFreqs.values().sum  
    val termScores = termFreqs.filter {  
        case (term, freq) => bTermIds.value.containsKey(term)  
    }.map{  
        case (term, freq) => (bTermIds.value(term),  
            bIdfs.value(term) * termFreqs(term) / docTotalTerms)  
    }.toSeq  
    Vectors.sparse(bTermIds.value.size, termScores)  
})
```

Compute the Singular Value Decomposition

- ▶ The `RowMatrix` class is used to represent entries of a **document-term matrix** in a sparse way. It is directly used as input for the SVD computation:

```
import org.apache.spark.mllib.linalg.distributed.RowMatrix  
  
vecs.persist()  
val mat = new RowMatrix(vecs)  
val k = 1000 // number of latent concepts  
              in the reduced matrix  
val svd = mat.computeSVD(k, computeU=true)
```

- ▶ Caution: Spark MLlib *unfortunately* swaps the common convention used to encode the term-document matrix A :
 - ▶ `svd.V` encodes the word-to-concept mappings
 - ▶ `svd.U` encodes the document-to-concept mappings
 - ▶ `svd.S` contains the singular values of A (in descending order)

Finding the Top Terms for Latent Concepts

We can now use V to inspect the **latent concepts** represented in our Wikipedia collection.

- ▶ First, we may want to **rank terms** by their similarity to the top concepts:

```
def topTermsInTopConcepts(  
    svd: SingularValueDecomposition[RowMatrix, Matrix],  
    numConcepts: Int, numTerms: Int): Seq[Seq[(String, Double)]] = {  
    val v = svd.V  
    val topTerms = new ArrayBuffer[Seq[(String, Double)]]()  
    for (i <- 0 until numConcepts) {  
        val offs = i * v.numRows  
        val termWeights = v.toArray.slice(offs, offs + v.numRows).zipWithIndex  
        val sorted = termWeights.sortBy(-_._1)  
        topTerms += sorted.take(numTerms).map {  
            case (score, id) =>  
                (bTermIds.value.find(_.value == id).getOrElse(("", -1)).value, score) }  
    }  
    topTerms }
```

Finding the Top Documents for the Latent Concepts

We can also use \mathbf{U} to inspect the **latent concepts** represented in our Wikipedia collection.

- ▶ Next, we may want to **rank documents** by these top concepts:

```
def topDocsInTopConcepts(  
    svd: SingularValueDecomposition[RowMatrix, Matrix],  
    numConcepts: Int, numDocs: Int): Seq[Seq[(Long, Double)]] = {  
    val u = svd.U  
    val topDocs = new ArrayBuffer[Seq[(Long, Double)]]()  
    for (i <- 0 until numConcepts) {  
        val docWeights = u.rows.map(_.toArray(i)).zipWithUniqueId()  
        topDocs += docWeights.top(numDocs).map {  
            case (score, id) => (id, score) } }  
    topDocs }
```

Print the Results

- ▶ Now we may finally **print the results of all our efforts** so far by using the two afore defined functions:

```
val topConceptTerms = topTermsInTopConcepts(svd, 4, 10)
val topConceptDocs = topDocsInTopConcepts(svd, 4, 10)
for ((terms, docs) <- topConceptTerms.zip(topConceptDocs)) {
    println("Concept terms: " + terms.map(_.toString).mkString(", "))
    println("Concept docs: " + docs.map(_.toString).mkString(", "))
    println()
}
```

- ▶ Note that MLlib stores V **locally at the driver node**, while U is stored in a **distributed manner across all executor nodes** (being sharded on its rows).

Querying the Latent Semantic Index

With the reduced decomposition of A into $U \times S \times V^T$ at hand, we can now perform the following tasks in the "latent" concept space instead of using the actual terms and documents:

- ▶ **Document-to-document similarities:**
 - ▶ Compute row similarities of $U \times S$
- ▶ **Term-to-term similarities:**
 - ▶ Compute row similarities of $V \times S$
- ▶ **Single-term-to-document similarities:**
 - ▶ Multiply $U \times S$ with the row vector of term i in V
- ▶ **Multi-term-to-document similarities:**
 - ▶ Transform $q' = q_{1 \times m} \times V_{m \times k}$
 - ▶ Multiply $U \times S$ with the transformed query vector q'^T

Document-Document Relevance

- ▶ Based on the reduced representation, we can find the **most similar documents** to a given query document:

```
import org.apache.spark.mllib.linalg.Matrices

def topDocsForDoc(normalizedUS: RowMatrix, docId: Long)
  : Seq[(Double, Long)] = {
  val docRowArr = row(normalizedUS, docId)
  val docRowVec = Matrices.dense(docRowArr.length, 1, docRowArr)
  val docScores = normalizedUS.multiply(docRowVec)
  val allDocWeights = docScores.rows.map(_.toArray(0)).
    zipWithUniqueId()
  allDocWeights.filter(!_._1.isNaN).top(10)
}

val US = multiplyByDiagonalMatrix(svd.U, svd.s)
val normalizedUS = rowsNormalized(US)
topDocsForDoc(normalizedUS, idDocs(doc), docIds)
```

Term-Term Relevance

- ▶ Similarly, we can also find the **top similar terms** to a given query term:

```
import breeze.linalg.{SparseVector => BSparseVector}
import breeze.linalg.{DenseVector => BDenseVector}
import breeze.linalg.{DenseMatrix => BDenseMatrix}

def topTermsForTerm(
    normalizedVS: BDenseMatrix[Double],
    termId: Int): Seq[(Double, Int)] = {
  val rowVec = new BDenseVector[Double](
    row(normalizedVS, termId).toArray)
  val termScores = (normalizedVS * rowVec).toArray.zipWithIndex
  termScores.sortBy(-_._1).take(10)
}

val VS = multiplyByDiagonalMatrix(svd.V, svd.s)
val normalizedVS = rowsNormalized(VS)
topTermsForTerm(normalizedVS, idTerms(term), termIds)
```

Term-Document Relevance

- As in an actual **Information Retrieval** setting, we can find the **most similar documents** for a given query term:

```
def topDocsForTerm(US: RowMatrix, V: Matrix, termId: Int)
  : Seq[(Double, Long)] = {
  val termRowArr = row(V, termId).toArray
  val termRowVec = Matrices.dense(termRowArr.length, 1, termRowArr)
  val docScores = US.multiply(termRowVec)
  val allDocWeights = docScores.rows.map(_.toArray(0)).
    zipWithUniqueId()
  allDocWeights.top(10)
}

topDocsForTerm(normalizedUS, svd.V, idTerms(term))
```

Multi-Term Queries (I)

- ▶ First, we transform a query vector into a compatible `SparseVector` representation using IDF values as term weights:

```
def termsToQueryVector(  
    terms: Seq[String],  
    idTerms: Map[String, Int],  
    idfs: Map[String, Double]): BSparseVector[Double] = {  
    val indices = terms.map(idTerms(_)).toArray  
    val values = terms.map(idfs(_)).toArray  
    new BSparseVector[Double](indices, values, idTerms.size)  
}
```

Multi-Term Queries (II)

- ▶ Finally, we multiply $\mathbf{U} \times \mathbf{S}$ with the transformed query vector \mathbf{q}'^T :

```
def topDocsForTermQuery(  
    US: RowMatrix,  
    V: Matrix,  
    query: BSparseVector[Double]): Seq[(Double, Long)] = {  
    val breezeV = new BDenseMatrix[Double](V numRows, V numCols, V toArray)  
    val termRowArr = (breezeV.t * query).toArray  
    val termRowVec = Matrices.dense(termRowArr.length, 1, termRowArr)  
    val docScores = US.multiply(termRowVec)  
    val allDocWeights = docScores.rows.map(_.toArray(0)).  
        zipWithUniqueId()  
    allDocWeights.top(10) }  
  
val queryVec = termsToQueryVector(terms, idTerms, idfs)  
topDocsForTermQuery(US, svd.V, queryVec)
```

Summary

- ▶ **LSA** and related techniques, such as PCA, pLSA, LDA, are very broadly applied techniques for **analyzing text data**.
 - ▶ Dimensionality reduction
 - ▶ Content filtering
 - ▶ Clustering
 - ▶ Visualization
 - ▶ Indexing & retrieval
- ▶ **LSA** also has a variety of applications also **outside text analysis**:
 - ▶ Face detection ("eigenfaces") to detect common patterns in human appearance
 - ▶ Detection of climate patterns, etc.