



Chapter 2

Downloading/Setting up Apache Spark



What is Homebrew?

- Homebrew is a popular package manager for macOS.
- It allows users to easily install, update, and manage software packages.

Benefits from Homebrew:

- Streamlines software installation: No need to download and compile software manually.
- Offers a vast library of software packages for developers and users.
- Maintained by a community of contributors, ensuring up-to-date packages.

Basic commands: `brew install`, `brew update`, `brew upgrade`, and more.



Step 1 - Downloading Apache Spark for Python

1. Install Homebrew

Simplify the installation and management of various software packages on your system.

Run the following command in your terminal:

```
/bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
```

2. Install Python with Homebrew

Ensure you have Python installed

Run the following command: brew install python

3. Install Apache Spark for Python

Enable integration of Apache Spark with Python

Run the following command: pip install pyspark

Update command: **pip install --upgrade pyspark**



Step 2 - Test

1. Run the following command on the terminal:

```
1 pyspark
2 # Used to start an interactive PySpark shell, which is an interactive
3 # Python shell with PySpark functionality enabled.
```

You should expect this:

```
1 ~ pyspark
2 Python 3.10.7 (v3.10.7:6cc6b13308, Sep  5 2022, 14:02:52) [Clang 13.0.0 (clang-1300.0.29.30)] on darwin
3 Type "help", "copyright", "credits" or "license" for more information.
4 Setting default log level to "WARN".
5 To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
6 23/10/01 16:54:08 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
7 Welcome to
8   __  __
9  / _ \/_\  _ \_ \_ / /_ \
10 / \_ \_ \_ \_ \_ \_ \_ \_ \_ \_ \
11 /_ \_ \_ \_ \_ \_ \_ \_ \_ \_ \_ \_ \
12 /_ \_ \_ \_ \_ \_ \_ \_ \_ \_ \_ \_ \
13
14 Using Python version 3.10.7 (v3.10.7:6cc6b13308, Sep  5 2022 14:02:52)
15 Spark context Web UI available at http://geralds-mbp.fritz.box:4040
16 Spark context available as 'sc' (master = local[*], app id = local-1696172049014).
17 SparkSession available as 'spark'.
```

To exit the shell, type: **quit()**

An example of using the Python Interpretive shell:

```
~ pyspark
Python 3.10.7 (v3.10.7:6cc6b13308, Sep  5 2022, 14:02:52) [Clang 13.0.0 (clang-1300.0.29.30)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/10/01 16:54:08 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Welcome to

    __
   /_/\_ \_ \_ \_ \_ \_ \_ \
  _\ \_ \_ \_ \_ \_ \_ \_ \_ \
 /_/\_ \_ \_ \_ \_ \_ \_ \_ \
version 3.5.0
/_\

Using Python version 3.10.7 (v3.10.7:6cc6b13308, Sep  5 2022 14:02:52)
Spark context Web UI available at http://geralds-mbp.fritz.box:4040
Spark context available as 'sc' (master = local[*], app id = local-1696172049014).
SparkSession available as 'spark'.
>>> strings = spark.read.text("README.md")
# The command reads the content of the "README.md" file into a PySpark DataFrame named "strings."
>>> strings.show(10, truncate=False)
# This command displays the first 10 rows of the "strings" DataFrame without truncating the text in each cell.

+-----+
|value|
+-----+
| # Apache Spark
| |
| |Spark is a unified analytics engine for large-scale data processing. It provides
| |high-level APIs in Scala, Java, Python, and R, and an optimized engine that
| |supports general computation graphs for data analysis. It also supports a
| |rich set of higher-level tools including Spark SQL for SQL and DataFrames,
| |pandas API on Spark for pandas workloads, MLLib for machine learning, GraphX for graph processing,
| |and Structured Streaming for stream processing.
| |
| |<https://spark.apache.org/>
+-----+
only showing top 10 rows

>>> strings.count()
125
```

Important Concepts

Application

It consists of a driver program and executors on the cluster.

Job

Parallel computation triggered by Spark actions like save() and collect().

Task

A single unit of work or execution that will be sent to a Spark executor.

SparkSession

Entry point for Spark API interaction in shell/applications.

Stage

Jobs split into dependent stages for parallel execution.

Transformations, Actions, and Lazy Evaluation

Spark operations on distributed data fall into two categories:

Transformations and Actions.

Transformations create a Spark DataFrame starting from another one, without modifying the original data.

Example: Operations like `select()` or `filter()` return a new DataFrame.

**Transformations are evaluated *lazily*

Meaning that results are not computed immediately, but they are recorded or remembered as a lineage.

Actions RDD actions are operations that return the raw values

Example: Spark actions include `count()` and `collect()`

Transformations	Actions
<code>orderBy()</code>	<code>show()</code>
<code>groupBy()</code>	<code>take()</code>
<code>filter()</code>	<code>count()</code>
<code>select()</code>	<code>collect()</code>
<code>join()</code>	<code>save()</code>

Example of Transformation

```
1 | >>> strings = spark.read.text("README.md")
2 | >>> filtered = strings.filter(strings.value.contains("Spark"))
3 | >>> filtered.count()
4 | 20
```

- It includes two transformations: `read()` (to read data) and `filter()` (to filter data).
- The key point is that nothing happens (in the terminal) when transformations are applied: execution is deferred until an action, such as `count()`, is called.
- In this example, the actual processing and filtering of data occur only when `filtered.count()` is executed in the Spark shell.

Narrow and Wide Transformations

Transformations can be classified into two groups:

Narrow and Wide Transformations

Narrow transformations are those where each output can be computed from a single input.

They are *efficient* and *fast*.


`filter()`, `map()`, `union()`, and `contains()`.

Wide transformations are those that require data from multiple input partitions to compute the output.


They are slower and less efficient as they require data shuffling to be performed.

`groupBy()`, `reduceByKey()`, `join()`, and `orderBy()`

Narrow Dependencies



Wide Dependencies




The Spark UI

Apache Spark provides a graphical user interface for monitoring applications at various stages. It can be found on the first message when you launch SPARK (<http://mf0023.homenet.telecomitalia.it:4040>).

The monitoring includes jobs, stages, and tasks.

Key Information Available

- Scheduler stages and tasks
- Memory Usage
- Environment Details
- Executor Information
- Spark SQL Queries



Code for Counting M&Ms

```
1 import sys
2 from pyspark.sql import SparkSession
3 from pyspark.sql.functions import count
4
5 # Check if this script is the main program
6 if __name__ == "__main__":
7     # Ensure that the correct number of command-line arguments are provided
8     if len(sys.argv) != 2:
9         print("Usage: python mnmcount.py <file>", file=sys.stderr)
10        sys.exit(-1)
11
12 # Create a SparkSession, which is the entry point to Spark functionality.
13 # If a SparkSession doesn't exist, it will create a new one.
14 spark = (SparkSession
15         .builder
16         .appName("PythonMnMCount") # Assign a name to the Spark application
17         .getOrCreate())
18
19 # Retrieve the M&M data set filename from the command-line arguments
20 mnm_file = sys.argv[1]
21
22 # Read the specified file into a Spark DataFrame. We use the CSV format
23 # and specify two options: "header" to indicate that the first row contains
24 # column names and "inferSchema" to automatically detect data types.
25 mnm_df = (spark.read.format("csv")
26           .option("header", "true")
27           .option("inferSchema", "true")
28           .load(mnm_file))
```

```
1 # We use the DataFrame high-level APIs to work with structured data. Note
2 # that we are exclusively using DataFrames and not RDDs (Resilient Distributed Datasets).
3 # In Spark, some functions return the same object, allowing us to chain function calls.
4
5 # Step 1: We select the fields "State," "Color," and "Count" from the DataFrame.
6 # Step 2: To group each state and its M&M color count, we use the groupBy() function.
7 # Step 3: We aggregate the counts of all colors and groupBy("State" and "Color.")
8 # Step 4: We order the results in descending order based on the "Total" count.
9 count_mmm_df = (mmm_df
10                 .select("State", "Color", "Count")
11                 .groupBy("State", "Color")
12                 .agg(count("Count").alias("Total")))
13                 .orderBy("Total", ascending=False))
14
15 # Show the resulting aggregations for all the states and colors;
16 # this displays the total count of each color per state.
17 # Note that "show()" is an action, which triggers the above query to be executed.
18 count_mmm_df.show(n=60, truncate=False)
19
20 # Print the total number of rows in the resulting DataFrame.
21 print("Total Rows = %d" % (count_mmm_df.count()))
```

```
1 # Here we focus on a single state.
2
3 # Step 1: Select all rows from the DataFrame.
4 # Step 2: Filter only those rows where the state is "CA."
5 # Step 3: Group the data by "State" and "Color" just like before.
6 # Step 4: Aggregate the counts for each color.
7 # Step 5: Order the results in descending order.
8
9 ca_count_mmm_df = (mmm_df
10                     .select("State", "Color", "Count")
11                     .where(mmm_df.State == "CA")
12                     .groupBy("State", "Color")
13                     .agg(count("Count").alias("Total")))
14                     .orderBy("Total", ascending=False))
15
16 # Show the resulting aggregation for the state of California (CA).
17 # Similar to the previous example, the "show()" function is an action
18 # that will execute the computation and display the results.
19 ca_count_mmm_df.show(n=10, truncate=False)
20
21 # Stop the SparkSession.
22 spark.stop()
```

Counting M&Ms for the Cookie Monster

1	+	-	-	-	-	+
2		State	Color	Count		
3	+	-	-	-	-	+
4	TX	Red	20			
5	NV	Blue	66			
6	CO	Blue	79			
7	OR	Blue	71			
8	WA	Yellow	93			
9	+	-	-	-	-	+
10	only showing top 5 rows					

Fig. 1

1	+	-	-	-	-	+
2		State	Color	sum(Count)		
3	+	-	-	-	-	+
4	CA	Yellow	100956			
5	CA	Brown	95762			
6	CA	Green	93505			
7	CA	Red	91527			
8	CA	Orange	90311			
9	CA	Blue	89123			
10	+	-	-	-	-	+

Fig. 3

1	+	-	-	-	-	+
2		State	Color	sum(Count)		
3	+	-	-	-	-	+
4	CA	Yellow	100956			
5	IA	Green	96486			
6	CA	Brown	95762			
7	TX	Green	95753			
8	TX	Red	95404			
9	CO	Yellow	95638			
10	NM	Red	94699			
11	UT	Green	94324			
12	NY	Green	94339			
13	NV	Orange	93929			
14	TX	Yellow	93819			
15	CO	Green	93724			
16	CO	Brown	93692			
17	CO	Green	93696			
18	NM	Brown	93447			
19	CO	Blue	93412			
20	WA	Red	93332			
21	WA	Brown	93882			
22	WA	Yellow	92928			
23	NM	Yellow	92747			
24	WA	Brown	92768			
25	TX	Orange	92215			
26	AZ	Brown	92287			
27	AZ	Green	91882			
28	NY	Red	91768			
29	AZ	Orange	91684			
30	UT	Green	91527			
31	WA	Orange	91521			
32	NV	Yellow	91390			
33	UT	Orange	91341			
34	NV	Green	91331			
35	NM	Orange	91251			
36	FL	Green	91168			
37	NY	Blue	90902			
38	UT	Red	90995			
39	CO	Orange	90971			
40	AZ	Yellow	90946			
41	TX	Brown	90736			
42	OR	Blue	90526			
43	OR	Orange	90411			
44	OR	Red	90286			
45	NM	Blue	90159			
46	AZ	Red	90042			
47	NV	Blue	89903			
48	UT	Blue	89977			
49	FL	Green	89975			
50	NA	Blue	89980			
51	OR	Green	89578			
52	CO	Red	89465			
53	NV	Red	89346			
54	UT	Yellow	89264			
55	OR	Brown	89136			
56	CA	Blue	89123			
57	UT	Brown	88973			
58	TX	Blue	88466			
59	UT	Green	88392			
60	OR	Yellow	88129			
61	NY	Orange	87956			
62	NY	Yellow	87800			
63	NY	Brown	86110			
64	+	-	-	-	-	+

Fig. 2