

# Generative modeling

Prototyping with Deep Learning

# Learning outcomes

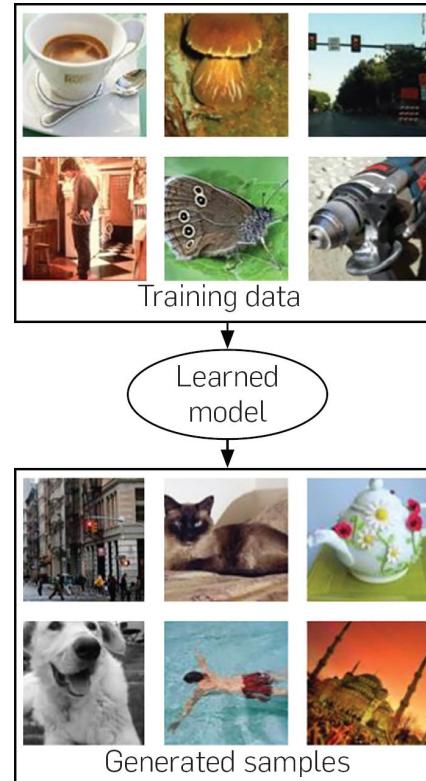
After this lesson you will be able to:

- Identify the need for generative models
- Recognize popular model architectures
- Understand the challenges of generative models

# Goal

Learn to generate *new data* from examples

How awesome does it sound?



# Recap: Learning paradigms



## Supervised learning

$$f(x) = y$$

Challenge: get high-quality labels

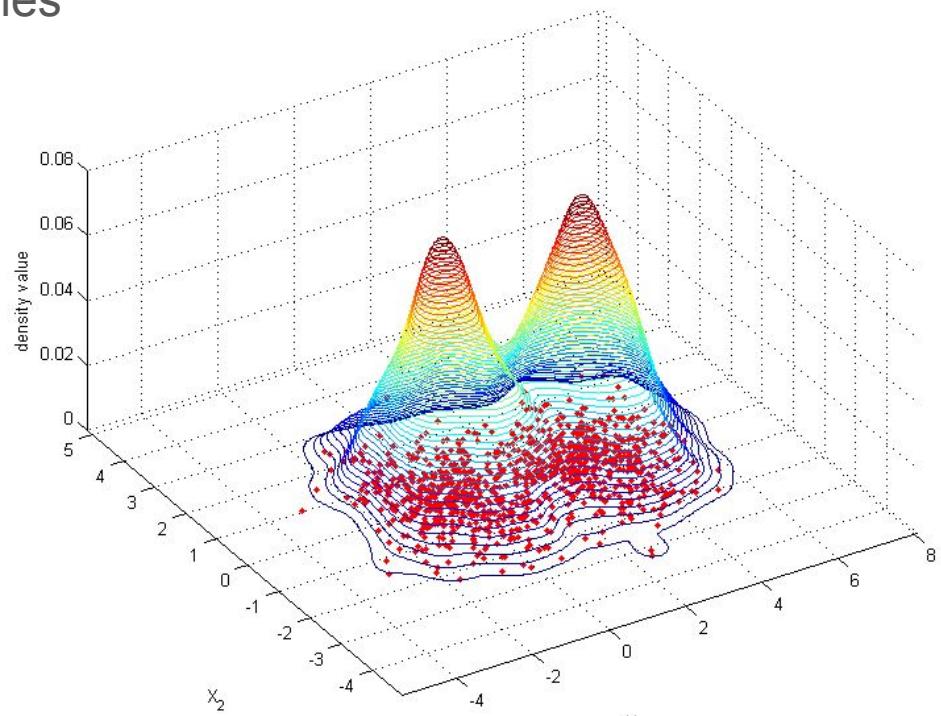
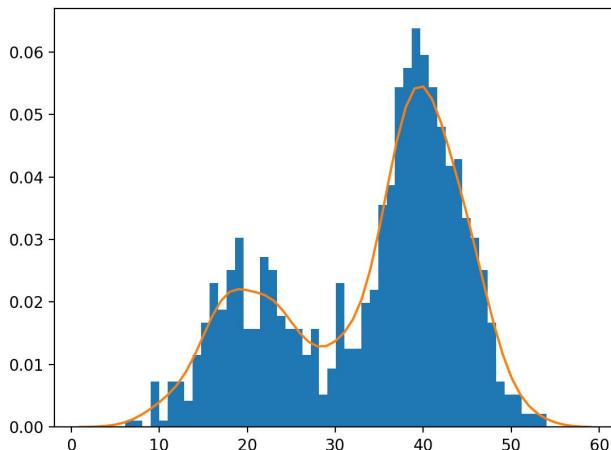
## Unsupervised learning

$$f(x) = x'$$

Challenge: understand the structure of the data

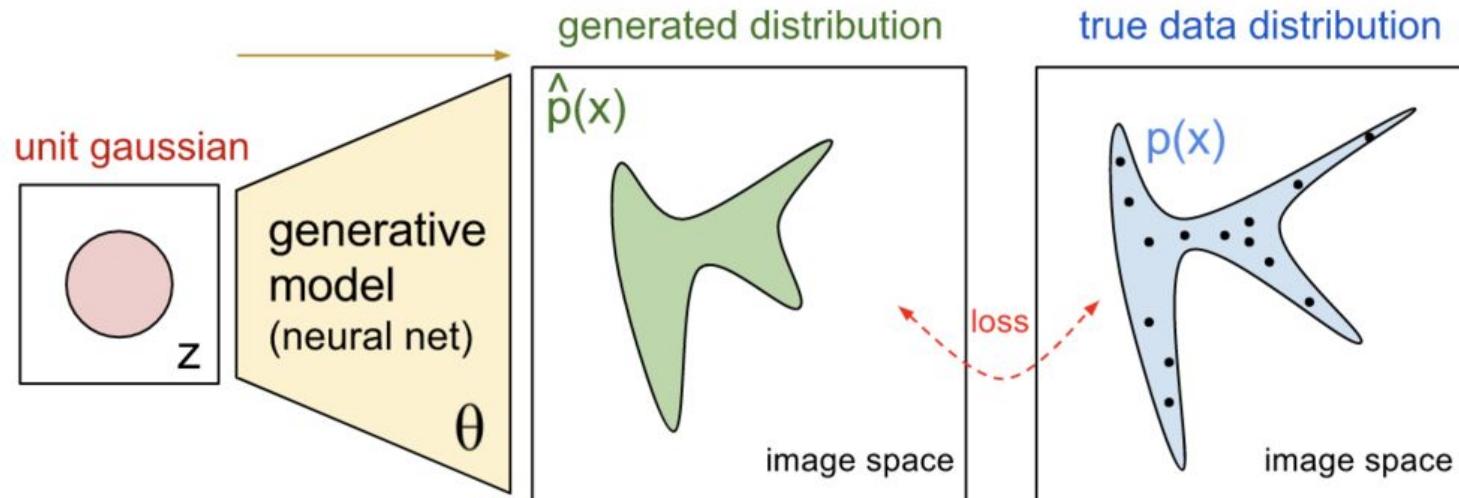
# Probability density estimation

## Parametric vs non-parametric approaches

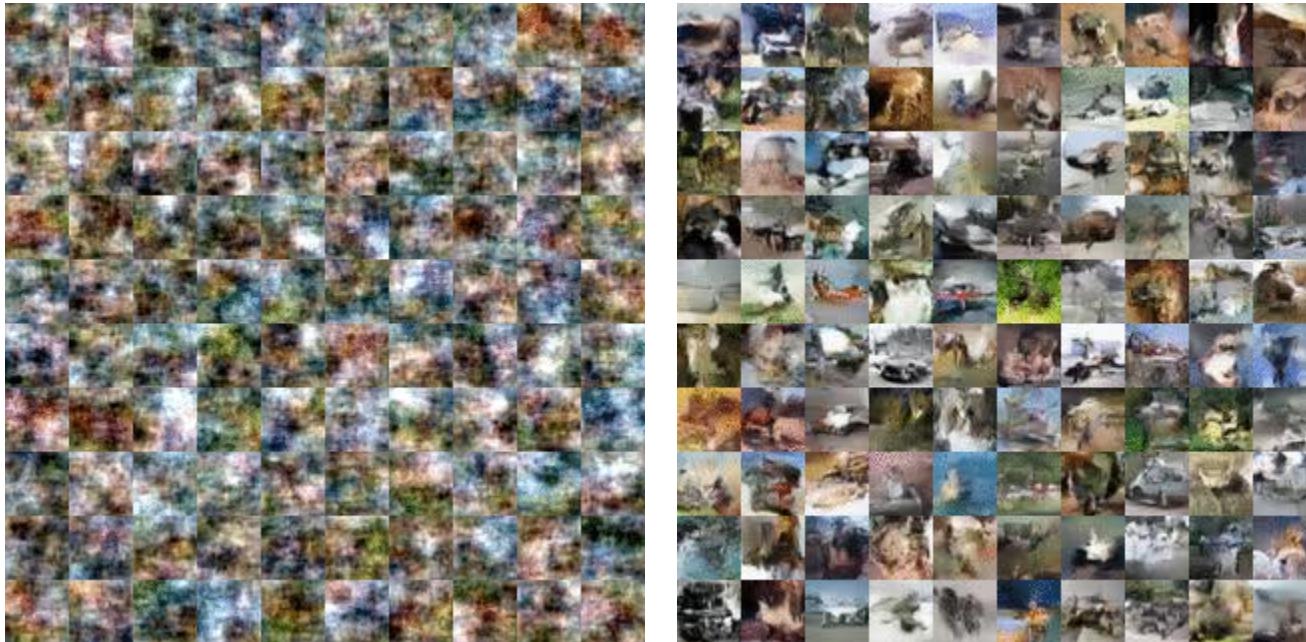


# Framework

Learn  $p_{\text{model}}(\mathbf{x}) \approx p_{\text{data}}(\mathbf{x})$ , however  $p_{\text{data}}(\mathbf{x})$  is unknown!



# Examples



<https://openai.com/blog/generative-models/>

# How to estimate density?

## Explicitly

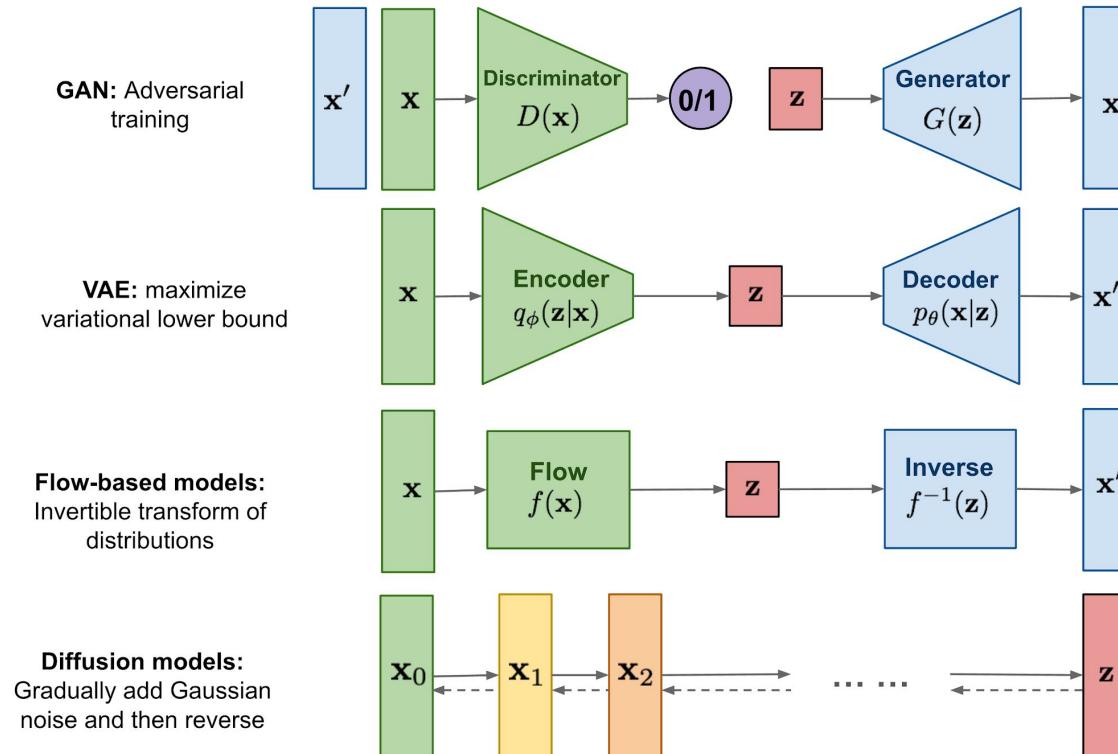
- Tractable density: Autoregressive, flow-based, and diffusion models
- Intractable density: VAEs

## Implicitly

- GANs

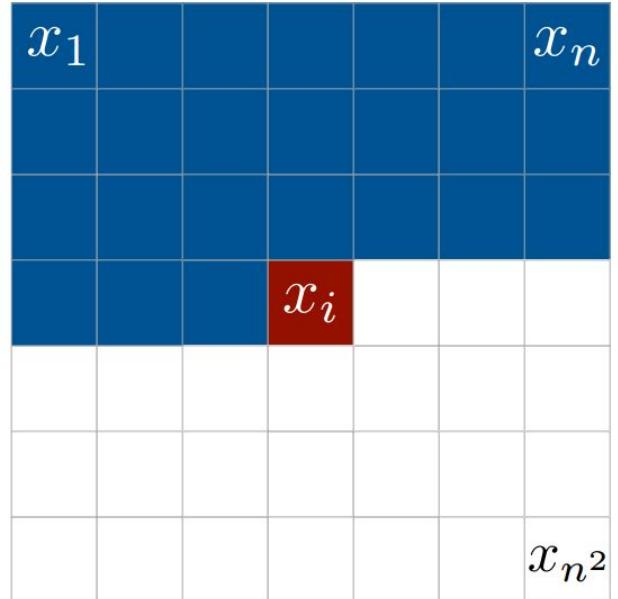
More approaches also available but they are not that much popular; see e.g. Fig.9 in  
<https://arxiv.org/pdf/1701.00160.pdf>

# Summary of approaches



# Autoregressive models

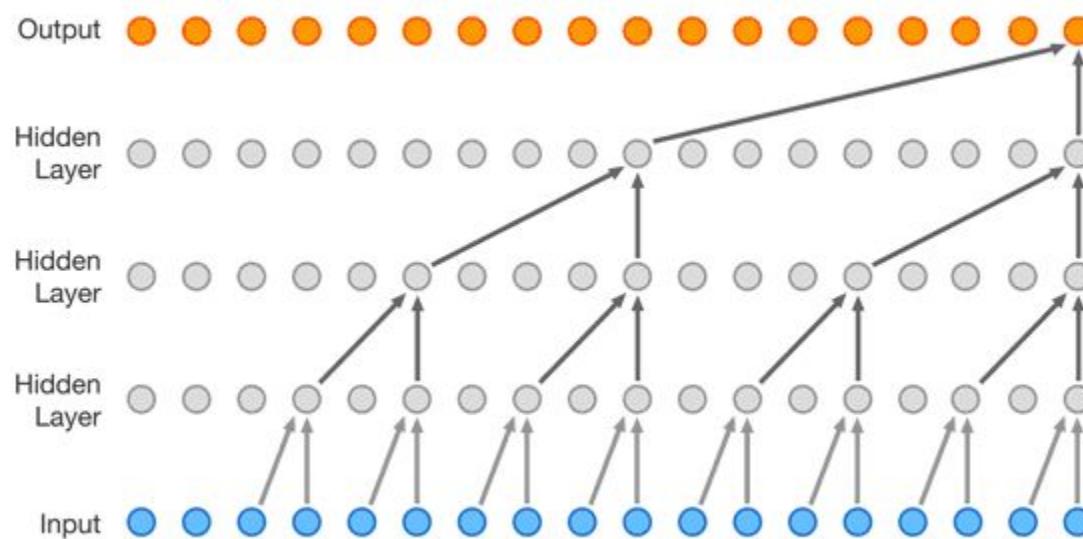
$$p_{\theta}(x) = \prod_{i=1}^n p_{\theta}(x_i | x_1, \dots, x_{i-1})$$



Examples: [PixelRNN](#) and [PixelCNN](#)

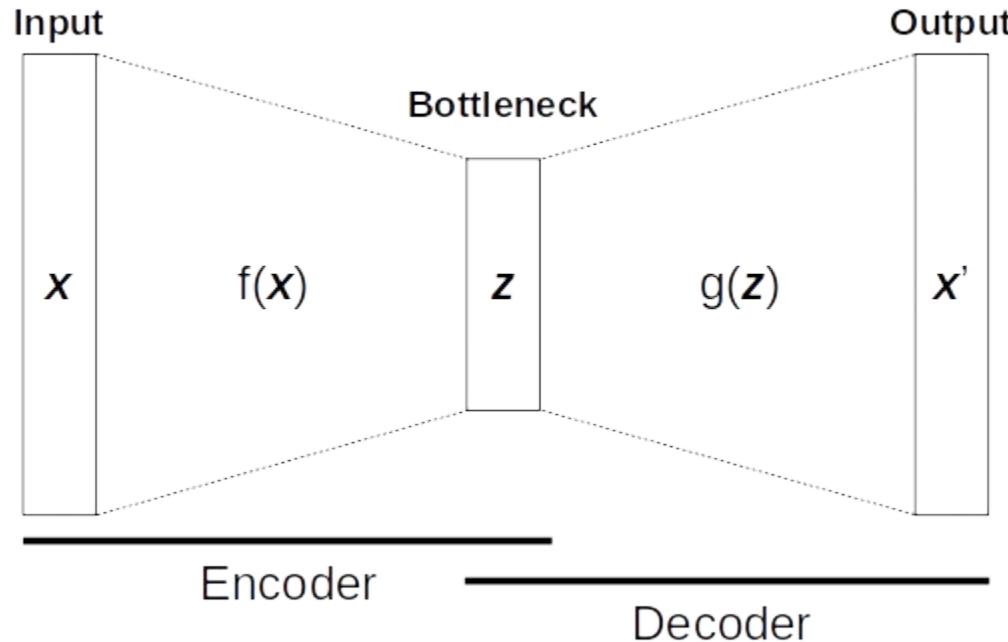
<https://towardsdatascience.com/32d19291173>

# Autoregressive models



<https://deepmind.com/blog/article/wavenet-generative-model-raw-audio>

# Recap: Autoencoders



# Variational inference

Intractable likelihood:  $p_\theta(x) = \int p_\theta(z)p_\theta(x|z)dz$

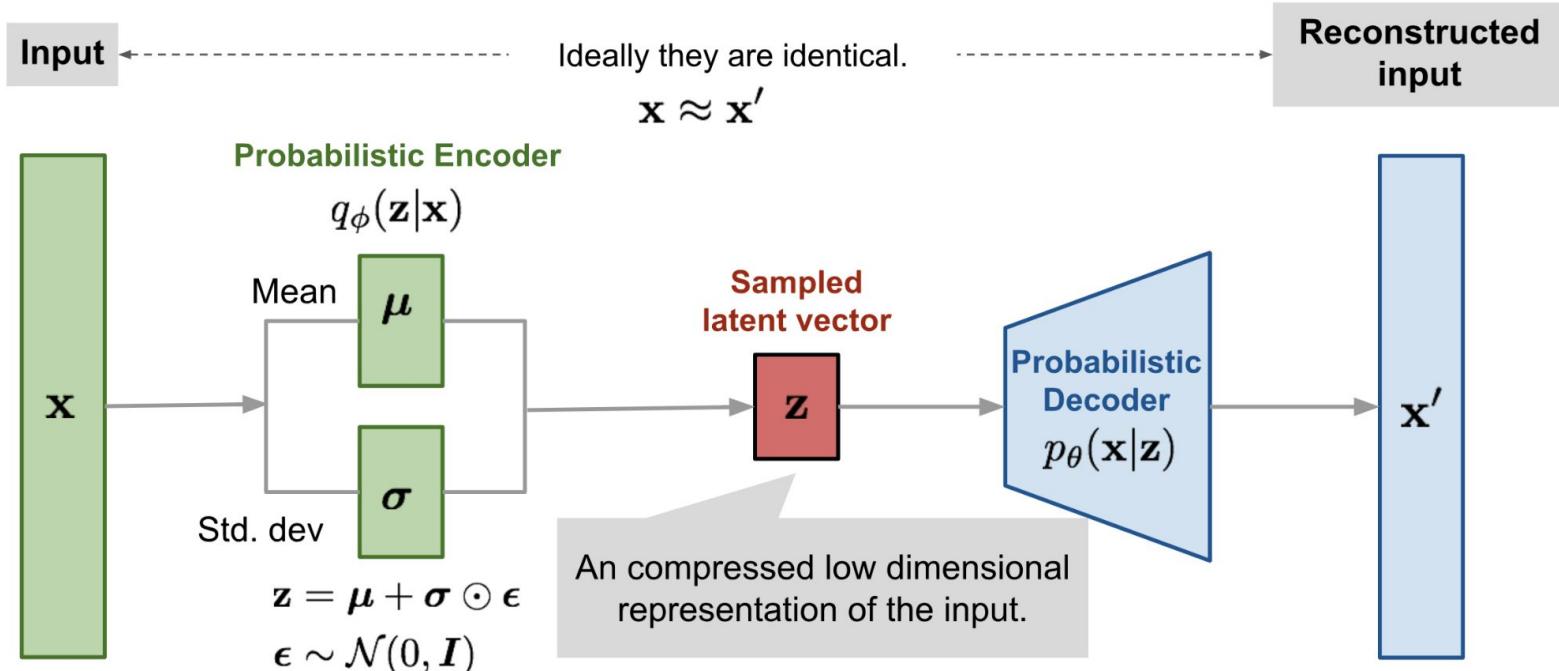
Posterior density also intractable:  $p_\theta(z|x) = \frac{p_\theta(x|z)p_\theta(z)}{p_\theta(x)}$

Solutions: (1) *probabilistic* encoder and decoder:  $q_\phi(z|x) \approx p_\theta(x|z)$   
(2) ELBO criterion:  $\log p(x) \geq \mathcal{L}(x; \theta, \phi)$

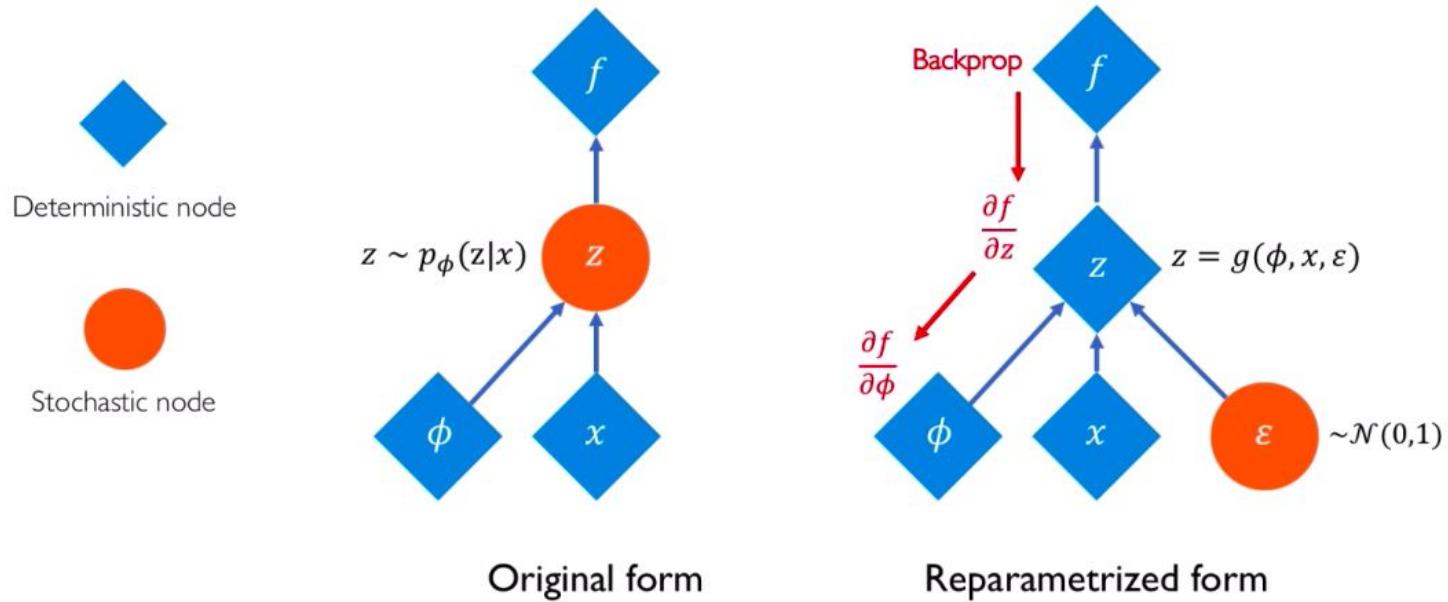
See also

<https://fdocuments.in/document/notes-on-variational-autoencoders-dmmmlvaepdf-notes-on-variational-autoencoders.html?page=1>

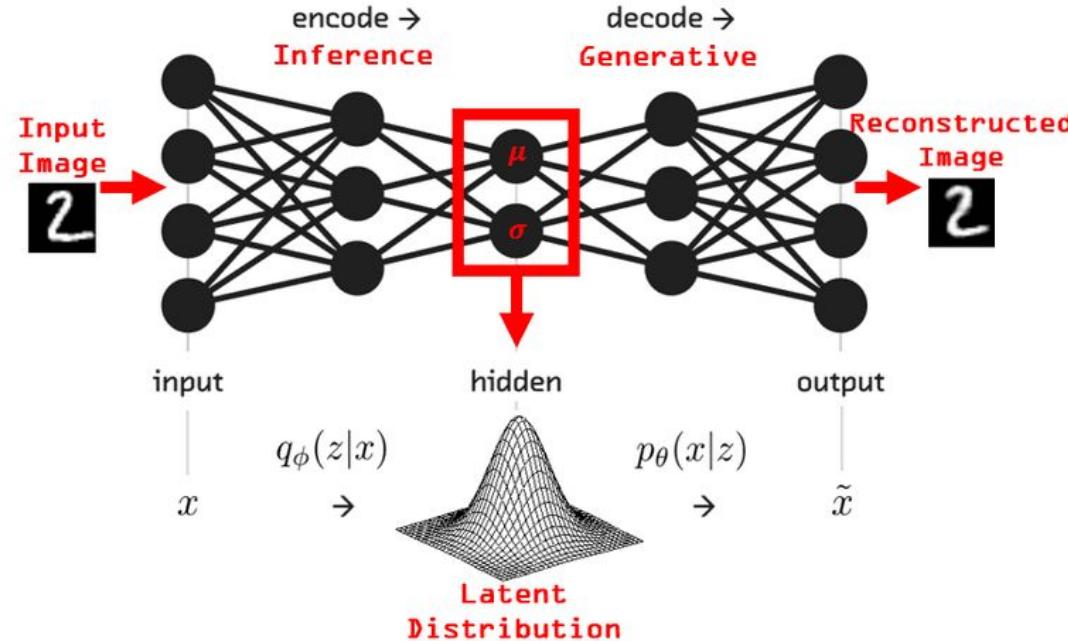
# Variational Autoencoder (VAE)



# Variational Autoencoder (VAE)



# Variational Autoencoder (VAE)



$$\begin{aligned}\mathcal{L}(\mathbf{x}; \theta) &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}, \mathbf{z}) - \log q(\mathbf{z}|\mathbf{x})] \\ &= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}|\mathbf{z})] - D_{KL}(q(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}))\end{aligned}$$

# Variational Autoencoder (VAE)

Reconstruction:



Generation:



# Variational Autoencoder (VAE)

Conditional reconstruction:



Conditional generation:



# Generative Adversarial Network (GAN)



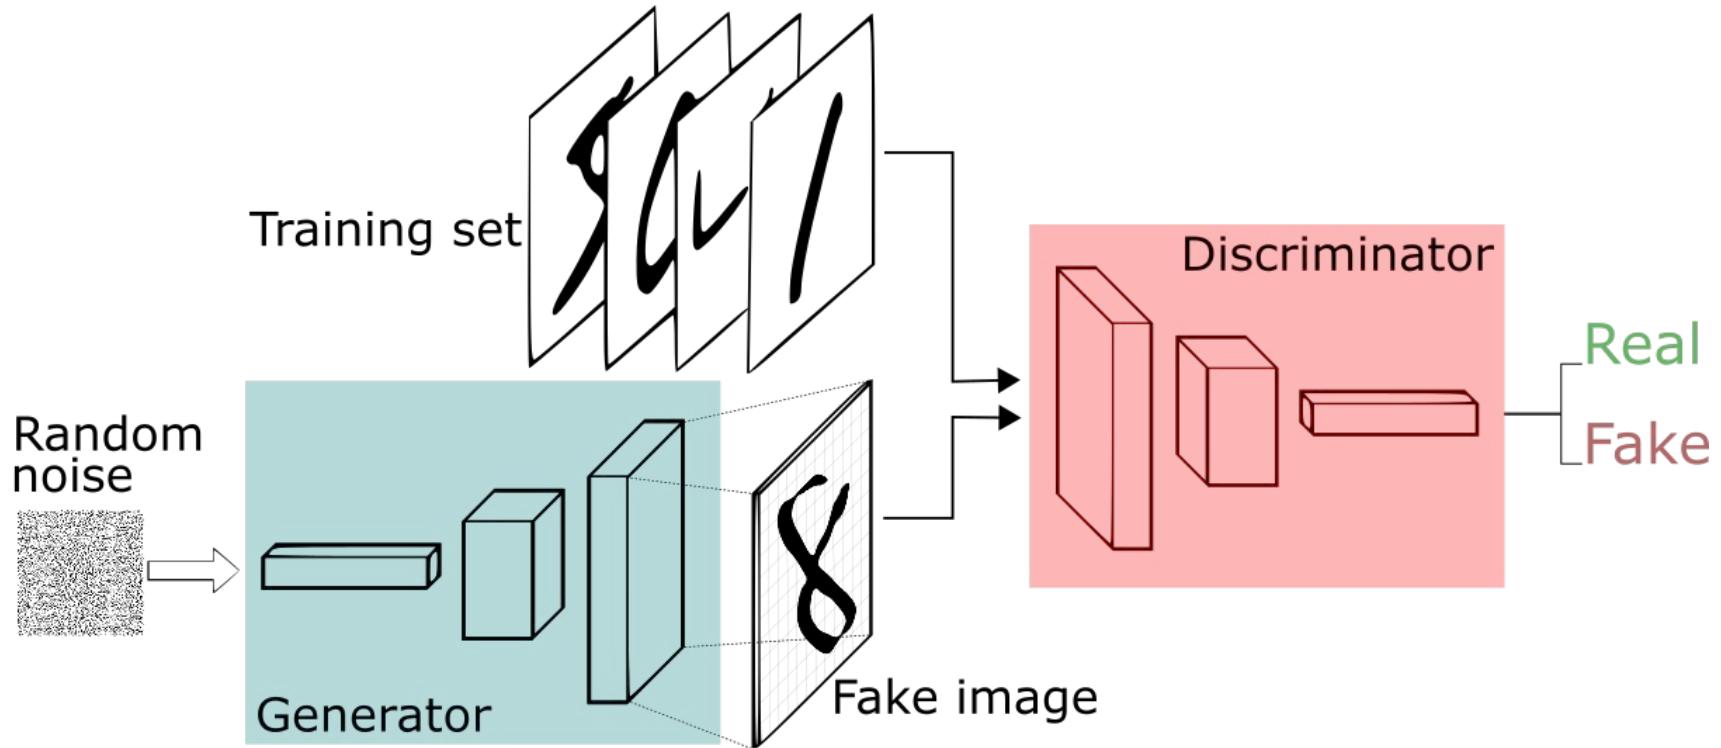
Forget about *explicit* density estimation

We care most about sampling!

Game-theoretic approach (Nash equilibrium)

- **Generator:** generate real-looking data
- **Discriminator:** distinguish between real and fake data

# Generative Adversarial Network (GAN)



# Generative Adversarial Network (GAN)

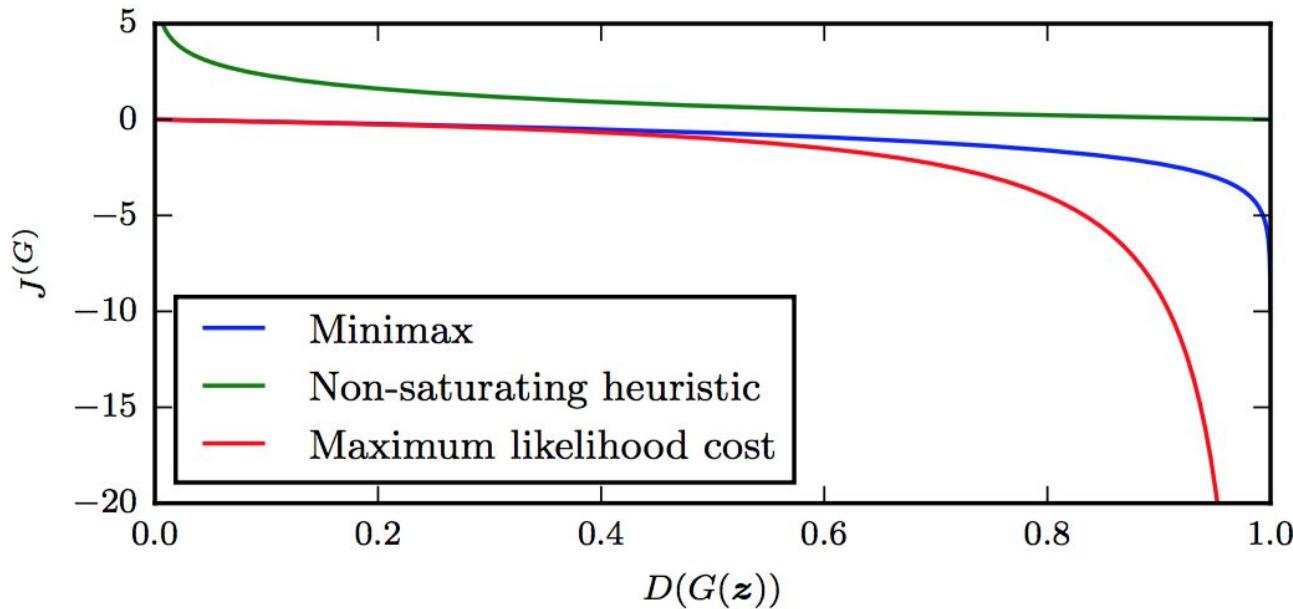
Minimax objective function:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p(x)} \log D(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D(G(z)))$$

Non-saturating version:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p(x)} \log D(x) - \mathbb{E}_{z \sim p(z)} \log D(G(z))$$

# Generative Adversarial Network (GAN)



<https://danieltakeshi.github.io/2017/03/05/understanding-generative-adversarial-networks/>

# Generative Adversarial Network (GAN)

Adversarial training:

1. Gradient ascent on discriminator:  $D(x) \rightarrow 1$  and  $D(G(z)) \rightarrow 0$
2. Gradient descent on generator:  $D(G(z)) \rightarrow 1$   
→ Better: gradient ascend on non-saturating generator loss

# Training GANs is (really) hard



<https://www.slideshare.net/EmanueleGhelfi/gan-theory-and-applications-143737572>

# Training GANs is (really) hard



Highly sensitive to hyperparameter selection

**Non-convergence:** model parameters oscillate (a lot)

**Mode collapse:** generator produces few varieties samples

**Diminished gradient:** discriminator quickly succeeds so generator cannot learn

# On Nash equilibrium

Min-max game:  $\min_x \max_y f(\textcolor{red}{x}, \textcolor{blue}{y}) \stackrel{\text{convex}}{\uparrow} \stackrel{\text{concave}}{\uparrow} \geq \max_y \min_x f(\textcolor{red}{x}, \textcolor{blue}{y})$

Saddle point  $(\textcolor{red}{x}^*, \textcolor{blue}{y}^*)$ :  $f(\textcolor{red}{x}^*, \textcolor{blue}{y}) \leq f(\textcolor{red}{x}^*, \textcolor{blue}{y}^*) \leq f(\textcolor{red}{x}, \textcolor{blue}{y}^*)$

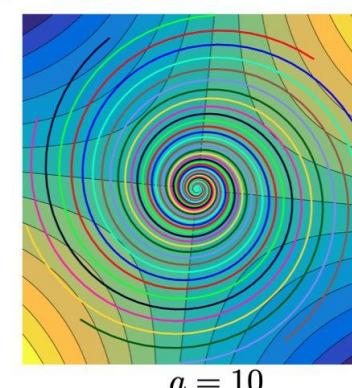
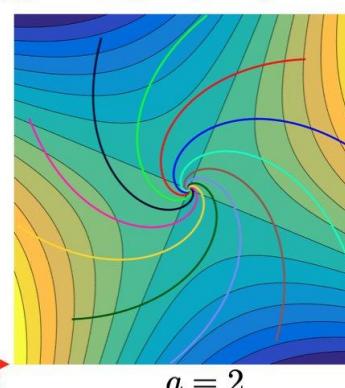
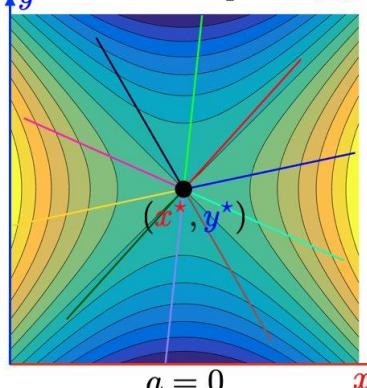
→ Strong duality:  $\min_x \max_y f(\textcolor{red}{x}, \textcolor{blue}{y}) = \max_y \min_x f(\textcolor{red}{x}, \textcolor{blue}{y})$

Gradient descent: 
$$\begin{cases} \textcolor{red}{x}_{k+1} = \textcolor{red}{x}_k - \tau \nabla_{\textcolor{red}{x}} f(\textcolor{red}{x}_k, \textcolor{blue}{y}_k) \\ \textcolor{blue}{y}_{k+1} = \textcolor{blue}{y}_k + \tau \nabla_{\textcolor{blue}{y}} f(\textcolor{red}{x}_k, \textcolor{blue}{y}_k) \end{cases}$$



John Forbes Nash

Example:  $f(\textcolor{red}{x}, \textcolor{blue}{y}) = \textcolor{green}{a} \textcolor{red}{x} \textcolor{blue}{y} + \textcolor{red}{x}^2 - \textcolor{blue}{y}^2$        $\textcolor{green}{a}$  = interaction



# Latent space interpolation in VAEs



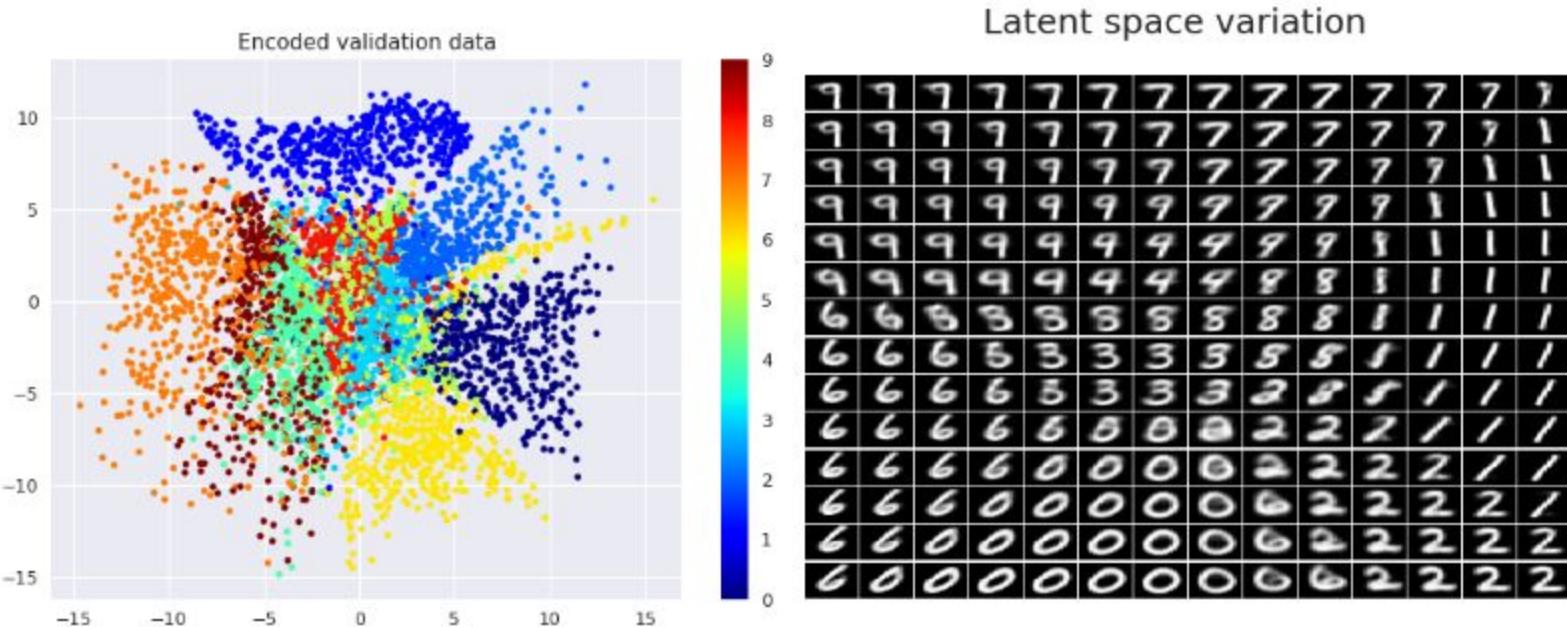
what can happen without regularisation



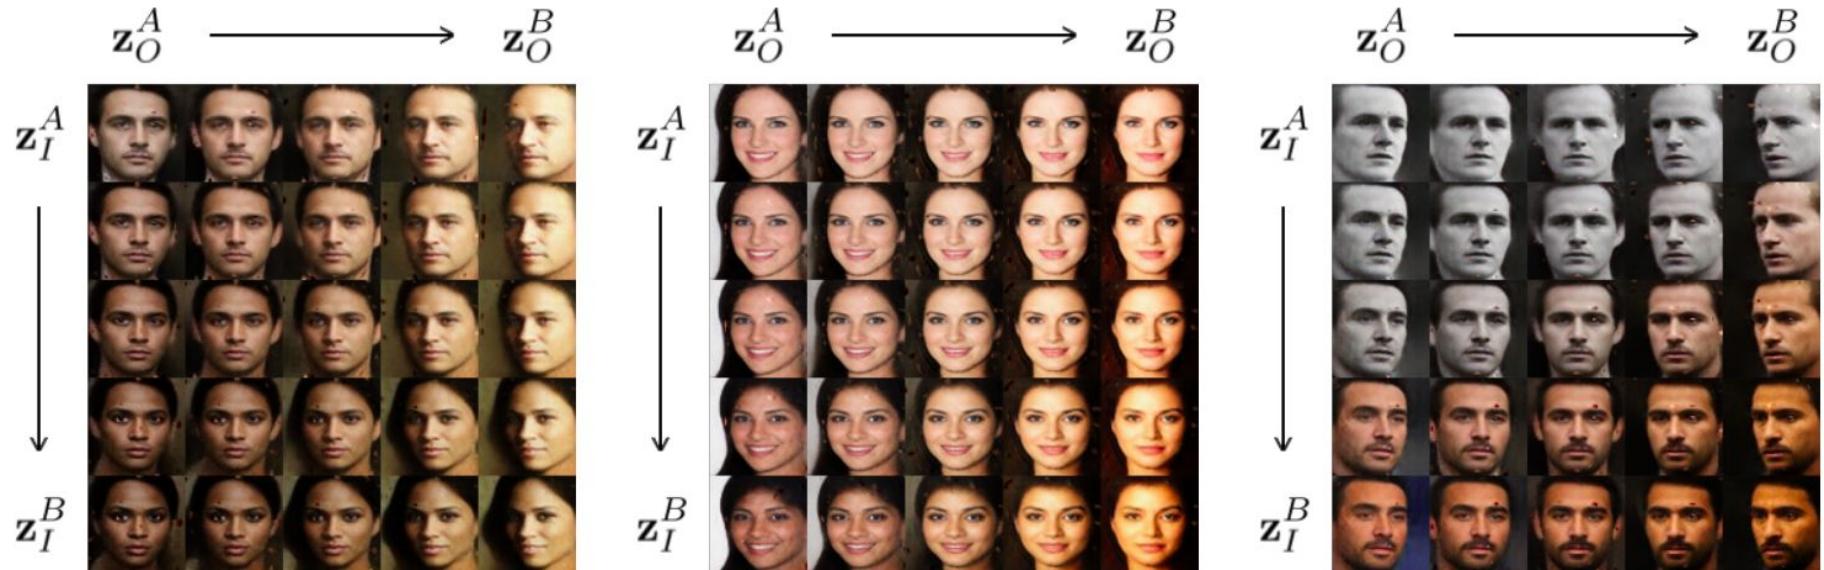
what we want to obtain with regularisation



# Latent space interpolation in GANs



# Semantic decomposition of latent spaces



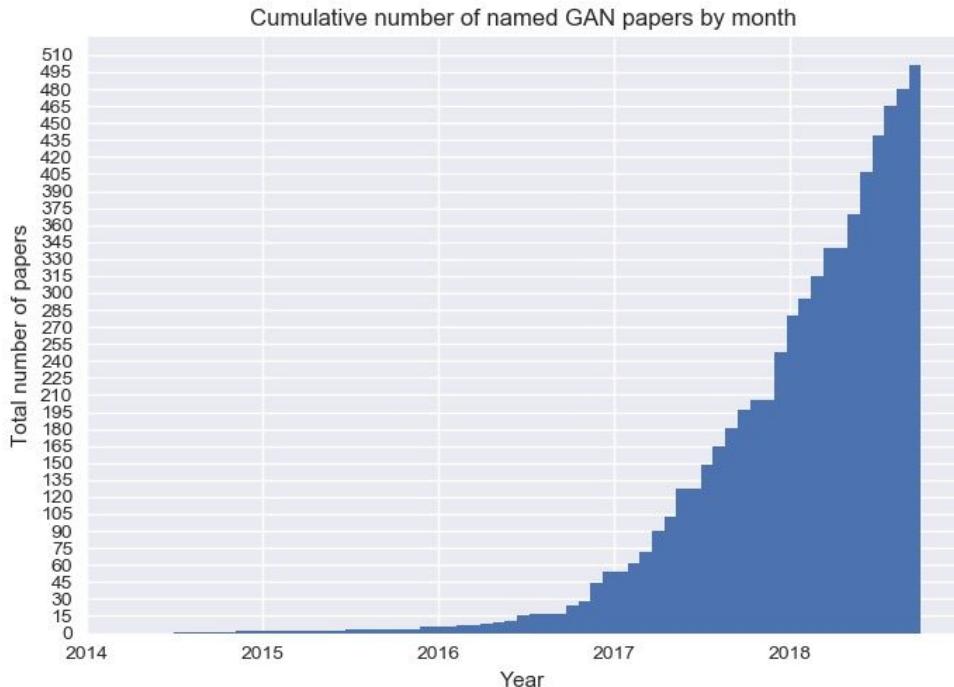
<https://arxiv.org/pdf/1705.07904v2.pdf>

# Latent space arithmetic

$$\text{cat} + (\text{dog} - \text{pig}) = \text{cat}$$
$$\text{pig} + (\text{cat} - \text{cat}) = \text{pig}$$

<https://arxiv.org/pdf/1704.03477.pdf>

# The GAN Zoo



<https://github.com/hindupuravinash/the-gan-zoo>

# Live GAN example applications

<https://thisxdoesnotexist.com/>

$x = \{ \text{person}, \text{cat}, \text{rental}, \text{waifu}, \text{url}, \text{startup}, \text{question}, \text{resume}, \text{emotion}, \dots \}$

## This X Does Not Exist

Using generative adversarial networks (GAN), we can learn how to create realistic-looking fake versions of almost anything, as shown by this collection of sites that have sprung up in the past month. Learn [how it works](#).



**This Person Does Not Exist**

The site that started it all, with the name



**This Cat Does Not Exist**

These purr-fect GAN-made cats will



**This Rental Does Not Exist**

Why bother trying to look for the perfect

# Classic GANs



pix2pix (2017)

CycleGAN (2017)

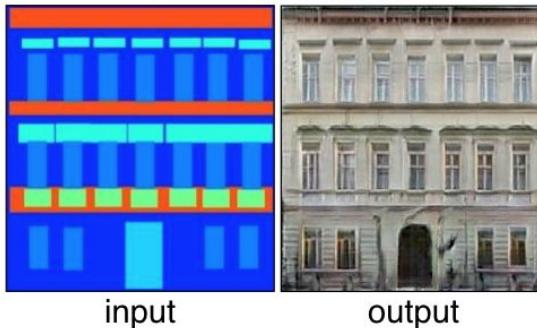
PGGAN (2018)

GauGAN (2019)

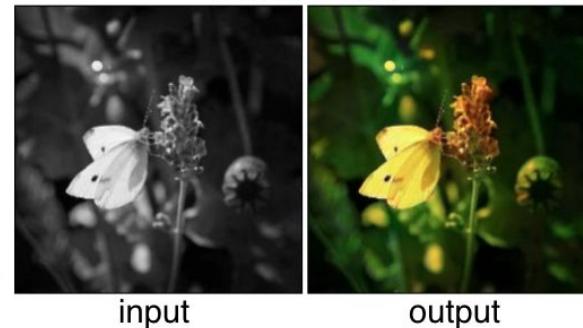
StyleGAN (2019, 2020, 2021)

# Classic GAN: pix2pix

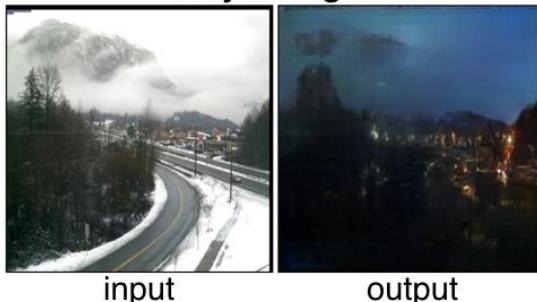
**Labels to Facade**



**BW to Color**



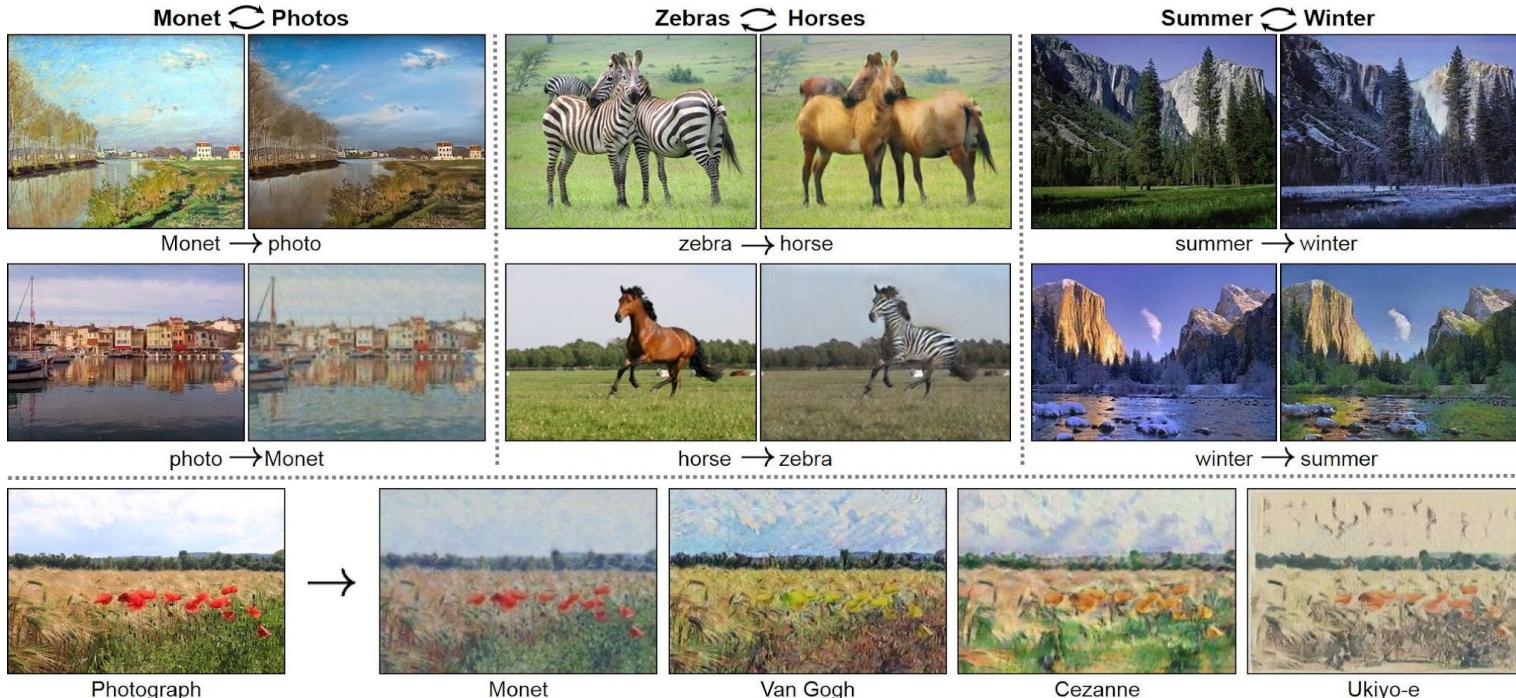
**Day to Night**



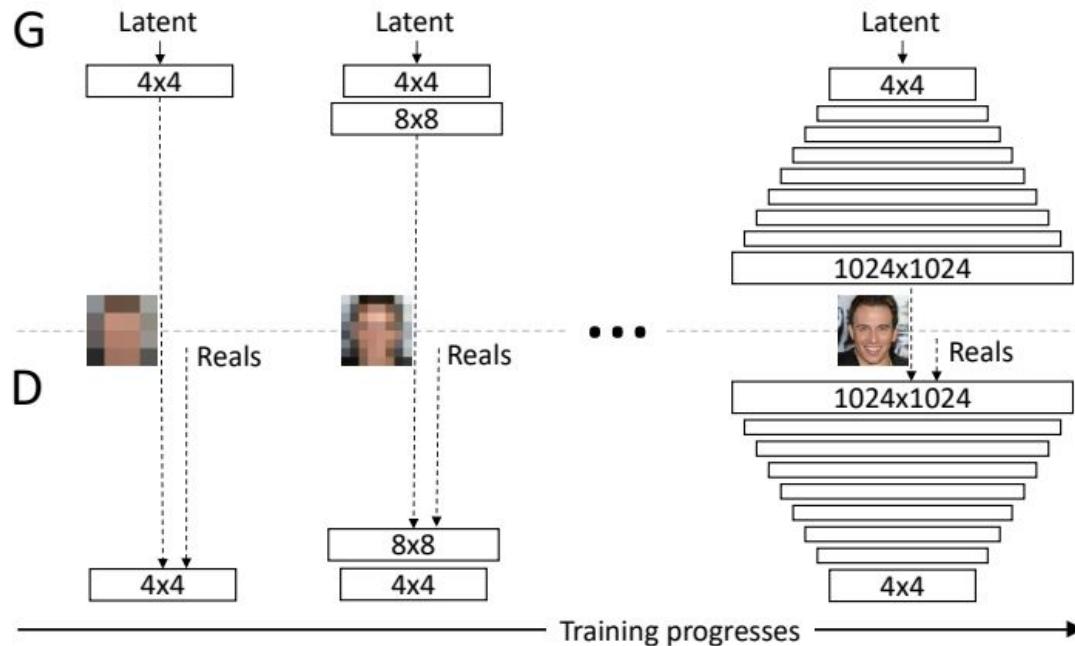
**Sketch to Photo**



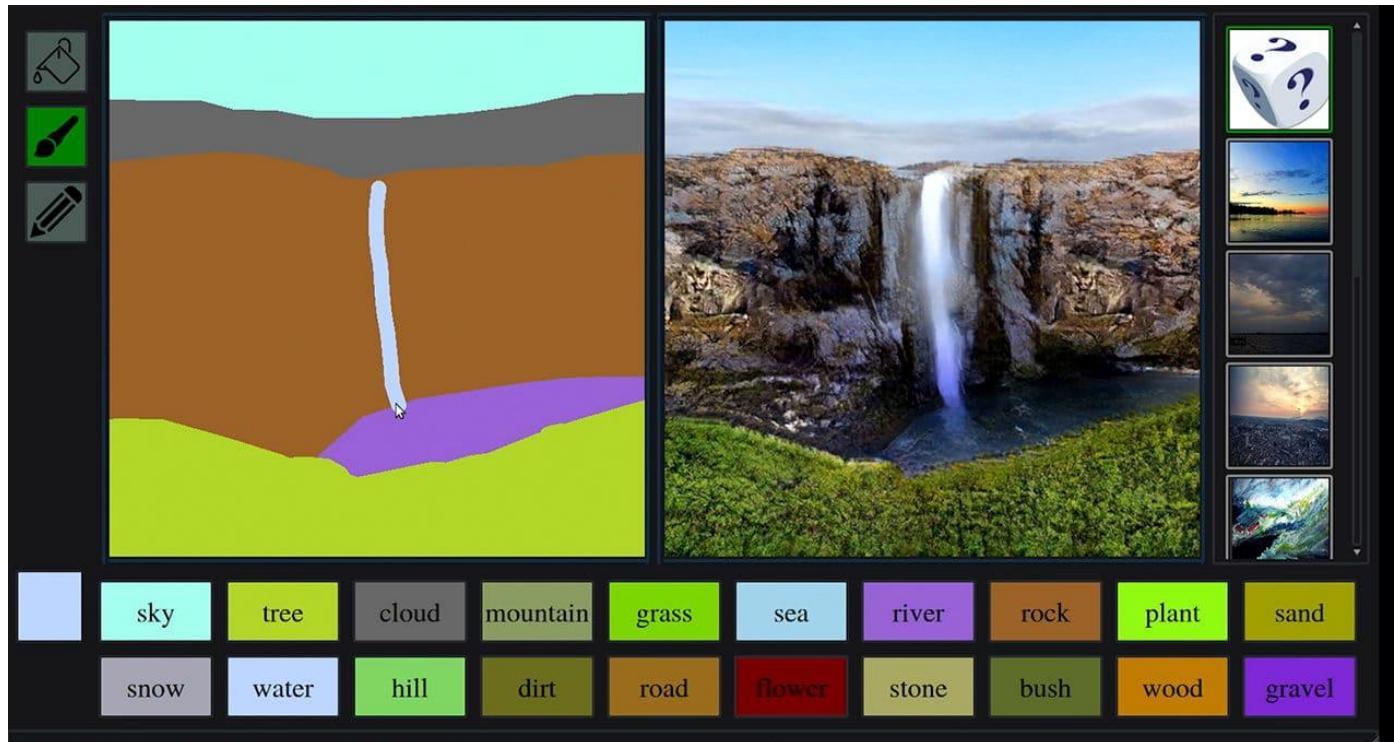
# Classic GAN: CycleGAN



# Classic GAN: PGGAN

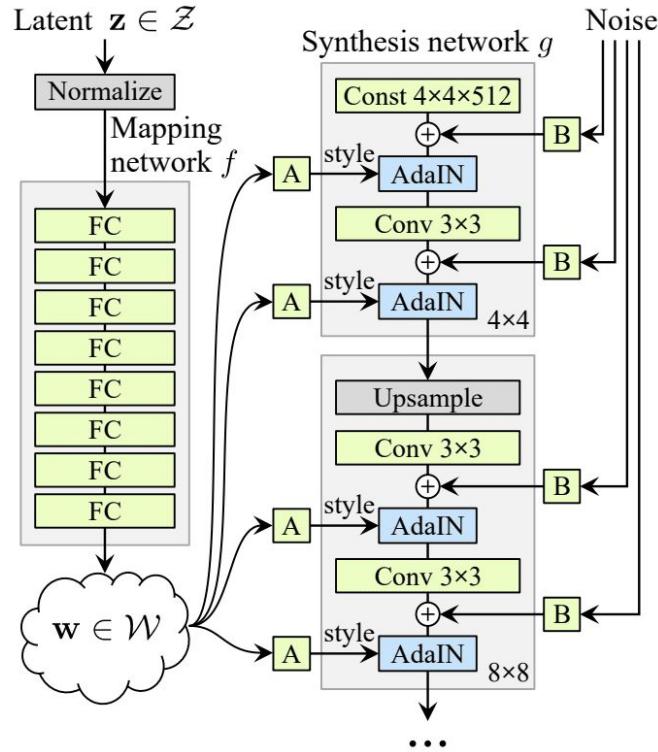


# Classic GAN: GauGAN



<https://blog.paperspace.com/nvidia-gaugan-introduction/>

# Classic GAN: StyleGAN



<https://machinelearningmastery.com/introduction-to-style-generative-adversarial-network-stylegan/>

# Current challenges in generative modeling



Weak evaluation measures: FID, IS, KID, AIS, MS-SSIM, etc.

Low coverage of latent spaces

Entangled latent features

Uncontrolled sampling

Sequential, discrete data is understudied

# Some tricks and tips

<https://towardsdatascience.com/c9071159628>

[https://medium.com/@jonathan\\_hui/819a86b3750b](https://medium.com/@jonathan_hui/819a86b3750b)

<https://medium.com/@utk.is.here/edd529764aa9>

<https://github.com/soumith/ganhacks>