

Reinforcement Learning

Prototyping with Deep Learning

Learning outcomes

After this lesson you will be able to:

- Understand the basics of Reinforcement Learning
- Understand the Q-Learning algorithm and the basics of Deep Q-Learning (DQN)
- Know some of the popular Deep Reinforcement Learning algorithms and their applications

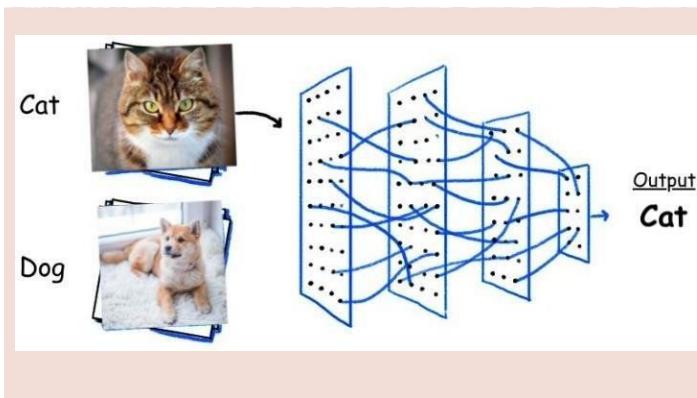
Recap: Classes of learning Problems

Supervised Learning

Data: (x, y)

Goal: Learn function to map

$$x \longrightarrow y$$

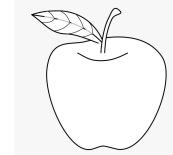


Unsupervised Learning

Data: x

x is data, No labels

Goal: Learn underlying structure



This thing is like the other thing

What is Reinforcement Learning?

Problems involving an **agent** interacting with an **environment** which provides (numeric) reward signals.

Data: State-action pairs

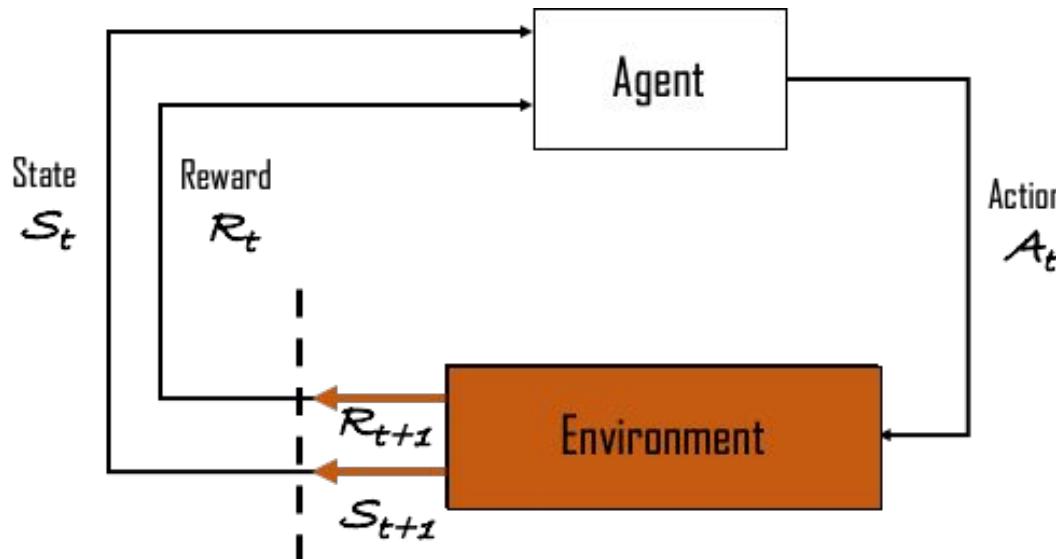
Goal: Learn how to take actions in order to maximize reward



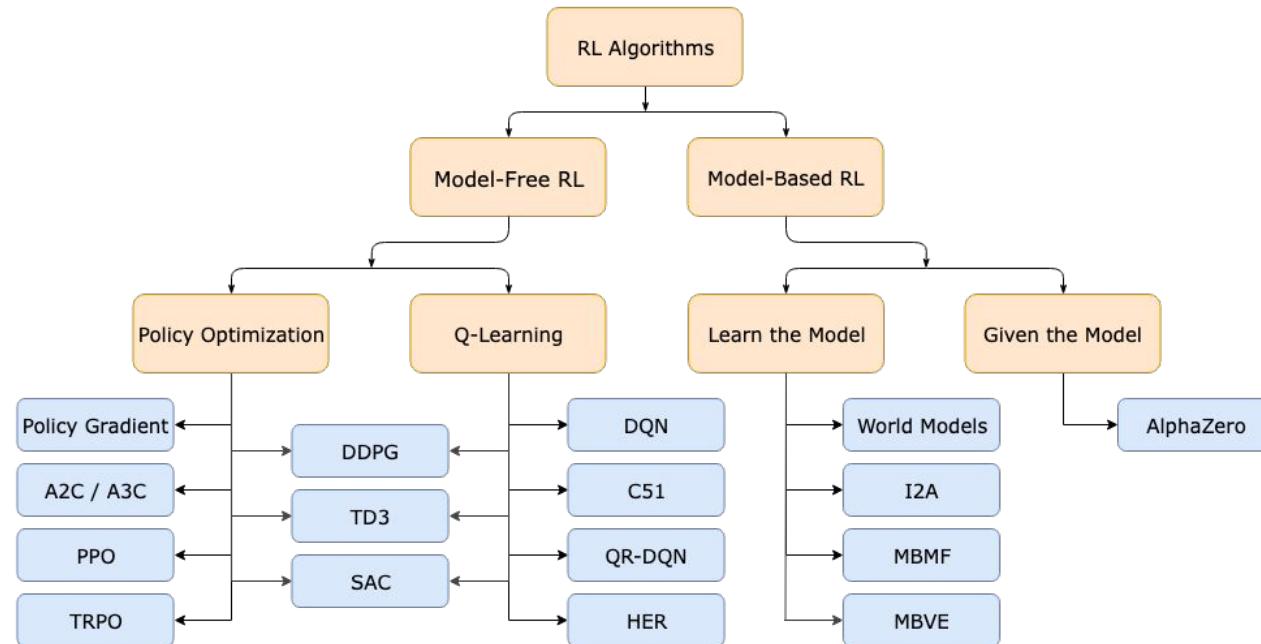
Eat this thing because it will keep you alive.

Classical RL

Agents interact with their environment through a sequence of **observations**, **actions** and **rewards**.

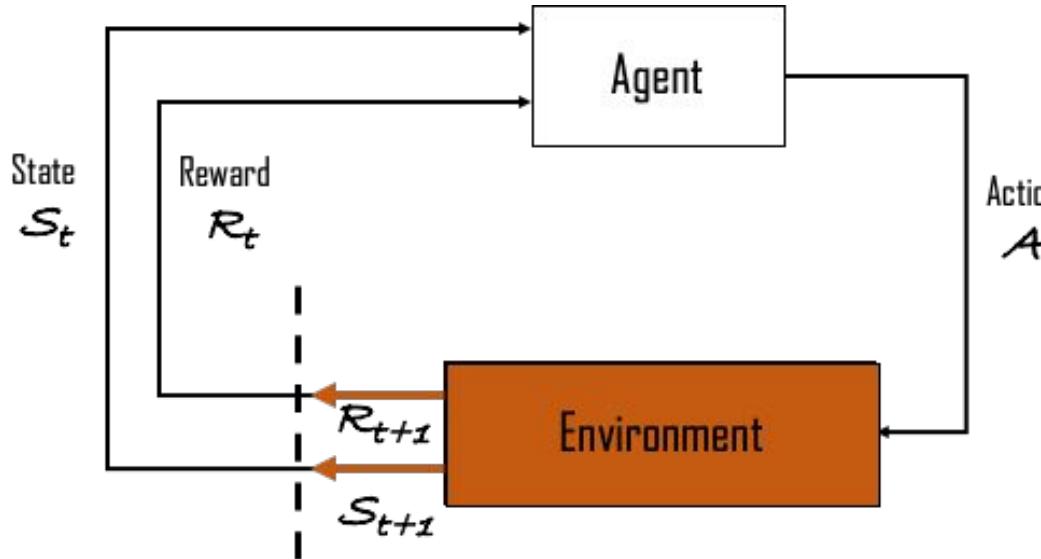


RL taxonomy



https://spinningup.openai.com/en/latest/spinningup/rl_intro2.html#citations-below

Classical RL



Reward: feedback that measures the success or failure of the agent's action

$$\text{Total Reward (Return)} \longrightarrow \sum_{i=t}^{\infty} r_i = r_t + r_{t+1} \dots + r_{t+n} + \dots$$

Total Reward
(Return) $\longrightarrow \sum_{i=t}^{\infty} r_i = r_t + r_{t+1} \dots + r_{t+n} + \dots$

Discounted
Total Reward
(Return) $R_t = \sum_{i=t}^{\infty} \gamma^i r_i = \gamma^t r_t + \gamma^{t+1} r_{t+1} \dots + \gamma^{t+n} r_{t+n} + \dots$

γ : discount factor ; $0 < \gamma < 1$, designed to make future rewards worth less

What is the probability that an agent will select a specific action in a given state?

- If an agent follows policy Π at time t , then $\Pi(a|s)$ is the probability that $A_t = a$ if $S_t = s$.
- This means that, at time t , under policy Π , the probability of taking action a in state s is $\Pi(a|s)$.

Note that, for each state $s \in S$, Π is a probability distribution over $a \in A(s)$

How do we measure the quality of actions?

Q-function

$$Q(s_t, a_t) = \mathbb{E}[R_t | s_t, a_t]$$

Q-function captures the **expected total future reward** an agent in **state s_t** can receive by executing a certain **action a_t**

How useful is a given action in gaining some future reward.

- *Good actions result high Q value*
- *Bad actions result low Q value*

Q-function

- The optimal Q-value function Q^* is the maximum expected cumulative reward achievable from a given (state, action) pair

$$Q^*(s, a) = \max_{\pi} \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t | s_0 = s, a_0 = a, \pi \right]$$

- The maximum sum of rewards r_t discounted by γ at each time step t , achievable by a policy $\Pi = p(a|s)$.
- The agent operates based on a policy Π to approximate Q-values (state-action pairs) that maximize a future reward.

This is done by enforcing the **Bellman equation**:

Reinforcement Learning

Bellman equation:

$$Q^*(s, a) = \mathbb{E}_{s' \sim \varepsilon} \left[r + \gamma \max_{a'} Q^*(s', a') | s, a \right]$$

Given any state-action pair **(s, a)** the maximum cumulative reward achieved is the sum of the reward for that pair **r** plus the value of the next state we end up with, **s'**.

The value at state **s'** is going to be the maximum over actions **a'** at **Q*(s',a')**.

- The objective of **RL** is to find an optimal **policy** in a sense that the expected return over all successive time steps is the **maximum** achievable.
- This can be done by learning the **optimal Q-values** for each state-action pairs.

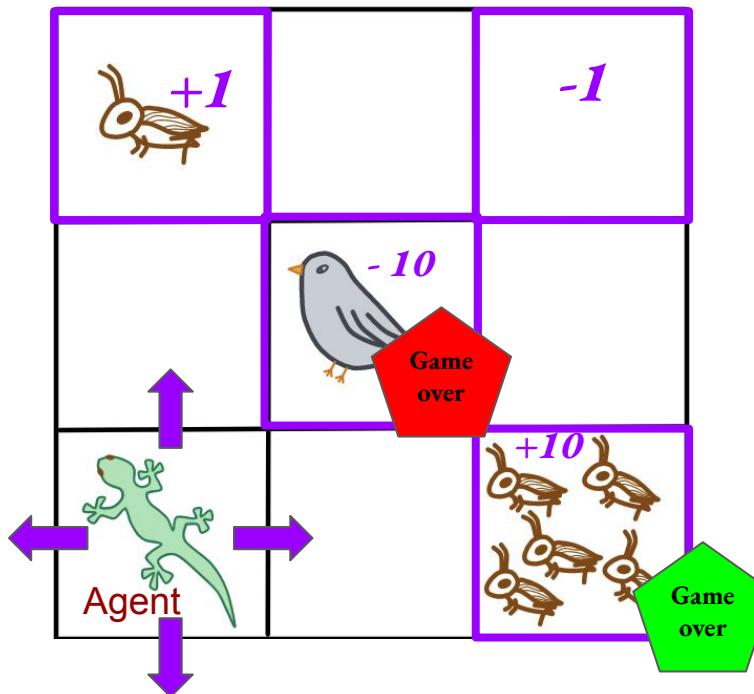
Q-Learning



Q-Learning



The lizard wants to eat as many crickets as possible in the least amount of time without stumbling across a bird, which will, itself, eat the lizard.

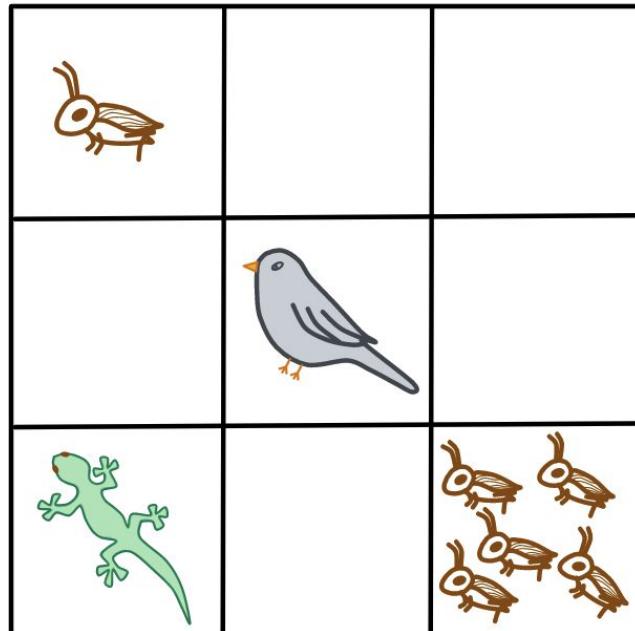


- States determined by individual tiles and where the agent is at a given time

Q-Learning



At the beginning the lizard has no idea of how good any action is at any given state



States

	Actions			
	Left	Right	Up	Down
1 cricket	0	0	0	0
Empty 1	0	0	0	0
Empty 2	0	0	0	0
Empty 3	0	0	0	0
Bird	0	0	0	0
Empty 4	0	0	0	0
Empty 5	0	0	0	0
Empty 6	0	X+10	0	0
5 crickets	0	0	0	0

Q-table

Q-Learning (Exploration Vs Exploitation)

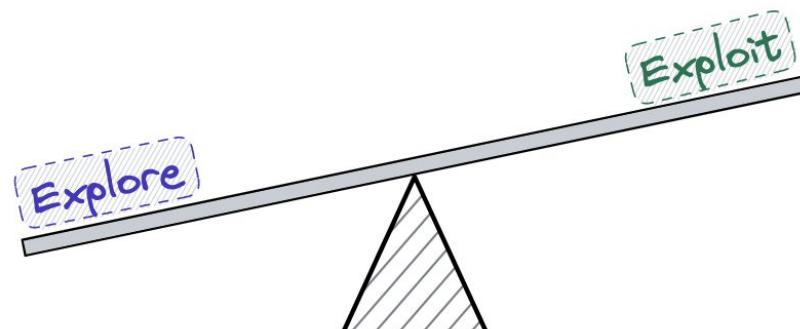


Exploration: the act of exploring the environment in order to find out information about it.

Exploitation: making use of the information that is already known about the environment in order to maximize the return. **Q-table look up.**

+1		
-1		


```
if random_num > epsilon:  
    # choose action via exploitation  
else:  
    # choose action via exploration
```



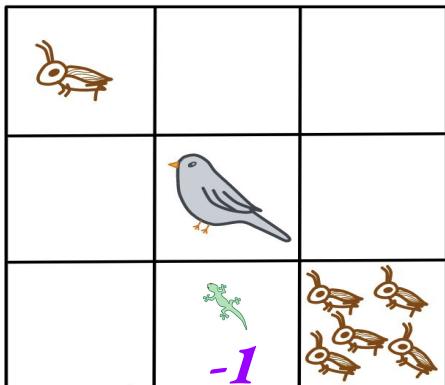
Greedy Epsilon strategy

Q-Learning (Updating Q-table)

- To update the Q-value We use the bellman equation
- We want to make the $Q(s,a)$ for any state action pair as close as possible to the right hand side of the Bellman equation.
- The formula to calculate new Q-values:

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \overbrace{\left(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\substack{\text{learned value} \\ \text{estimate of optimal future value}}} \right)}^{\text{learned value}}$$

Q-Learning (Updating Q-table)



discount rate $\gamma = 0.99$
 learning rate $\alpha = 0.7$

States	Actions			
	Left	Right	Up	Down
1 cricket	0	0	0	0
Empty 1	0	0	0	0
Empty 2	0	0	0	0
Empty 3	0	0	0	0
Bird	0	0	0	0
Empty 4	0	0	0	0
Empty 5	0	X -0.7	0	0
Empty 6	0	0	0	0
5 crickets	0	0	0	0

$$\begin{aligned}
 Q^{new}(s, a) &= (1 - \alpha) \underbrace{Q(s, a)}_{\text{old value}} + \alpha \overbrace{\left(R_{t+1} + \gamma \max_{a'} Q(s', a') \right)}^{\text{new value}} \\
 &= (1 - 0.7)(0) + 0.7 \left(-1 + 0.99 \left(\max_{a'} Q(s', a') \right) \right) \\
 &= (1 - 0.7)(0) + 0.7(-1 + 0.99(0)) \\
 &= 0 + 0.7(-1) \\
 &= -0.7
 \end{aligned}$$

$$\begin{aligned}
 &= \max(Q(\text{empty6}, \text{left}), Q(\text{empty6}, \text{right}), Q(\text{empty6}, \text{up}), Q(\text{empty6}, \text{down})) \\
 &= \max(0, 0, 0, 0) \\
 &= 0
 \end{aligned}$$

Reinforcement Learning

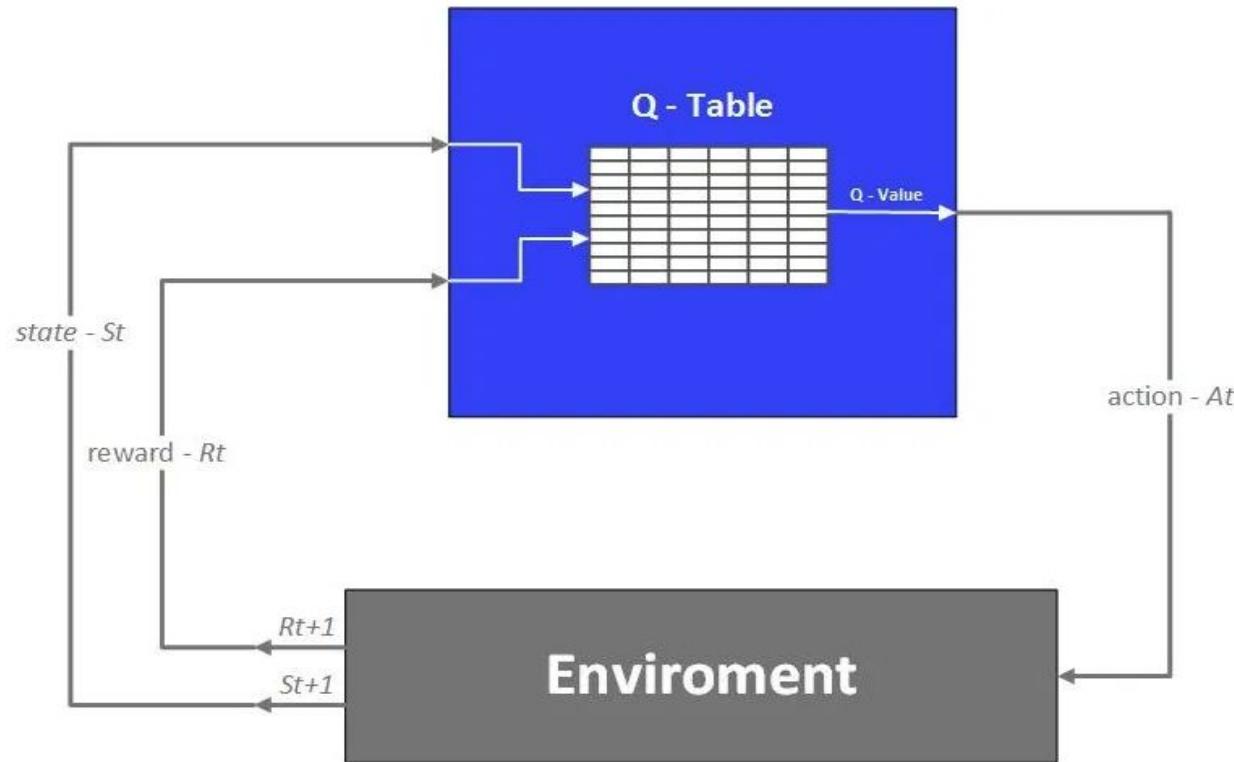
Learning: Refine Q table by approximating the optimal Q-values

Q-Table		Actions				
States		Action 1	Action 2	...	Action n-1	Action n
	State 1	0.789112	0.745642	...	0.212485	0.256545
	State 2	5.123455	5.11565	...	5.156545	4.155612

	State n-1	2.156454	2.15567	...	2.144423	2.454658
	State n	6.156212	6.154556	...	6.145441	6.444444

- Initially the agent will do a bit of *exploration*. It selects random values. Over time when it reaches rewards it will slowly learn to exploit those rewards for future action.

Reinforcement Learning



Reinforcement Learning



Objective: Complete the game with the highest score

State: Raw pixel inputs of the game state

Action: Game controls e.g. Left, Right, Up, Down

Reward: Score increase/decrease at each time step

What is the problem with this?

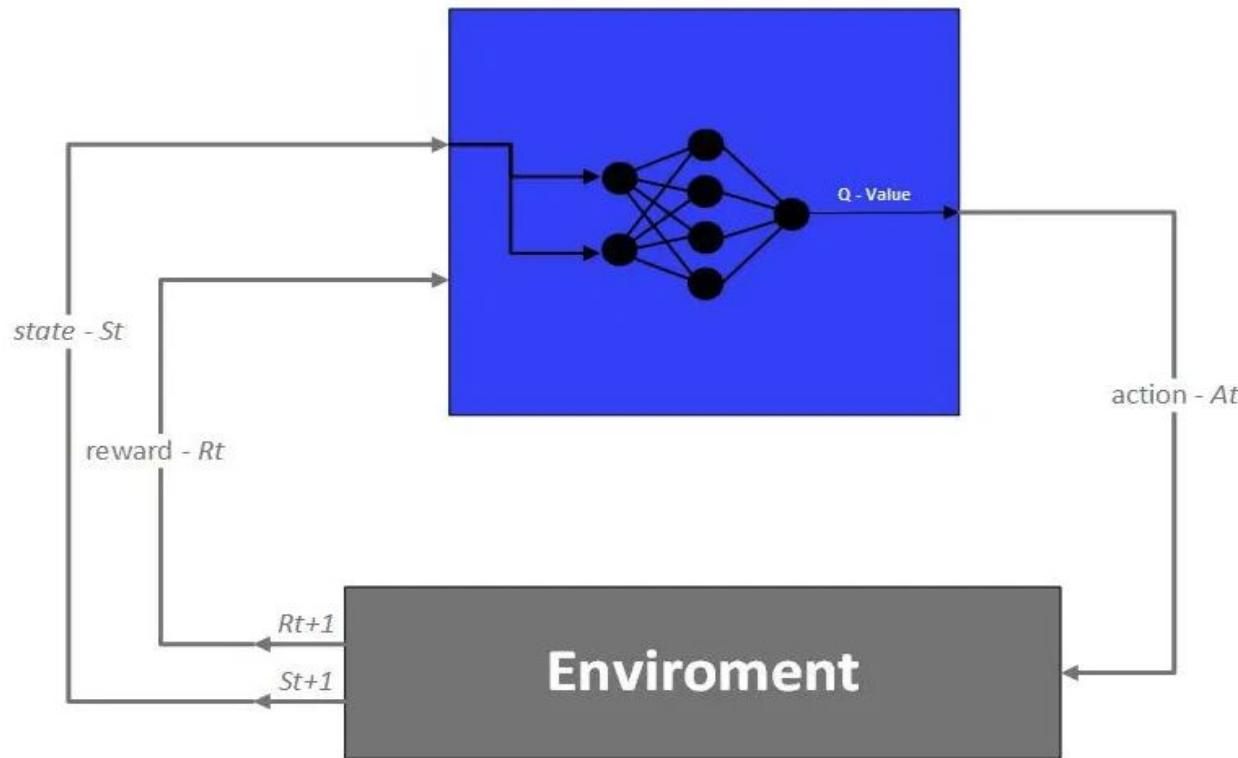
Not scalable: we must compute $Q(s, a)$ for every state-action pair.

Computationally expensive to compute for the entire state space, perhaps infeasible.

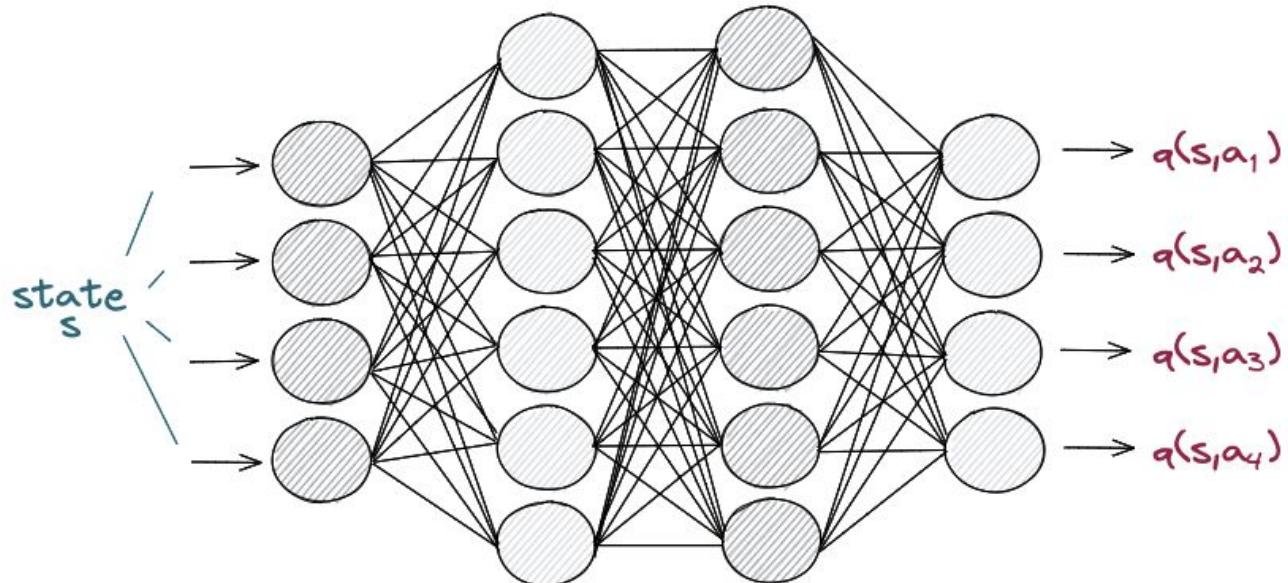
Solution: Use function approximator to estimate the value of $Q(s, a)$,

- Neural Network → **Deep Q-learning**

Deep Reinforcement Learning



Deep Reinforcement Learning



Deep Reinforcement Learning

Forward pass: loss function tries to Minimise the error of the bellman equation.

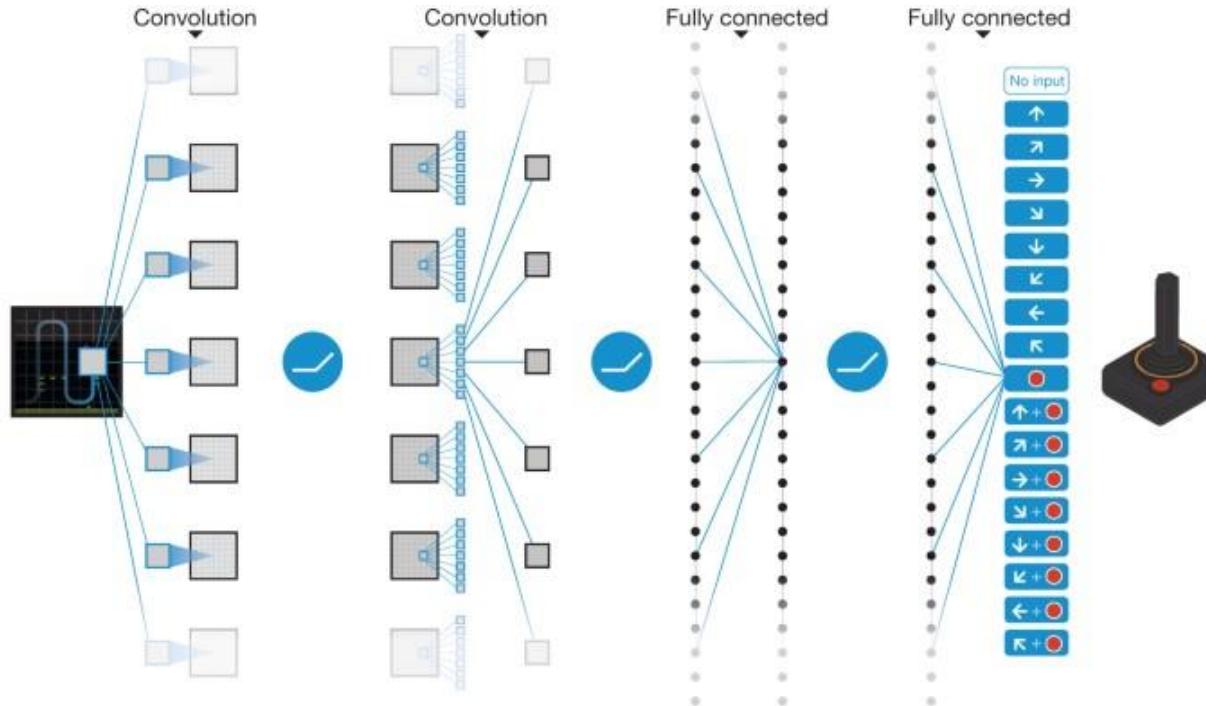
Backward pass: gradient update with respect to the Q-function parameters θ .

$$Q(s,a;\theta) \approx Q^*(s,a)$$

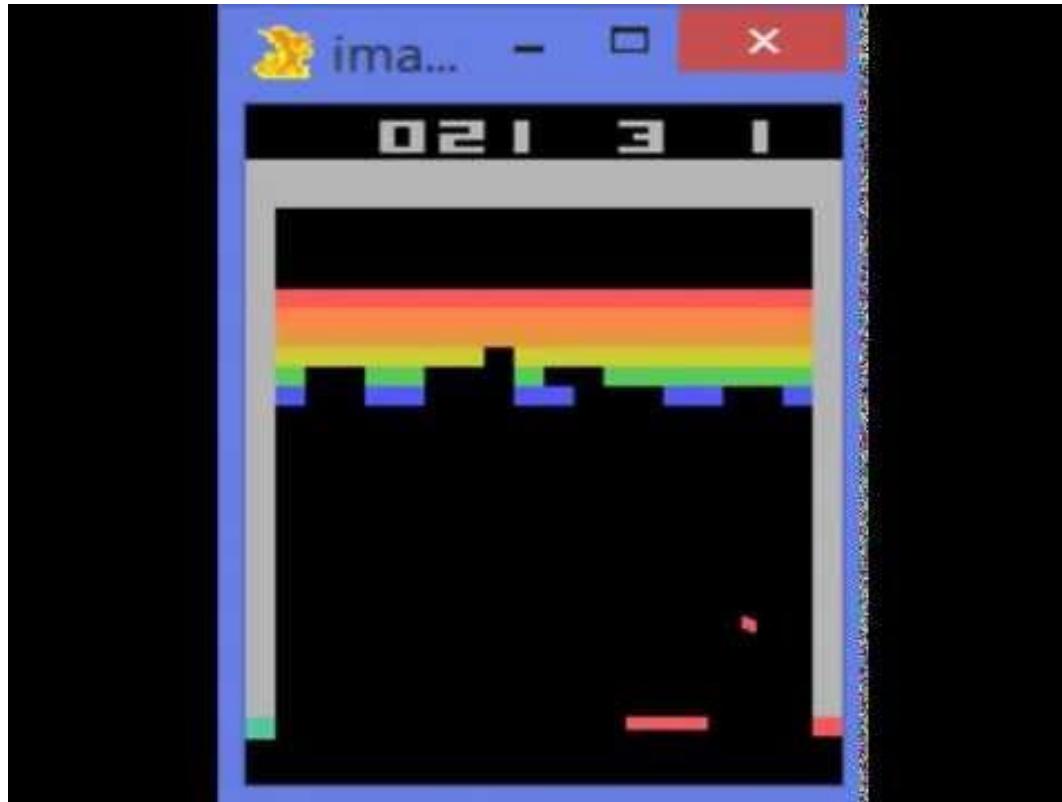
The optimal policy $\Pi^*(s) = \text{argmax } Q(s,a)$

Deep Reinforcement Learning

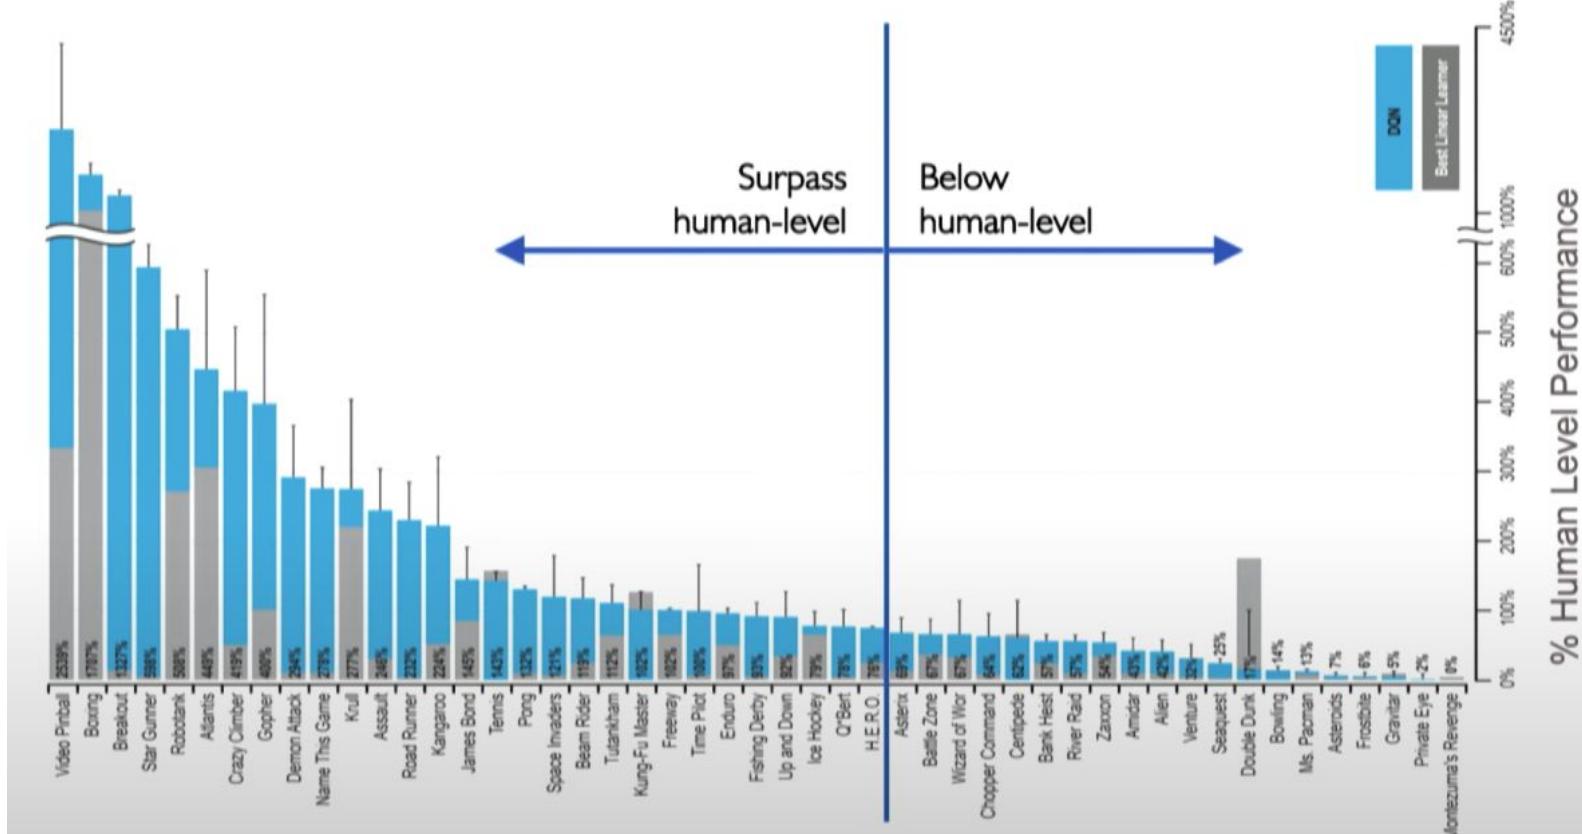
- DeepMind Human-level control through deep reinforcement learning



Deep Reinforcement Learning



Deep Reinforcement Learning



Deep Reinforcement Learning

Limitations of Q-learning

- Can not handle **continuous action spaces**: limited to scenarios where we can define the action space in **discrete** and **small** pieces

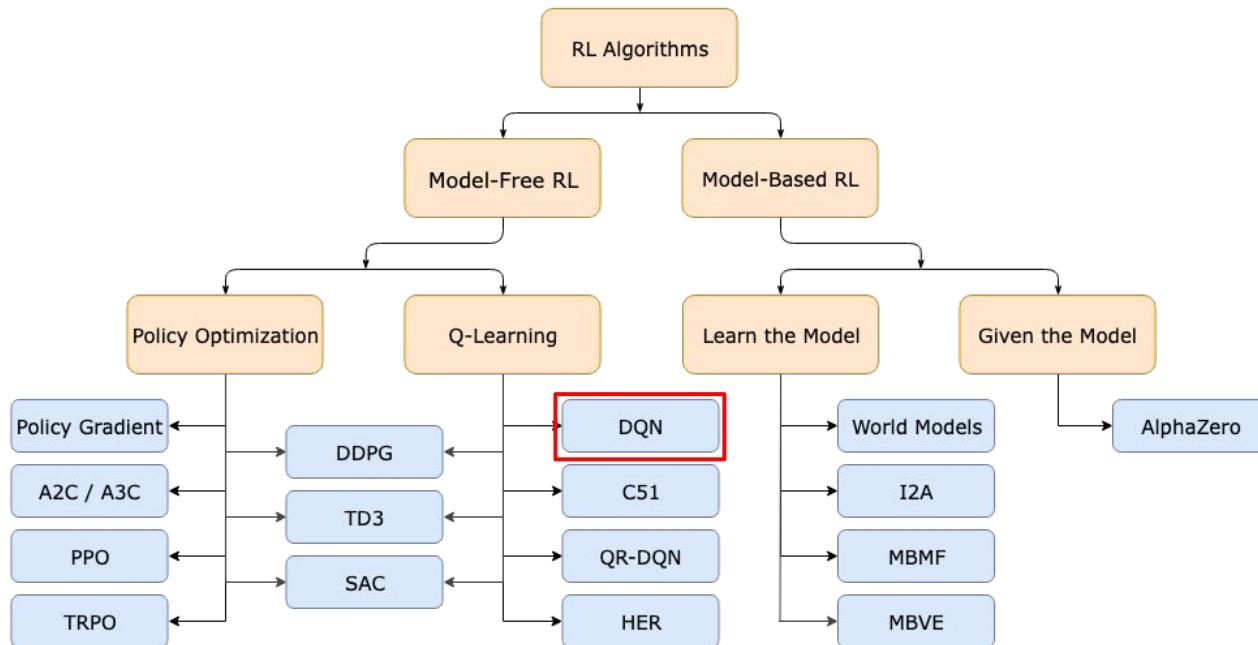


Limitations of Q-learning

- Policy is deterministically computed from the Q function by maximising the reward : can not learn stochastic policies.

$$\Pi^*(s) = \operatorname{argmax} Q(s,a)$$

Deep Reinforcement Learning



Value Learning

- Find $Q(s, a)$
- $a = \text{argmax } Q(s, a)$

Policy Learning

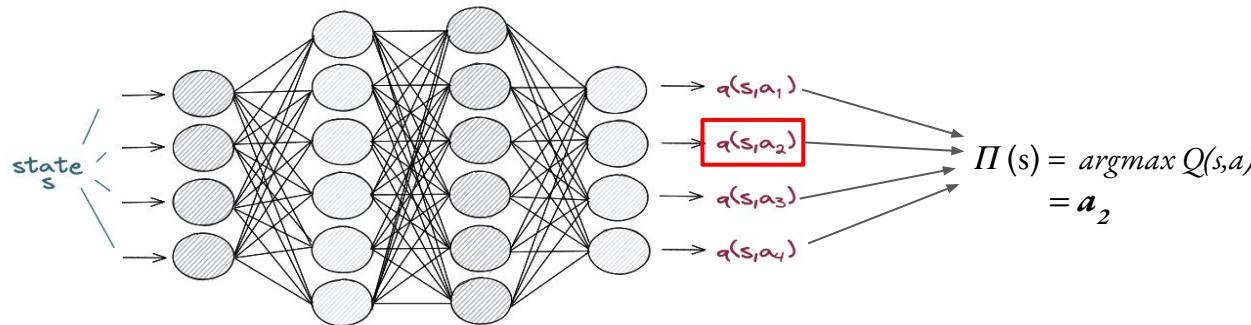
- Find $\pi(s)$
- Sample $a \sim \pi(s)$



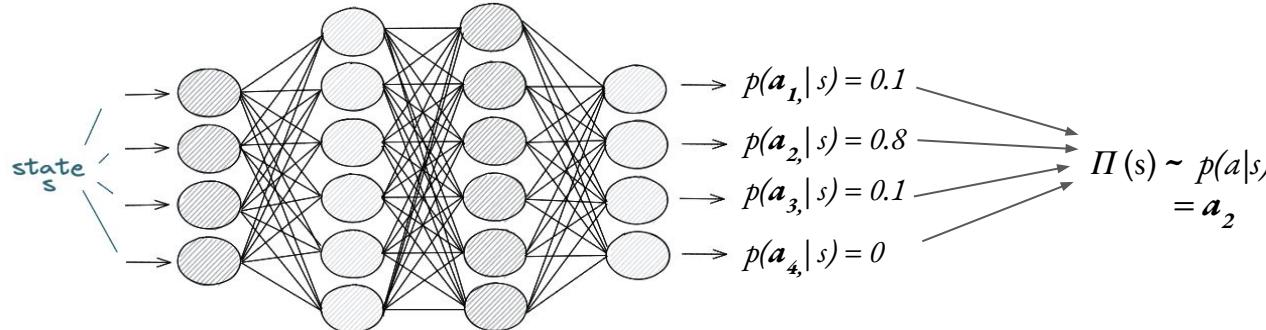
Learning a policy can be much simpler

Deep Reinforcement Learning

DQN: Approximate Q-function then use it to infer the optimal policy, $\pi(s)$.



Policy gradient: directly optimise the policy, $\pi(s)$.



Deep Reinforcement Learning

Policy gradient: Handle continuous action spaces.

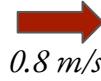
Discrete action space



$$p(a|s)$$

State (s) →

Continuous action space



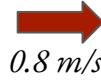
$$\begin{aligned} &\text{Mean, } \mu \\ &\text{Variance, } \sigma^2 \end{aligned}$$

$$p(a|s)$$

1

Faster
Right

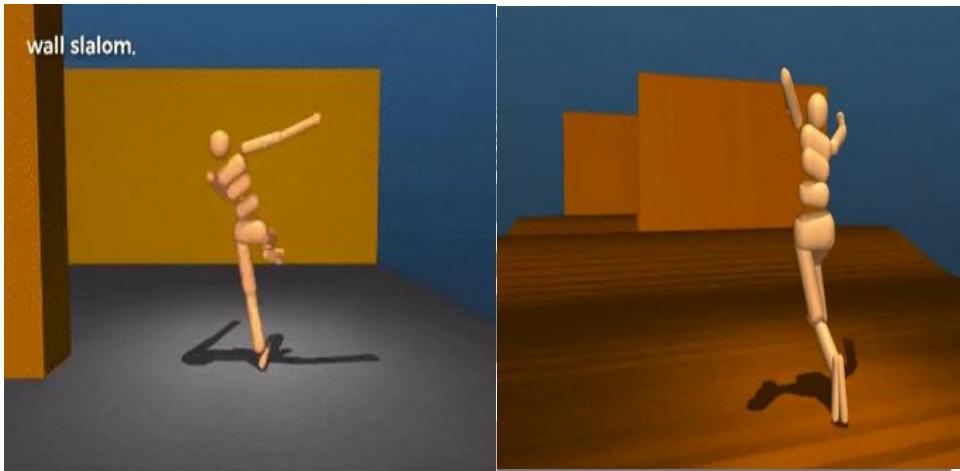
Faster
Left



Left Stay Right



Deep Reinforcement Learning



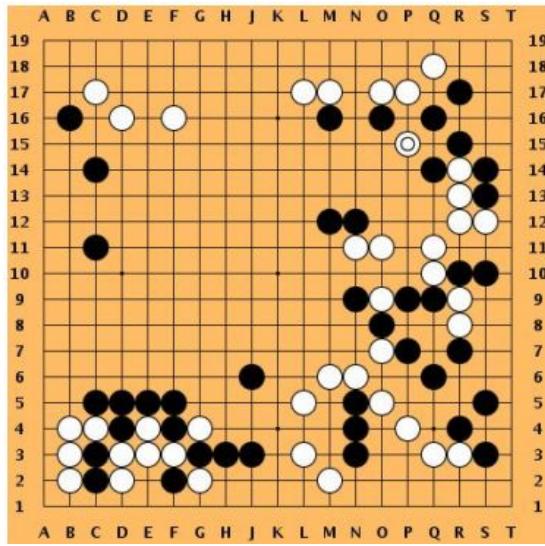
Objective: Make the robot move forward

State: Angle and position of the joints

Action: Torques applied on joints

Reward: 1 at each time step upright +
forward movement

Deep Reinforcement Learning



Objective: Win the game!

State: Position of all pieces

Action: Where to put the next piece down

Reward: 1 if win at the end of the game, 0 otherwise

Deep Reinforcement Learning

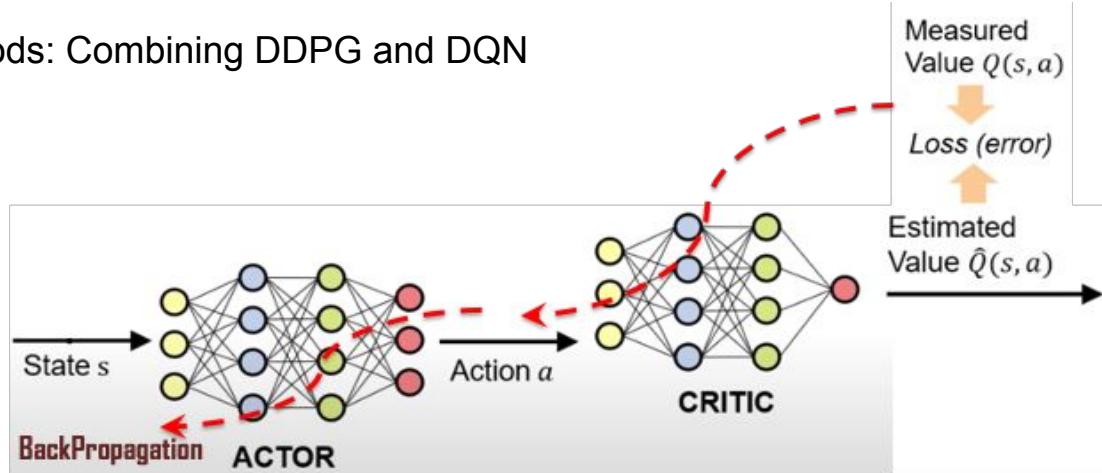
Continuous control with deep reinforcement learning



Actor critic method works well when we have both an **infinite input space** and infinite output space

Deep Reinforcement Learning

Actor-Critic Methods: Combining DDPG and DQN

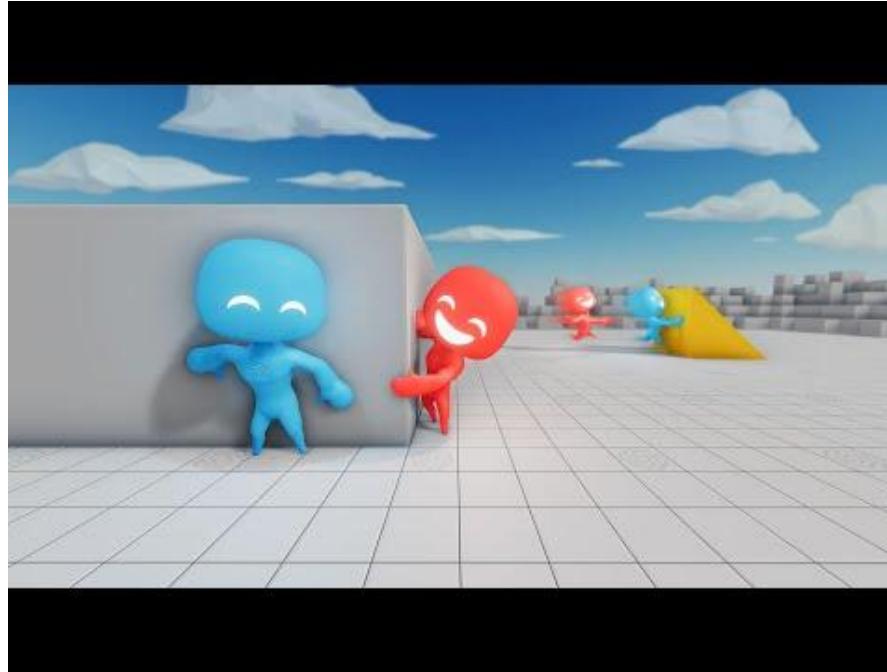


Actor approximates policy $p(a|s)$

Critic approximates Q-values

- Quite natural in the human's world **Child as an actor and parent as a Critic** .

Deep Reinforcement Learning



- Multi-Agent Actor-Critic for Mixed Cooperative Competitive Environments
- Emergent Tool Use From Multi-Agent Autocurricula

Hands-on

- **OpenAI gym:** a toolkit for developing and comparing reinforcement learning algorithms. <https://gym.openai.com/>
- **Stable Baselines3 (SB3):** a set of reliable implementations of reinforcement learning algorithms in PyTorch.

<https://stable-baselines3.readthedocs.io/en/master/>