

Deployment

Prototyping with Deep Learning

Learning outcomes

- Identify best practices in model deployment
- Save, load, and use your DL models

What can we do with Jupyter notebooks?



Add [interactive widgets](#)

Create [slideshow presentations](#)

Write [technical blogs](#)

... But software **deployment** is something else

Goals

Use trained models in **production** environments


Also monitor models usage over time

Also improve models with new data

Why are ML systems hard?


Some key issues:

- Entanglement
- Data dependencies
- Configuration
- Data preparation
- Tracking model errors
- Separation of expertise




[https://christophergs.github.io/machine%20learning/2019/03/17/
how-to-deploy-machine-learning-models/](https://christophergs.github.io/machine%20learning/2019/03/17/how-to-deploy-machine-learning-models/)

MLOps pipeline




The Hidden Technical Debt in ML Systems



<https://papers.nips.cc/paper/5656-hidden-technical-debt-in-machine-learning-systems.pdf>

Biases everywhere



Key design principles


Plan for **reproducibility**

ML as part of your **build pipeline**


Plan for **extensibility**

Plan for **modularity**

Testing, testing, testing



Some TF tools



<https://blog.tensorflow.org/2019/09/tensorflow-20-is-now-available.html>

Saving trained models

In Keras

```
model.save(...)
```

Only architecture: `model.to_json()` or
`model.to_yaml()`

Only weights: `model.save_weights()`

In Pytorch

```
torch.save(model.state_dict(), ...)
```

<https://keras.io/getting-started/faq/#how-can-i-save-a-keras-model>

https://pytorch.org/tutorials/beginner/saving_loading_models.html

What can (and should) be saved?



Model architecture

Model weights

Training configuration: loss function, epochs

Optimizer: learning rate, state (to resume training)

Model file formats



.pt or .pth

Pytorch

.pb

Protobuf (TF native)

.tflite

Tensorflow, on-device

.h5

HDF, cross-platform

.json

TF.js

Don't use Python's pickle!

Model format conversion

<https://github.com/tensorflow/tfjs/tree/master/tfjs-converter>

```
~$ tflite_convert --graph_def_file=model.pb --output_file=model.tflite
~$ tflite_convert --keras_model_file=model.h5 --output_file=model.tflite
~$ tensorflowjs_converter --input_format keras model.h5 /path/to/dir
```

Name	Headers	Preview	Response	Timing
<input type="checkbox"/> messages.json	General			
<input type="checkbox"/> model.json				
<input type="checkbox"/> group1-shard1of1				
<input type="checkbox"/> group2-shard1of1				

Request URL: http://localhost:1234/json/model.json
Request Method: GET
Status Code: 200 OK

Model serving with Flask (web APIs)



```
from flask import Flask, request, jsonify

from tensorflow.keras.models import load_model
# from torch import load as load_model

model = load_model('model.h5', compile=False)

app = Flask(__name__)

@app.route('/predict', methods=['POST'])
def predict():
    # TODO: Read data from `request` object.
    return jsonify(model.predict(...))

# Init service.
app.run(port=50005)
```

<https://flask.palletsprojects.com/en/2.1.x/>

Model serving with Gradio (graphical UIs)



```
import gradio as gr
import numpy as np
from flash.image import ImageClassifier

# 1. Load Model
model = ImageClassifier.load_from_checkpoint("image_classification_model.pt")

# 2. Define Classification Function
def classify(img):
    img = np.transpose(img, (2, 0, 1))
    return model.predict(img, data_source="numpy")[0]

#Init Gradio
image = gr.inputs.Image(shape=(299, 299))
label = gr.outputs.Label(num_top_classes=1)
gr.Interface(fn=classify, inputs=image, outputs=label, capture_session=True).launch()
```

https://gradio.app/getting_started/

Monitoring



<https://wandb.ai/site>

Some resources

Best practices:

<https://opendatascience.com/best-practices-for-deploying-machine-learning-in-the-enterprise/>

Challenges: <https://towardsdatascience.com/443af67493cd>

CI/CD in ML: <https://martinfowler.com/articles/cd4ml.html>

DL in production (book):

<https://github.com/The-AI-Summer/Deep-Learning-In-Production>