

MS_DS-19 - Introduction to Deep Learning for Image Analysis and Computer Vision with Examples in Medical Applications

Lecture 2 – Digital Images and Basic Operations

Andreas Husch, PhD

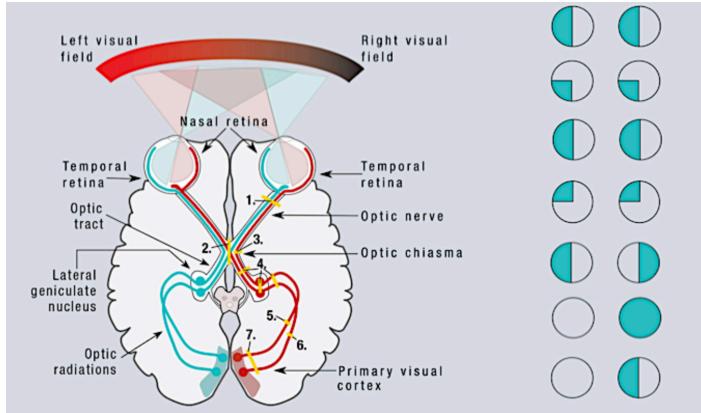
Luxembourg Centre for Systems Biomedicine
Interventional Neuroscience Group

andreas.husch@uni.lu



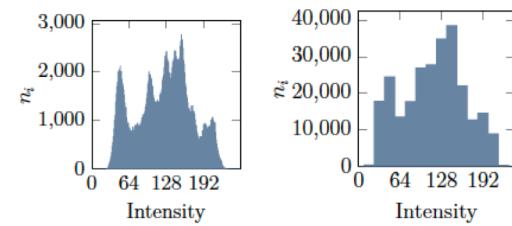
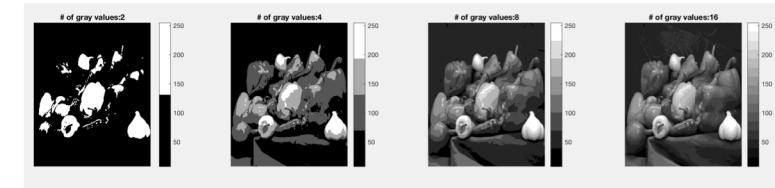
Recap last lecture

- Human Visual System
- Images as Functions
 - Multiple equivalent discrete representations
- Indexed Images
 - Colormaps (Look-up-table)
- Sampling and Quantization
 - Resolution
 - Bit-Depth
- Histograms
 - Look-up-tables
 - Window/Level



$$B_t = (r, g, b)(x, y), \text{ with } r, g, b \in \{0, \dots, 255\}$$

$$B_c = b(x, y, c) \text{ with } b \in \{0, \dots, 255\} \text{ and } c \in \{1, 2, 3\}$$



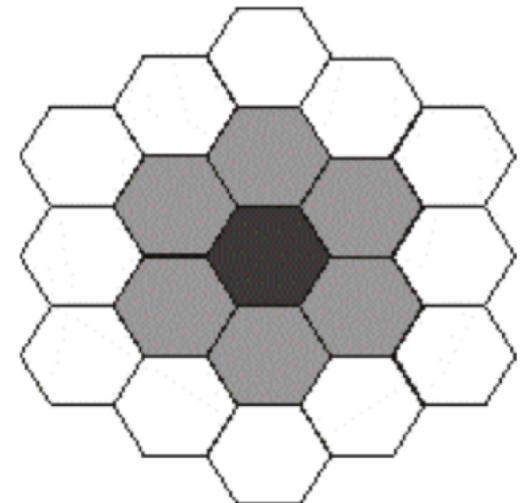
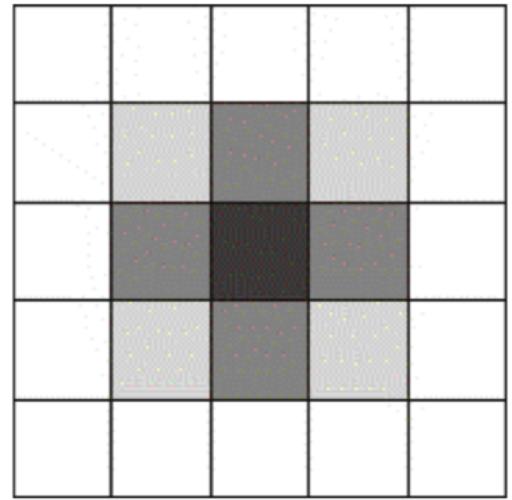
(a) Example grayscale image.

(b) Histogram of the example image with 256 bins.

(c) Histogram of the example image with 16 bins.

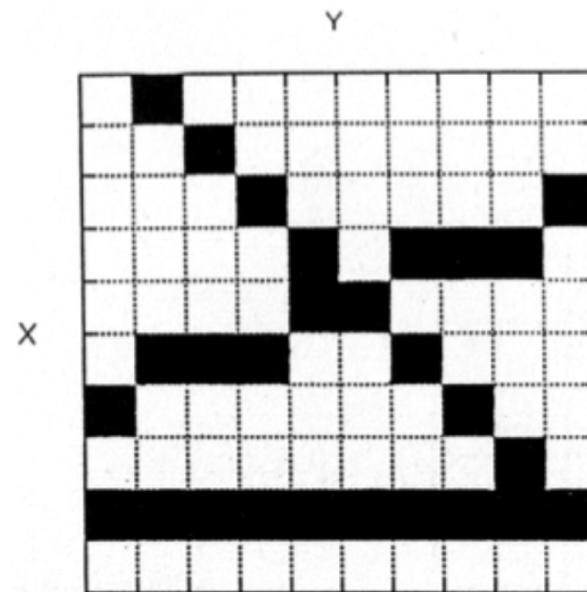
Image Grids and Neighborhood relations

- **Cartesian grids**
 - **Square** grid: simplest and most common form; regular grid over the image.
 - **Non-Square** (rectangular) grid: **anisotropic pixels**: $\Delta x \neq \Delta y$
 - frequent in medical imaging! (why?)
- **Non-Cartesian**, for example **Hexagonal grid**: Each pixel has exactly six neighbors.
 - Neighboring pixels have exactly one common edge.
 - Discrete geometry is simpler (distances, contour length, area of objects, etc.).



Digital Image Grids

- Strengths and weaknesses of discrete geometries: e.g. on the problem of connected components



Neighborhood (in 2D)

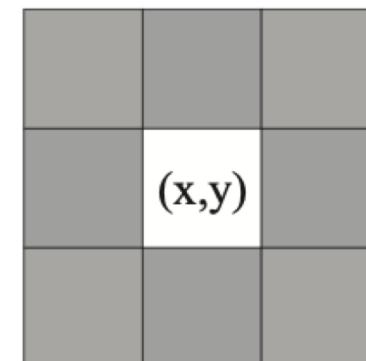
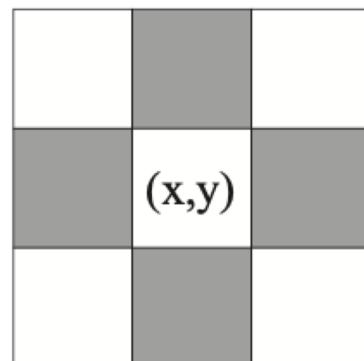
4-neighborhood: The four direct neighbors of a central pixel

- central pixel (x,y) form the 4-neighborhood $N_4(x,y)$. Coordinate list:

$$N_4 = \{(x,y - 1), (x - 1, y), (x, y + 1), (x + 1, y)\} .$$

8-neighborhood: The four direct neighbors plus the four diagonal neighbors of a central pixel (x,y) form the 8-neighborhood $N_8(x,y)$

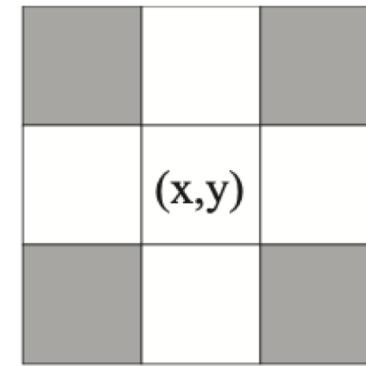
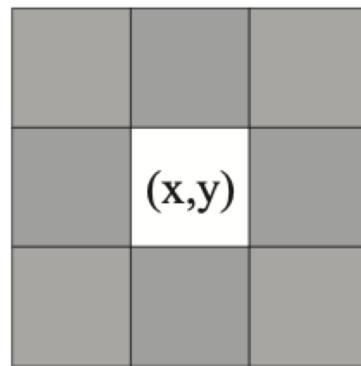
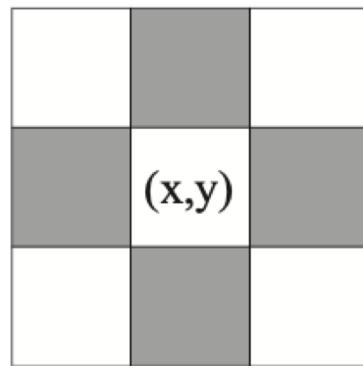
$$N_8 = \{(x - 1, y - 1), (x - 1, y), (x - 1, y + 1), \\ (x, y - 1), (x, y + 1), (x + 1, y - 1), (x + 1, y), (x + 1, y + 1)\}$$



Neighborhood (in 2D)

D-neighborhood: The **diagonal** neighbors of a central pixel (x,y) for the D-neighborhood $ND(x,y)$

$$ND = \{(x - 1, y - 1), (x - 1, y + 1), (x + 1, y - 1), (x + 1, y + 1)\}$$



Neighborhood – other dimensions

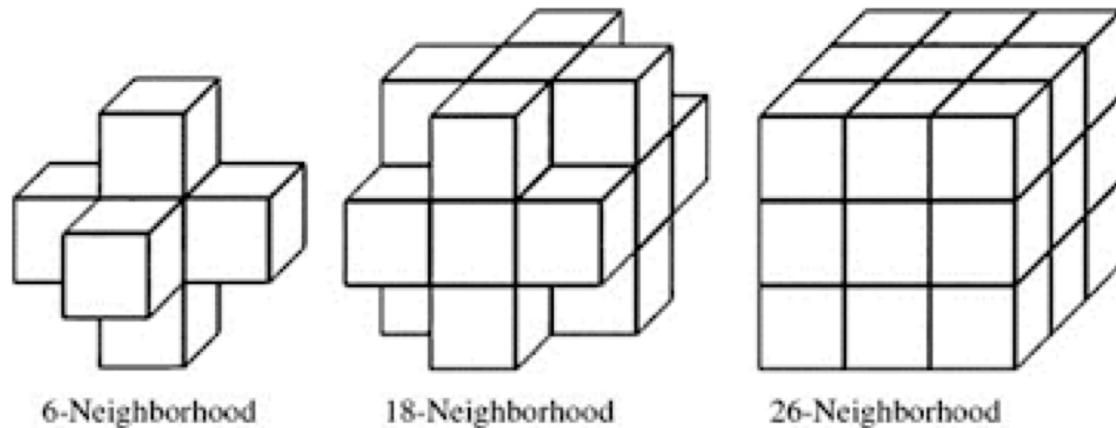
How does the equivalent of a 4-Neighborhood look

- In 1D?
- In 3D?

How does the equivalent of a 8-Neighborhood look

- In 1D?
- In 3D?

What about applications
in Deep Learning?



Distance

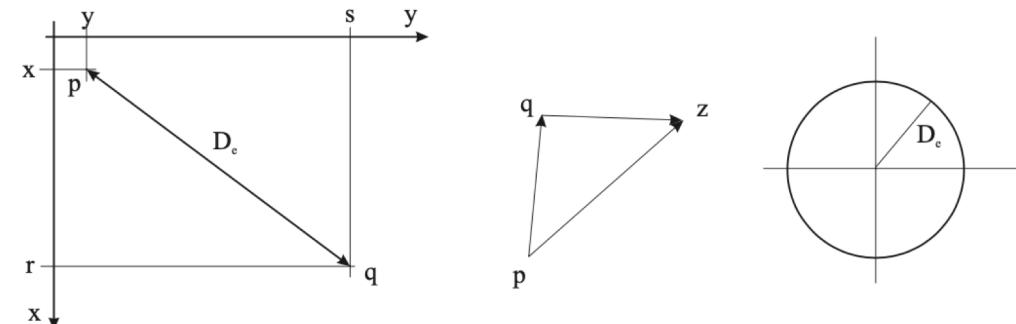
- Geometric aspects: Distance measures

- Measure of the distance $d(p,q)$ of two points $p=(x,y), q=(r,s)$ in a square grid.

- Simple metric: Euclidean distance with

$$d_e(p, q) = \sqrt{(r - x)^2 + (s - y)^2}$$

- Direct neighbors $q \in N_4$ of p have distance $d_e = 1$ (standardized grid size with edge length of one unit length)
 - Diagonal neighbors $q \in ND$ then have Euclidean distance $d_e = \sqrt{2}$.



Distances

- (Discrete) metrics must satisfy the triangle inequality;
- with the image points $p = (x,y)$, $q = (r,s)$ and $z = (u,v)$ we obtain a proper distance function or metric D_x when

$$D_x(p, q) \geq 0$$

$$D_x(p, q) = 0 \text{ if } p = q$$

$$D_x(p, q) = D_x(q, p)$$

$$D_x(p, z) \leq D_x(p, q) + D_x(q, z)$$

Distances

- This allows simpler (integer number) metrics to be defined for image processing:
- **D₄-distance**: D₄ distance or absolute metric of points p=(x,y),q=(r,s) is the shortest N₄ path between p and q:

$$d_4(p,q)=|r-x|+|s-y| .$$

Points of the same distance lie on the contour of a rhombus

- **D₈-distance**: D₈ distance or checkerboard metric ('city block') of points p,q is shortest N₈ path between p and q:

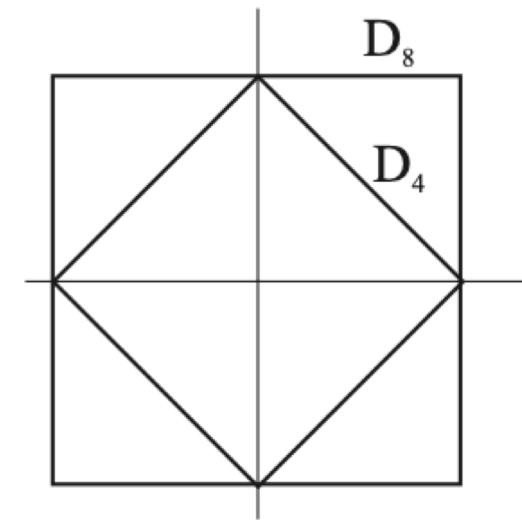
$$d_8(p,q)=\max(|r-x|,|s-y|) .$$

Points of equal distance lie on the contour of a square

D₈ and D₄

4	3	2	3	4
3	2	1	2	3
2	1	0	1	2
3	2	1	2	3
4	3	2	3	4

2	2	2	2	2
2	1	1	1	2
2	1	0	1	2
2	1	1	1	2
2	2	2	2	2



Topological aspects

- Topological aspects of the neighborhood of a pixel:
- **Neighboring pixels** are examined to determine **relationships**: *is a pixel “similar” and “neighboring”?*
- **Similarity conditions** can be defined by a predicate P (e.g. P : gray values f in an interval $V = [f_{min}, \dots, f_{max}]$).
 - In practise, one compares a pixel p with the pixels q from the neighborhood $N(p)$: is similarity criterion fulfilled? ($P = \text{true} \mid f(p), f(q) \in V$). If yes, then the found pixels p, q form a N -connection (**connected component**).
- Connections can be defined by means of so-called adjacencies ($r = 4, 8, m$):

Adjacency

- 4-adjacency: Two pixels p and q with values within V are 4-adjacent if q is in $N_4(p)$.
- D-adjacency: Two pixels p and q with values within V are adjacent if q is in $ND(p)$.
- 8-adjacency: Two pixels p and q with values within V are 8-adjacent if q lies in $N_8(p)$.
- m-adjacency: Two pixels p and q with values within V are adjacent if
 - 1. q is in $N_4(p)$ or
 - 2. q is in $ND(p)$ and the set $N_4(p) \cap N_4(q)$ has no pixels with values from V
- The m adjacency avoids inconsistencies (e.g., multiple assignments in the case of 8 adjacency).

Adjacencies of pixels

1	...	1	0
	:		
0	1	0	
	:		
0	1	...	1

0	1	...	1
	:		
0	1	0	
	:		
1	0	0	

0	1	...	1
	:		
0	1	0	
	:		
1	0	0	

Connected components

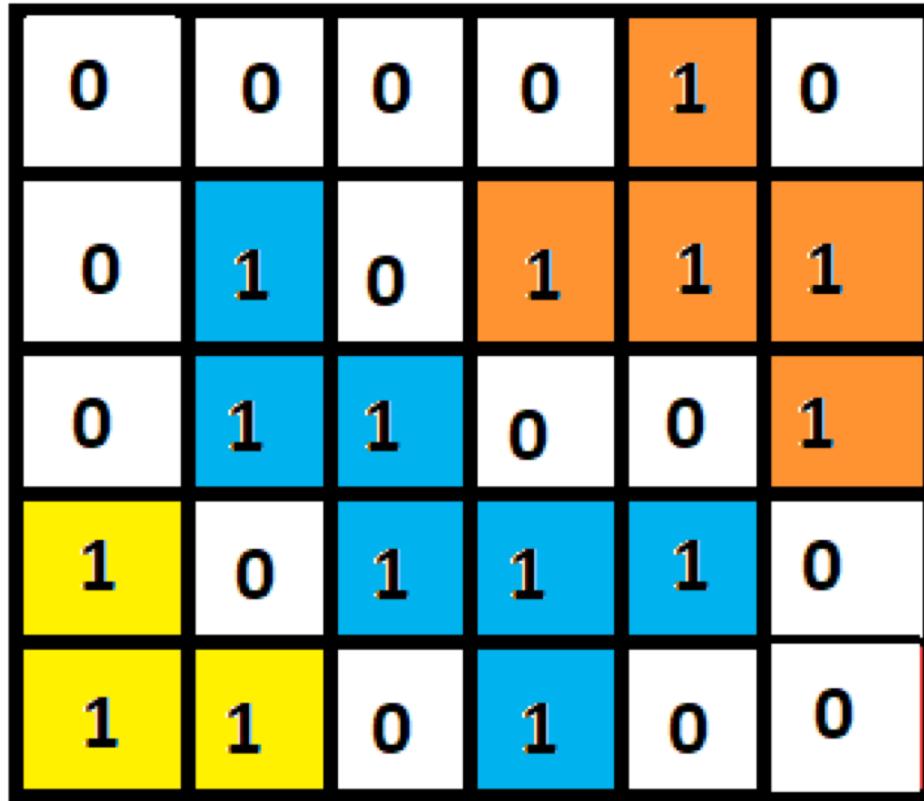
Digital path: A digital path $S = \overline{pq}$ between two points $p=(x,y)$ and $q = (r,s)$ is given by the sequence of certain points:

$$S = \overline{pq} = [(x_0, y_0), (x_1, y_1), \dots, (x_n, y_n)]$$

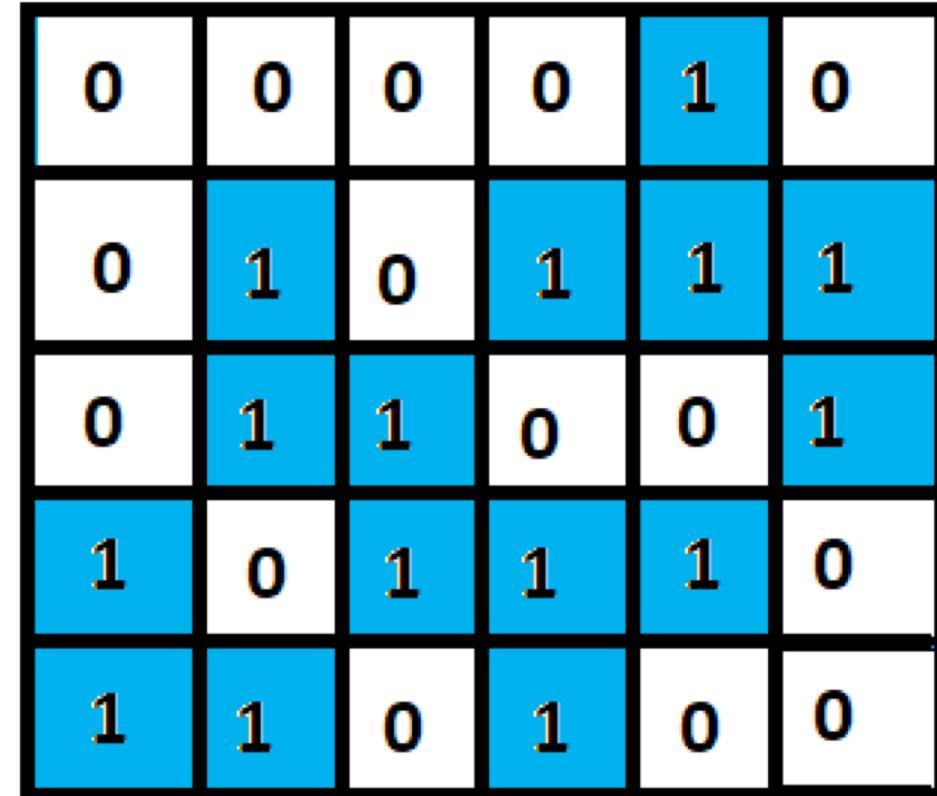
where $p=(x,y)=(x_0,y_0)$ and $q = (r,s)=(x_n,y_n)$ and the pixels (x_{i-1},y_{i-1}) and (x_i,y_i) are adjacent (4,8 or m) for $1 \leq i \leq n$.

Connected component: let S be a set of pixels of an image, then two pixels p,q are connected in S if there is a path \overline{pq} that is entirely in S . For each pixel p in S , the set of all pixels in S that are connected to it is called a connectivity component of S .

Connectivity induced by a neighborhood relation



4-connectivity map



8-connectivity map

Connected components II

- **Connectivity relation Z**

$$Z : pZq \Leftrightarrow \exists \overline{pq}$$

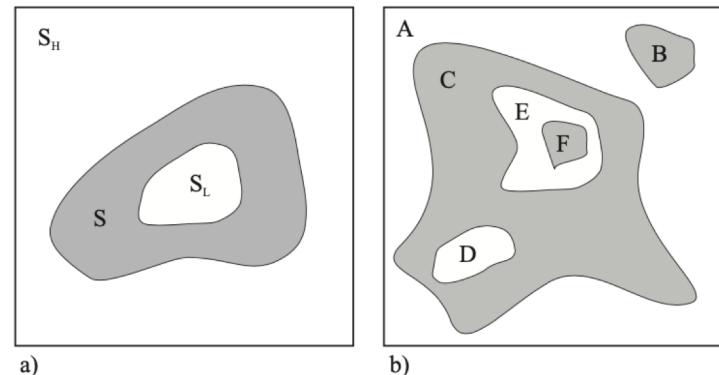
$$S_Z : \{p, q \mid pZq\}$$

$$pZq : \text{ if } \exists \overline{pq} \text{ with } \overline{pq} \in S$$

- Z is an equivalence relation; thus, pZp (reflexive) holds,
- $pZq \rightarrow qZp$ (symmetric) and
- $pZq \wedge qZu \rightarrow pZu$ (transitive).
- If there is only one correlation component, then it is called a correlation set.

Euler number and adjacency tree I

- Simple description of segmented images: Objects and background each represented as **connected components**.
- An image then contains the components
 - S (object or silhouette),
 - S_H (background, touching the edge)
 - S_L (hole in the object).
- It holds: $\neg S = S_B \cup S_{\#l}$, image $B = S \cup \neg S$
(segmentation is a *partition*), $S \subset B$ and $S \cap \neg S = 0$.



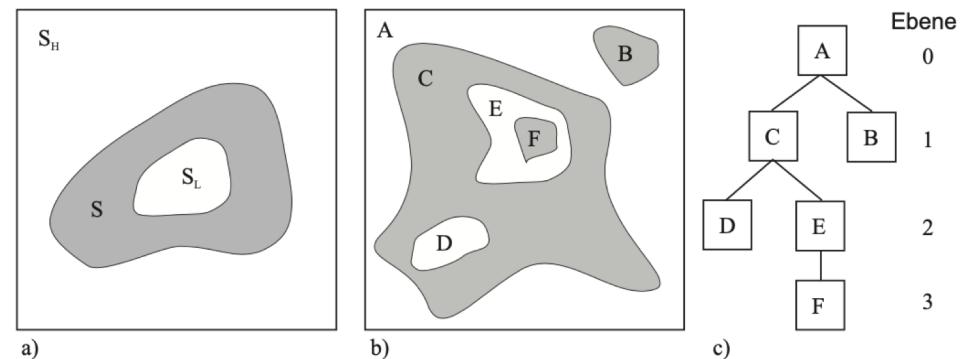
Euler number and adjacency tree II

- **Surrounding relation U:** The relation U says that the silhouette S "surrounds" the hole S_L "surrounds":

- $S \cup S_L \cup S$ surrounds S_L ,
 - U is a partial relation,
- $RUS \rightarrow S \neg U R$
- $RUS \wedge SUT \rightarrow RUT$ (transitive) .

- **Adjacency tree:** the relation U can be **described as a tree** with the following level numbering.

- Level 0 : Image B,
- odd levels : components S_i
- even levels : holes S_{L_i}



Euler number and adjacency tree III

- **Euler number:** A characteristic quantity is the Euler number E, which is the difference of the number C of components minus the number H of holes:

$$E = C - L$$

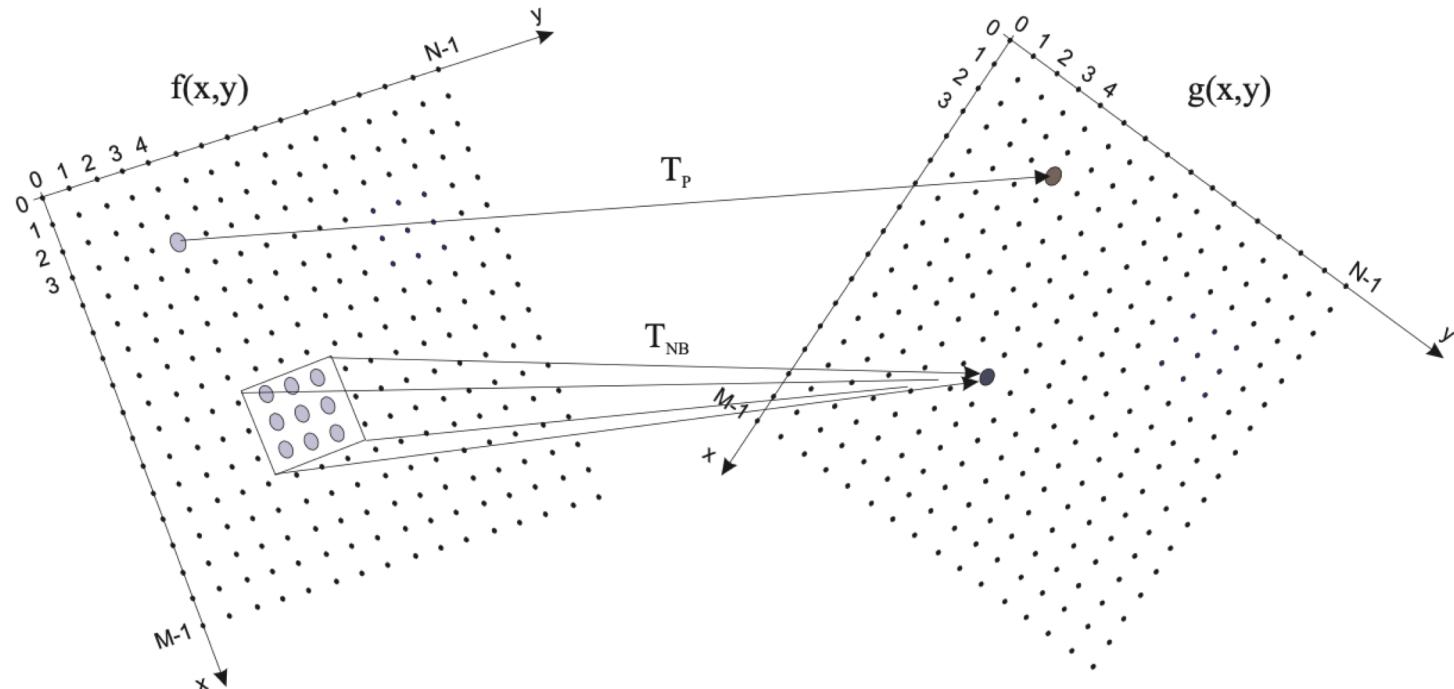
- E is a topological property.

Elementary image operations

- Examples of image operations
 - point and local image operations,
 - notice that local operations might be applied to the global image ☺
 - -> cf. convolutional neuronal networks
 - simple gray value manipulations (intensity transformations),
 - local convolution operation as generic image operation
 - -> cf. convolutional neuronal networks
 - simple geometric transformations
- Note: Transformations from spatial to frequency domain (FFT) and subsequently application of operators in frequency domain are possible but out of scope of this class. Images are Signals ☺ (and we can consider Signals as Images)

Point and neighborhood operations

- **Point operations** T_p and **neighborhood operations** T_{NP} are very important operation groups.
 - Mainly linear or nonlinear value range transformations
- **Geometric operations** are often a mixture of the two.

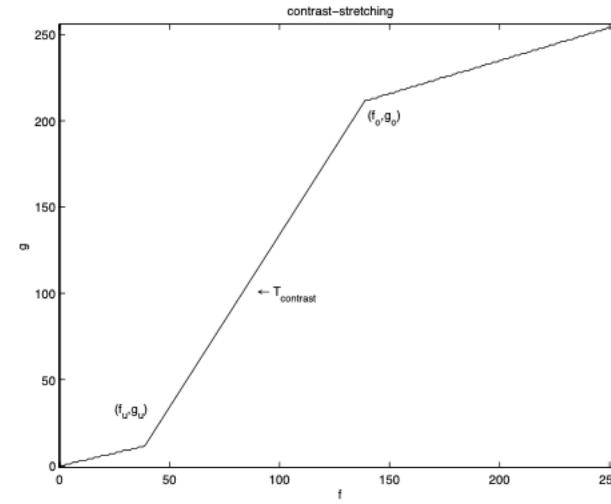
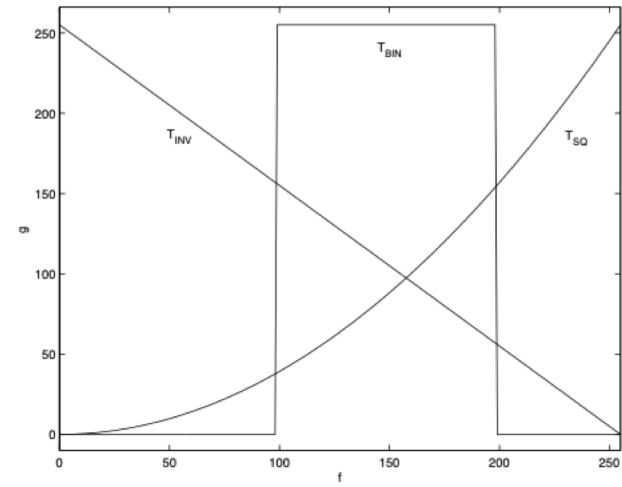


Point operations I

- **Point operations:** Value **range** (intensity) transformations in the form

$$T_p : f(r, s) \mapsto g(x, y)$$

with T_p : any linear or non-linear transformation $g = T_p(f)$



- Typical applications are **image enhancement** (mostly *subjective* but also *objective*) and also **segmentation** (binary/enumerable image generation).

Point operations II

- Examples:
- **Inversion**: Brightness values mirrored at the center of the value range:

$$T_{INV} : g := (L - 1) - f \quad , f \in [0, L - 1]$$

- **Binarization**: Brightness values in the interval $[f_u, f_o]$ are set to "1" - otherwise to "0".

$$T_{BIN} : g := \begin{cases} 0 & | \quad 0 \leq f < f_u \\ 1 & | \quad f_u \leq f \leq f_o \\ 0 & | \quad f_o < f \leq L - 1 \end{cases}$$

Point operations II

- or simply also by means of threshold Θ :

$$T_{BIN} : g := \begin{cases} 1 & | \quad f \geq \Theta \\ 0 & | \quad \text{sonst} \end{cases}$$

- **Point-operations** consider each “point” of the image range *individually*

Point operations III

- **Contrast enhancement:** Gray values are stretched (*spread*) or compressed in certain areas.

$$T_{contrast} : g := \begin{cases} c_1 * f & | \\ c * f & | \\ c_2 * f & | \end{cases} \begin{array}{l} f \leq f < f_u \\ f_u \leq f \leq f_o \\ f_o < f \leq L - 1 \end{array}$$

with $c_1, c_2 < 1, c > 1$

(general: $c > 1$: spreading; $c < 1$: compression)

Local or neighborhood operations I

- Local or so called **neighborhood operations** consider the function values $f(q)$ of the points q in the *neighborhood* of a pixel $p = (x,y)$ to generate a new function value g at the point p' (usually: $p' = p$).
- for this the definition of a neighborhood $N(p)$ is needed(e.g.N4,N8 or arbitrary forms for N)

$$T_{NB} : \mathbf{f}(N(x,y)) \mapsto g(x,y)$$

Local or neighborhood operations I

- **Example 1:** Average over neighborhood $N(p)$ as N_4 neighbor points to the point p

$(x, y) = T_{NB_M}(f(N_4(x, y)))$ respectively

$g(x, y) = 1/4 \sum_{\forall(r,s) \in N_4(x,y)} f(r, s)$ and tendered to

$g(x, y) = 1/4(f(x, y - 1) + f(x, y + 1) + f(x - 1, y) + f(x + 1, y))$

Local or neighborhood operations II

- **Example 2:** Pairwise gray value *difference* over the horizontal and vertical neighboring points in N_4 , respectively.

$$g(x,y) = T_{NB_D}(f(N_4(x,y)))$$

$$g(x,y) = (f(x,y-1) - f(x,y+1)) + (f(x-1,y) - f(x+1,y))$$

- Replacing the signs by factors (± 1) , we obtain T_{NB_D} in sum form to:

$$g(x,y) = (+1)f(x-1,y) + (-1)f(x+1,y) +$$

$$(+1)f(x,y-1) + (-1)f(x,y+1)$$

or compact to

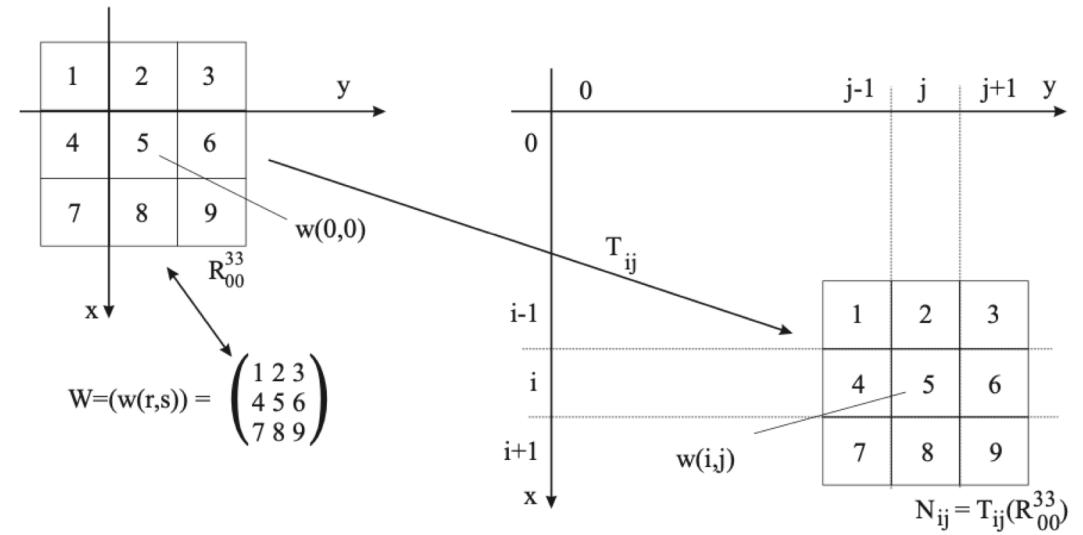
$$g(x,y) = \sum_{\forall(r,s) \in N_4(x,y)} w(r,s) \cdot f(r,s)$$

with $w(r,s)$: Weight factor or coefficients for neighboring point

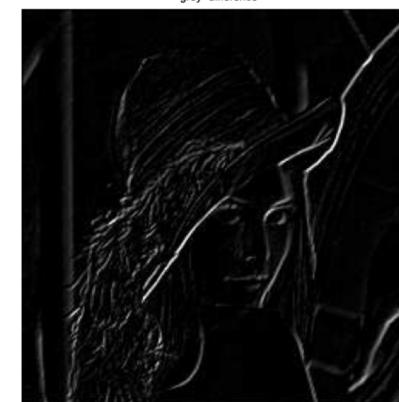
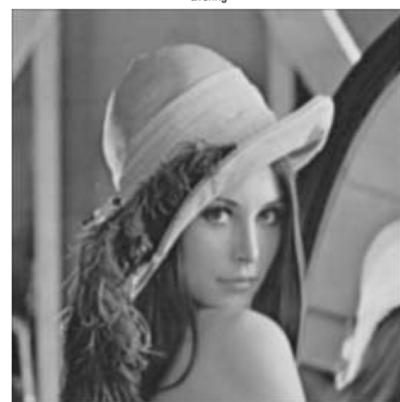
$$q = (r,s) \in N_4, (p) = N_4(x,y))$$

Local or neighborhood operations III

- Base neighborhood with weight values



- Result of the local gray value averaging / difference (see above).



Local or neighborhood operations IV

Typical applications for neighborhood operations

- **local image enhancement**,
 - e.g. *noise reduction* by *local averaging*
 - e.g. *contrast enhancement* based on local gradient
 - e.g. *edge enhancement* by means of Laplace operator
- all forms of **local filtering** (linear, nonlinear) in the spatial domain
- any (special) forms of local operators for segmentation, context analysis etc.

Filtering in local area I

- **Operation of the local area filters:**

A filter mask $w(r,s)$ slided over the image $F=f(x,y)$ column by column and row by row.

For each pixel (x,y) the filter response $g(x,y)$ is calculated linking the filter coefficints $w(r,s)$ with the (underlying) pixels $f(r,s)$

- For linear filters the filter response $G=g(x,y)$ results in

$$g(x,y) = \sum_{\forall(r,s) \in N(x,y)} w(r,s)f(r,s)$$

- This can also be formulated explicitly for a mask W of dimension $m \times n$ via formulate a double sum

$$g(x,y) = \sum_{r=-a}^a \sum_{s=-b}^b w(r,s)f(x+r, y+s)$$

- with $a = (m - 1)/2$ and $b = (n - 1)/2$ if m, n odd

Filtering in local area II

- The **operation is performed point by point**, i.e. for an image of the dimension $M \times N$ there are MN operations (computational time in $O(MN)$!)
- The shifted element of the local neighborhood operation, describing the function $w(x,y)$ is **referred to as filter, mask, kernel, windows or template**.
- The discrete elements describing the filter are called **coefficients**
- Q: *Do you see a connection to convolutional neural networks?*

Filtering in local area III

- Examples:

Contour enhancement with modified Laplace operator (TL):

$$W_{T_L} := \begin{pmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$



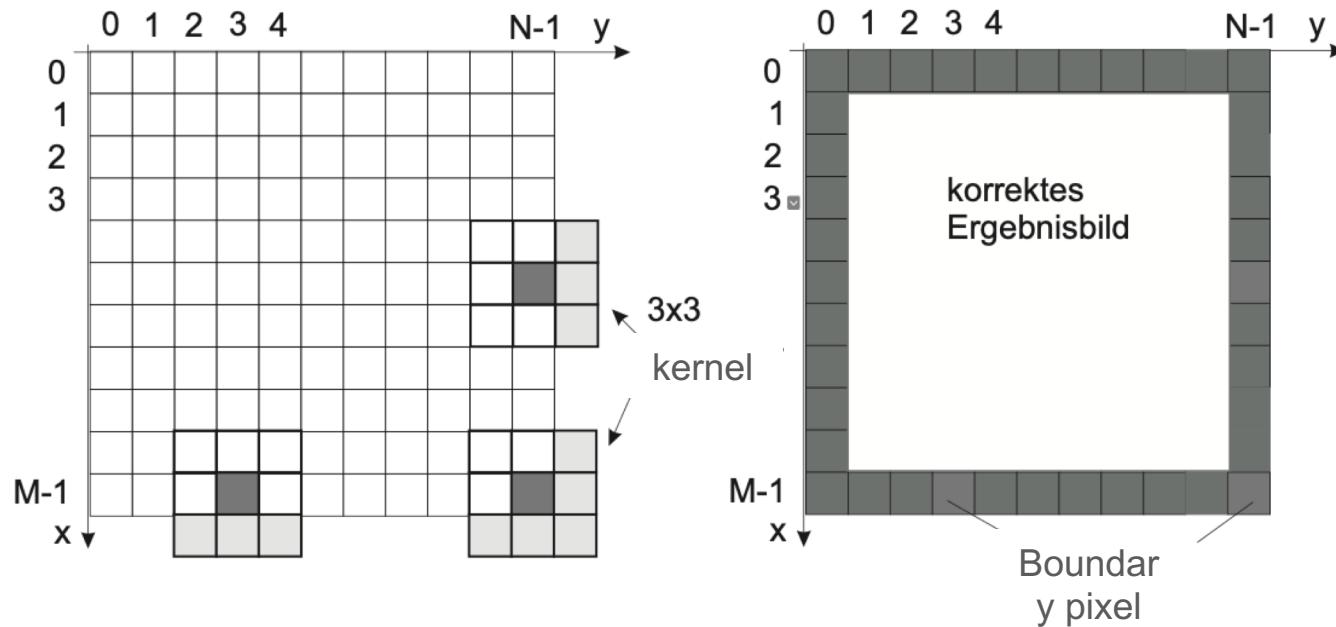
Weighted smoothing (\approx Gauss):
Filter masks WTG

$$W_{T_G} := \frac{1}{16} \begin{pmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{pmatrix}$$



Filtering in local area IV

- The local operations result in **boundary pixels problems**: overlapping of W with F at all positions q in $M \times N$ image not possible!



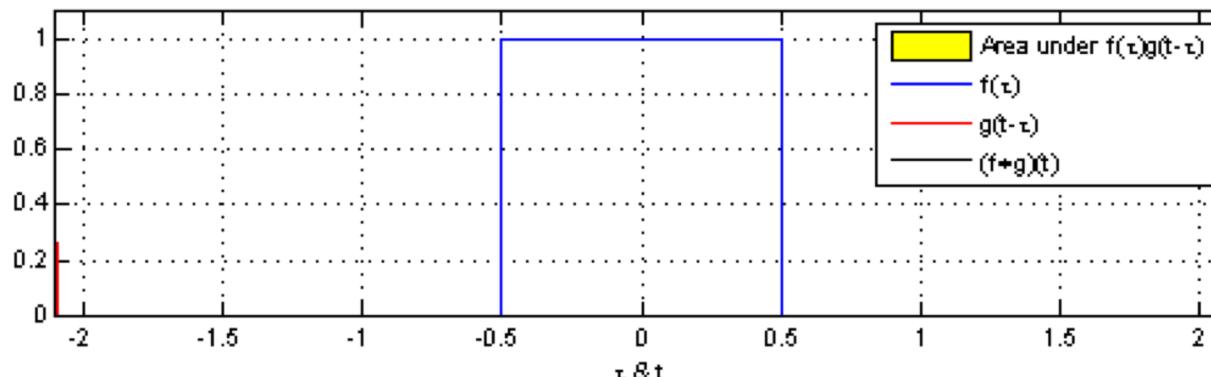
- Solutions:** *Padding, Mirroring of F at the boundaries, and more. Application dependent!*
- Q: Do you see a connection to convolutional neural networks?*

Convolution theorem I

- Linear filtering, as discussed before, can be described as the convolution of two functions $f(x,y)$ and $g(x,y)$ of the dimension $M \times N$ with the resulting convolution sum

$$f(x,y) * g(x,y) = \frac{1}{MN} \sum_{r=0}^{M-1} \sum_{s=0}^{N-1} f(r,s)g(x-r,y-s)$$

- this corresponds in the *frequency domain* to the point-wise multiplication of the Frequency functions $F(u,v)$ with $G(u,v)$



Convolution theorem I

- **Convolution theorem:** Convolution of two functions $f(x), w(x)$ in the *spatial domain (time domain)* corresponds to a multiplication of their Fourier transforms $F(u), G(u)$ in the *frequency domain*

$$f(x) \star g(x) \Leftrightarrow F(u) \cdot G(u)$$

- Mirroring of $g(x,y)$ is often ignored in BV practice, since $g(x,y)$ is mostly symmetric

End Lecture 2 (short) – following slides context was covered in exercise, slides here for reference and self study

Image enhancement (introduction)

- **Subjective image enhancement** - for the human viewer, so that he or she can, for example, more easily and reliably grasp the image content.
- **Objective image enhancement** - for machine evaluation, (e.g. noise suppression or contour sharpening).

Example (subjective)

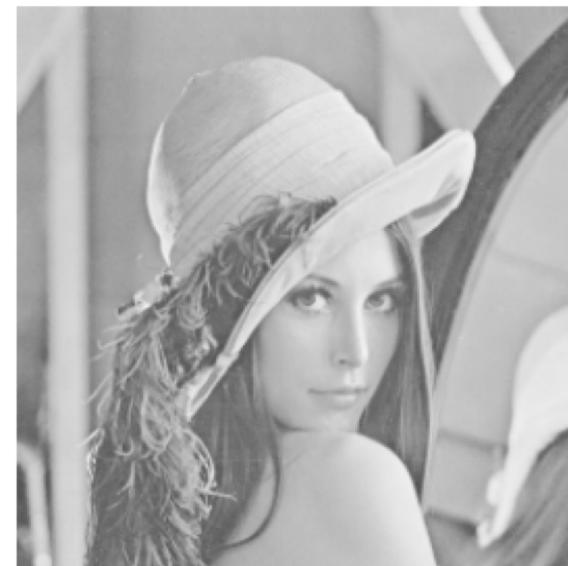
- **Subjective** image enhancement, e.g. gamma correction

$$f(x, y) = c \cdot b^\gamma(x, y) \quad c, \gamma \in \mathbf{R} > 0$$

original



LUT-transform $\gamma = 0.5$



Q.: How could we implement this in a line of MATLAB code?

Example (objective)

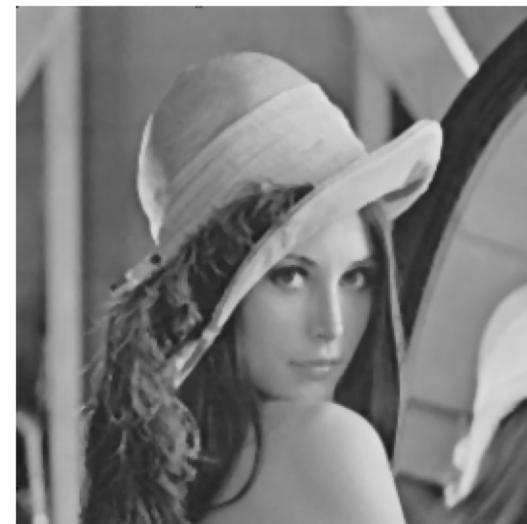
- **Objective Image enhancement**, e.g. nonlinear noise reduction

$$f(x, y) = \text{MEDIAN}(b(r, s)_{\forall (r, s) \in N(x, y)}) \quad \text{with} \quad N(x, y) = R_{xy}^{33}$$

original + salt & pepper – noise



noisy image median filtered



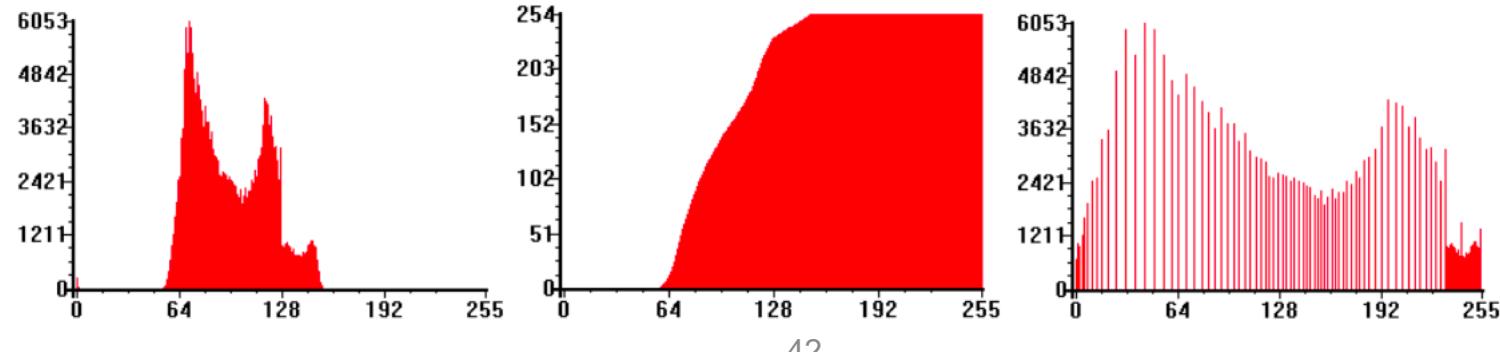
- Figure : Noise reduction (left: image with salt & pepper interference, right: median filtered image).

Histogram Equalization

- **Histogram Equalization:** Idea - using the full intensity range
- **Example:** Subjective Effect on Example Image



- **Algorithm:** Transformation Function from Histogram



Histogram Equalization I

- Given: *probability density* for f : $p_f(f)$, transformation function T :
 - Let T be unique and monotonically increasing between 0 and 1.
 - Let be $f, g \in R$ with $0 \leq g, f \leq 1$, then g is obtained to:

$$g = T(f) = \int_0^f p_f(\gamma) d\gamma \quad \text{mit } (0 \leq f \leq 1)$$

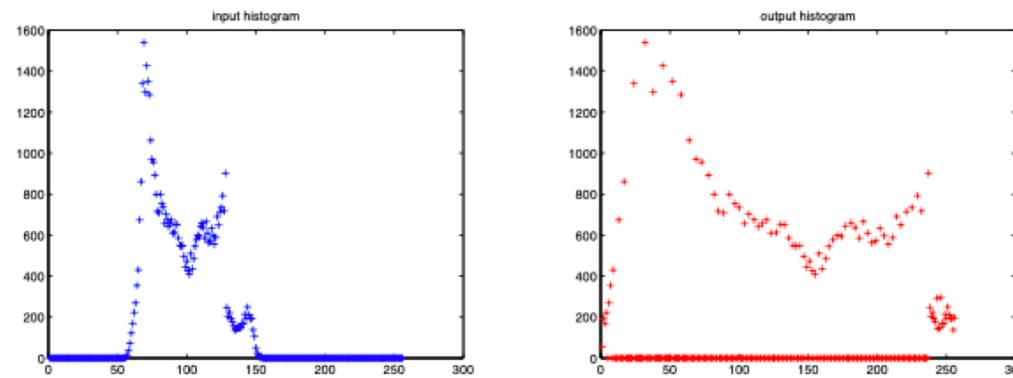
Digital practice I

- Instead of the **distribution density**, one considers the **relative frequency** of the gray values f_k :

$$p_f(f_k) = \frac{n_k}{n} \quad k = 0, 1, \dots, L-1$$

n_k = Frequency of gray values f_k

n = $M * N$ (Number of pixels in image)



Discrete histogram equalization (histo-leveling)

- The discrete gray values $g \in \mathbb{N}_+$, $g_{\min} \leq g \leq g_{\max}$ are obtained by means of a unary transformation $T_{HE} : f \rightarrow g$
- Result image $g(x,y)$ using point operation over all image coordinates (x,y) :

$$T_{HE} : g(x,y) = g(f(x,y)) = \text{ROUND} \left[(g_{\max} - g_{\min}) \cdot \sum_{\gamma=f_{\min}}^f h_f(\gamma) \right] + f_{\min}$$

- With a gray value range $0 \leq b,g \leq 255$ simplifies to

$$T_{HE} : g(x,y) = g(f(x,y)) = \text{ROUND} \left[255 \cdot \sum_{\gamma=0}^{f(x,y)} h_f(\gamma) \right]$$

- Rounding needed to map real values back to integer indexes for the “table lookup” of the CDF!

End Lecture 2 (short) – including topic introduced in exercise