

“Big O Notation”



Dr. Oscar Castro, Dr. Pierrick Pochelu

Summary

- General presentation
- Demo with Jupyter Notebook
- Data Structure and some ML algorithms
- Conclusion & next steps with Big O

Introduction

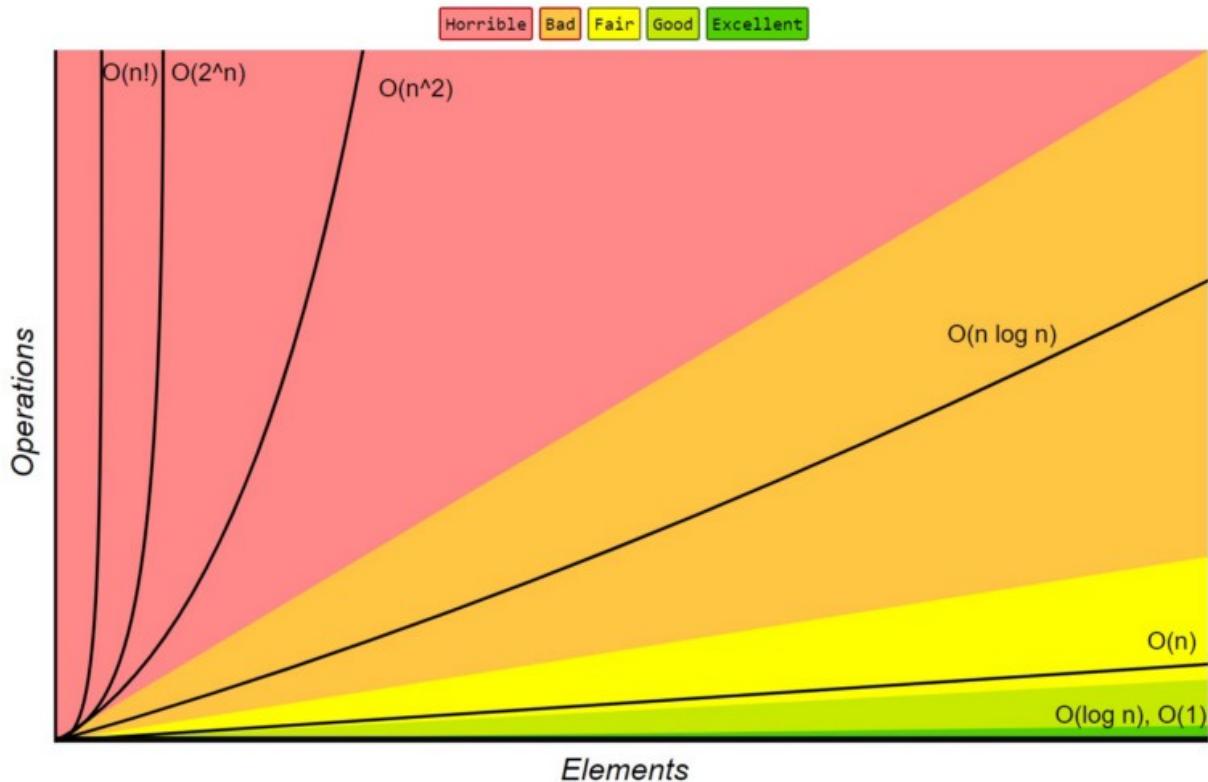
What is Big O Notation?

- "Big": Represents the *upper bound* of an algorithm's growth rate for large input sizes.
- "O": Stands for *Order of Magnitude*, describing how the runtime or space requirements increase with input size.

Why is Big O Important?

- **Analyzes Performance Trends:**
 - Worst-case performance scenario as input size grows.
 - Provides a **theoretical analysis**, (does not provide an accurate measurement #iterations, #seconds).
 - May be used for different aspects: computing time, memory, data transfer, ...
- **Evaluates Algorithm Efficiency:**
 - Predicts how algorithms will scale with larger inputs.
 - Helps in comparing different algorithms to choose the best one.
 - Optimizing existing algorithms (bottleneck identification, improvement)
- **Key Questions It Helps Answer:**
 - Is the algorithm's performance **feasible** for large data?
 - Does it justify the **computational cost**?
 - How much **computing power** (e.g., # of cores) is required for reasonable performance?

Big-O Complexity Chart



In this order:

- O(1) constant (or “O(C)”)
- O(log(n)) logarithmic
- O(n) linear
- O(n^C) polynomial (“quadratic” when C=2, “cubic when C=3”)
- O(C^n) exponential
- O(n!) factorial

$\Rightarrow O(n^m)$

Putting in perspective Big O trends

We vary `n` and observe the effect on complexity estimation

Constant (C)	n	$\log_2(n)$	n	$n * \log_2(n)$	n^2	2^n	n!
1	10	3.3219	10	33.2193	100	1024	3,628,800
1	20	4.3219	20	86.4386	400	1,048,576	243,290,200,817,664,000

Big O Notation in ML projects: When and Why ?

1. Large-Scale Datasets

- **What happens if I increase the number of data samples?**

- Can the model process the dataset within a reasonable time?
- Should we subsample the data for faster processing?

- **What happens if I increase the number of features?**

- Analyzes the impact of high dimensionality on computational cost.
- Guides the need for dimensionality reduction (e.g., PCA) or feature selection.

2. Large-Scale Models

- **What happens if I increase the model size?**

- Should we use a more efficient model or opt for distributed training?
- Is the computational cost of increasing model parameters justifiable?

3. Specific Performance Needs

- **Real-time applications** (e.g., real-time inference for video streaming, online decision-making):

- Can the algorithm meet latency requirements?

- **Low-capacity hardware** (e.g., drones, mobile devices):

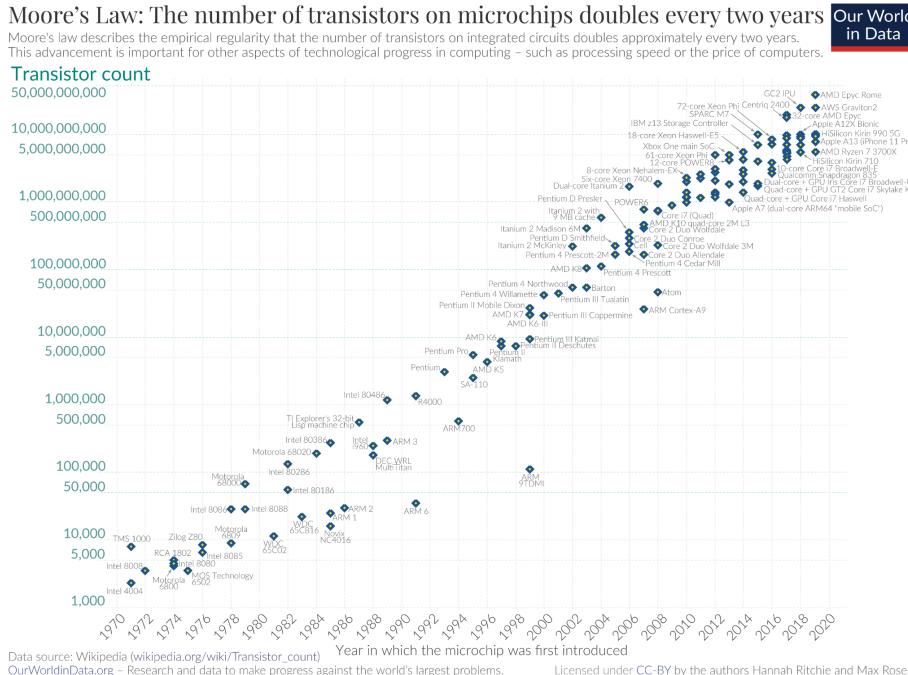
- Is the data and model lightweight enough to run efficiently on limited hardware?

Big O notation a timely topic

2 current trends:

- **End of Moore's Law:** CPU speeds are no longer increasing at the exponential rate they once did.
- **Modern applications (Deep Learning, Big Data)** require massive computational resources.

Conclusion: Optimizing and Parallelizing **algorithms** is the only way to advance modern applications.



```
/* demo */
```

Built-in Python data structures

List

Copy	$O(n)$
Append[1]	$O(1)$
Pop last	$O(1)$
Pop intermediate[2]	$O(n)$
Insert	$O(n)$
Get Item	$O(1)$
Set Item	$O(1)$
Delete Item	$O(n)$
Iteration	$O(n)$
Get Slice	$O(k)$
Del Slice	$O(n)$
Set Slice	$O(k+n)$
Extend[1]	$O(k)$
Sort	$O(n \log n)$
Multiply	$O(nk)$
x in s	$O(n)$
min(s), max(s)	$O(n)$
Get Length	$O(1)$

Double-Ended queue
("from collections import deque")

Copy	$O(n)$
append	$O(1)$
appendleft	$O(1)$
pop	$O(1)$
popleft	$O(1)$
extend	$O(k)$
extendleft	$O(k)$
rotate	$O(k)$
remove	$O(n)$
Get Length	$O(1)$

Set

x in s	$O(1)$
Union s t	$O(\max(\text{len}(s), \text{len}(t)))$
Intersection s&t	$O(\min(\text{len}(s), \text{len}(t)))$
Multiple intersection s1&s2&..&sn	
Difference s-t	$O(\text{len}(s))$
s.difference_update(t)	$O(\text{len}(t))$
Symmetric Difference s^t	$O(\text{len}(s))$
s.symmetric_difference_update(t)	$O(\text{len}(t))$

Dictionary
("dict" keyword)

k in d	$O(1)$
Copy[3]	$O(n)$
Get Item	$O(1)$
Set Item[1]	$O(1)$
Delete Item	$O(1)$
Iteration[3]	$O(n)$

ML complexity

Conclusion:
Different
algorithms have
different time &
memory
complexity
⇒ Use it at your
advantage

Time Complexity of 10 Most Popular ML Algorithms		Training	Inference
	Linear Regression (OLS)	$O(nm^2 + m^3)$	$O(m)$
	Linear Regression (SGD)	$O(n_{epoch}nm)$	$O(m)$
	Logistic Regression (Binary)	$O(n_{epoch}nm)$	$O(m)$
	Logistic Regression (Multiclass OvR)	$O(n_{epoch}nmc)$	$O(mc)$
	Decision Tree	$O(n \cdot \log(n) \cdot m)$ $O(n^2 \cdot m)^*$ Worst case	$O(d_{tree})$
	Random Forest Classifier	$O(n_{trees} \cdot n \cdot \log(n) \cdot m)$	$O(n_{trees} \cdot d_{tree})$
	Support Vector Machines (SVMs)	$O(n^2m + n^3)$	$O(m \cdot n_{SV})$
	k-Nearest Neighbors	—	$O(nm)$

Memory

$O(m^2)$

$O(m^2)$

$O(m^2)$

$O(m^2)$

$O(n)$

$O(ntrees * n)$

$O(m^2)$

$O(n*m)$

n : samples m : dimensions n_{epoch} : epochs c : classes d_{tree} : depth
 n_{SV} : Support vectors k : clusters i : iterations

Conclusion & next steps with Big O

1. Useful for describing ML algorithms

- Describes how an algorithm's performance scales with input size
- Useful to take quick decisions on data structure, ML algorithm
- Useful also to describe complex cases where different factor

2. When a code is well optimized, Big O notation is not enough:

- Ignoring constant factors and lower-order terms
- Ignoring language performance , and individual operations (example: FP64 operations slower than INT32)
- The performance of algorithms can be affected by various CPU architectures, memory hierarchies, and parallel processing capabilities
- Big O notation with HPC methodologies
 - Big O notation of #messages in addition of time and memory consumption
 - Code Profiling: Computing time, Memory, ...
 - Strong Scalability Analysis (Fixing the input size, increase the #cores)
 - Weak Scalability Analysis (Increase the input size, increase the #cores)